Cisco live!

6-9 March 2018 · Melbourne, Australia

VRF, MPLS and MP-BGP Fundamentals

Jason Gooley, CCIEx2 (RS, SP) #38759

Twitter: @Jason_Gooley

LinkedIn: http://www.linkedin.com/in/jgooley

BRKCRT-2601

Cisco live!

Agenda

- Introduction to Virtualisation
- VRF-Lite
- MPLS & BGP Free Core
- Multiprotocol BGP (MP-BGP)
- Conclusion
- Q & A

Cisco Spark

Questions?
Use Cisco Spark to chat with the speaker after the session

How

1. Find this session in the Cisco Live Mobile App
2. Click “Join the Discussion”
3. Install Spark or go directly to the space
4. Enter messages/questions in the space

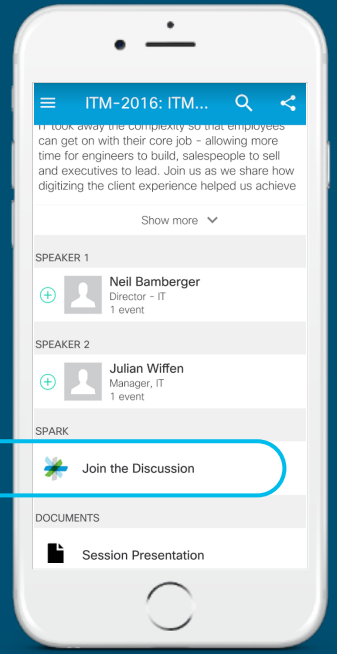

cs.co/ciscolivebot#BRKCRT-2601

E.g: session ID = BRKCRT-2601

3 Networks Walk into a…

What is a VRF?

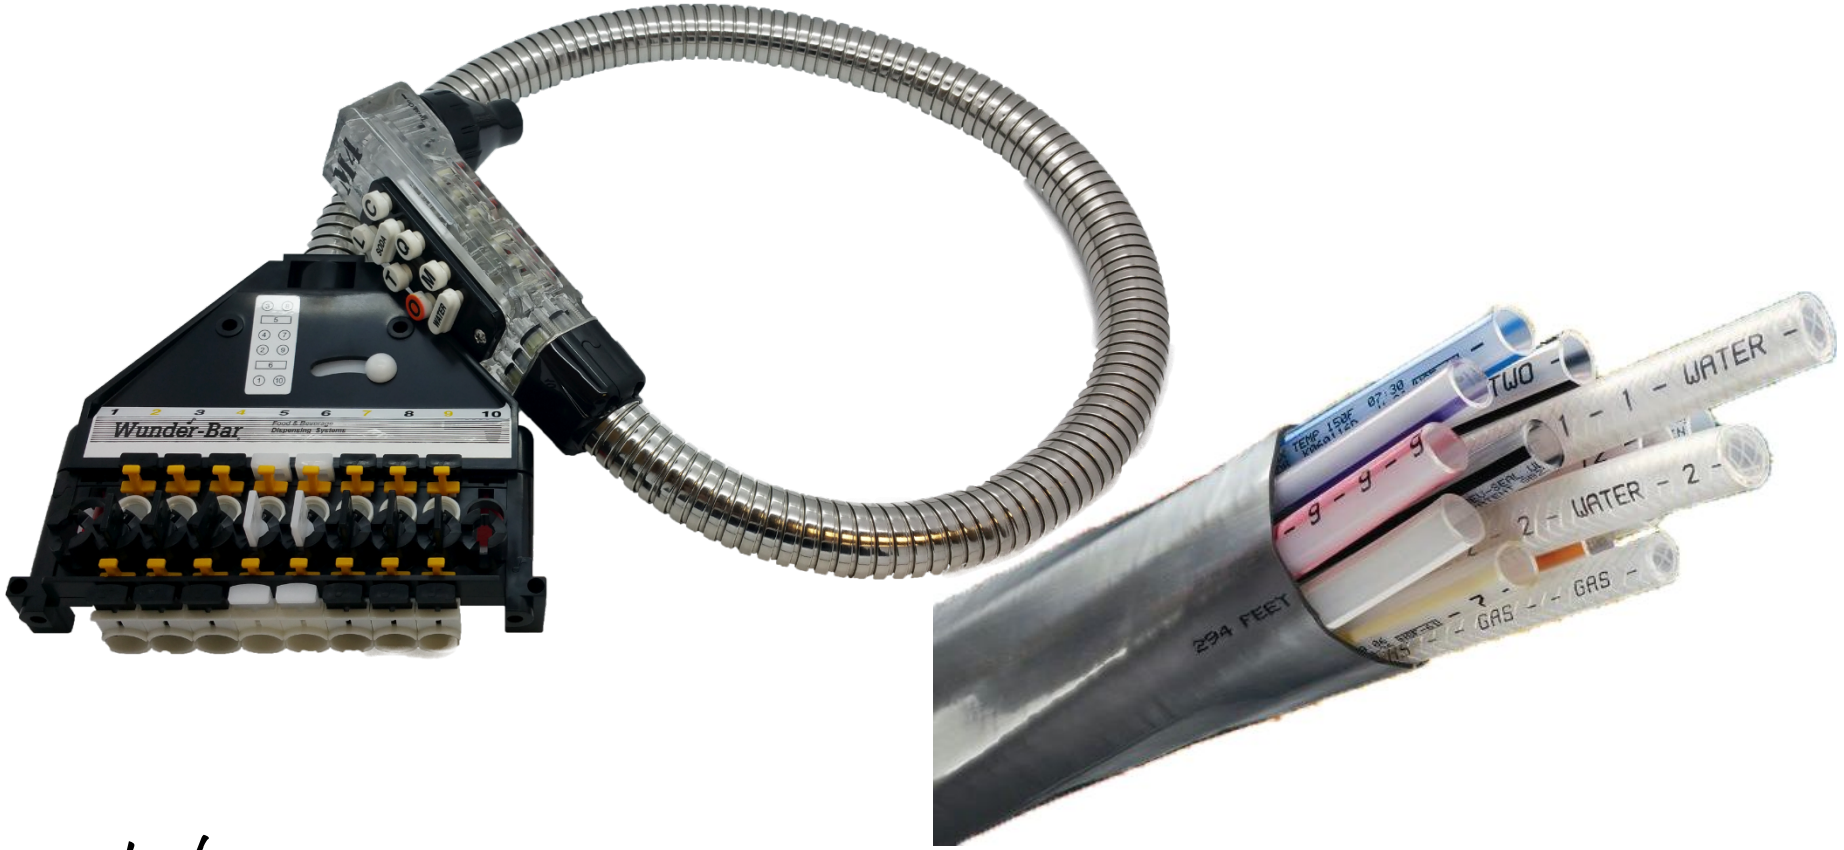

Enterprise Network Virtualisation

Key Building Blocks

Device Partitioning

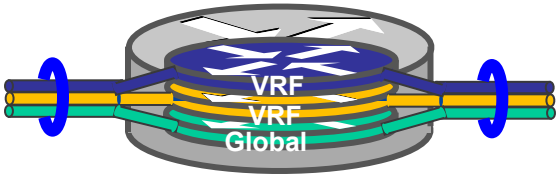

“Virtualising” the Routing and Forwarding of the Device

Virtualised Interconnect

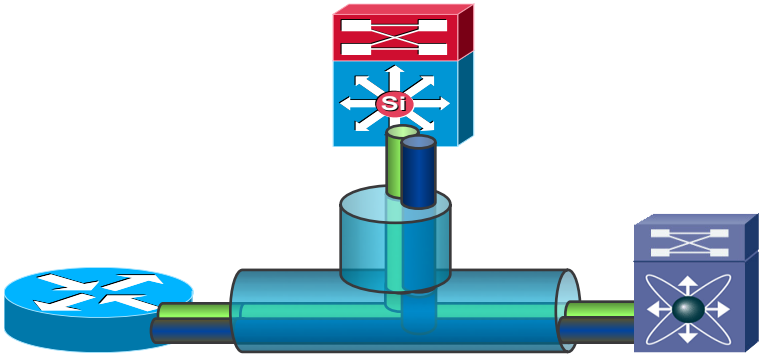

Extending and Maintaining the “Virtualised” Devices/Pools over Any Media

Device Partioning

Layer 2 vs. Layer 3 Virtualisation

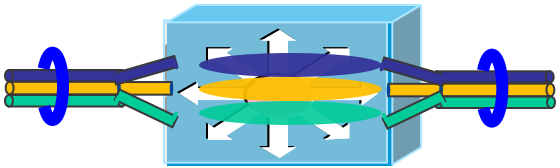

VLAN—Virtual LAN

VRF—Virtual Routing and Forwarding

Device Partioning

Layer 2 vs. Layer 3 Virtualisation

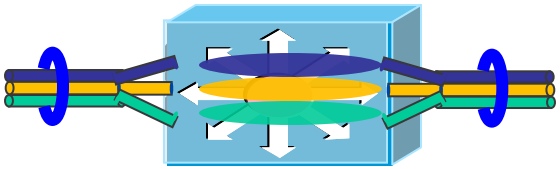

VLAN—Virtual LAN

- Virtualise at Layer 2 forwarding
- Associates to one or more L2 interfaces on switch
- Has its own MAC forwarding table and spanning-tree instance per VLAN
- Interconnect options?

  VLANs are extended via a physical cable or virtual 802.1q trunk

VRF—Virtual Routing and Forwarding

- Virtualise at Layer 3 forwarding
- Associates to one or more Layer 3 interfaces on router/ switch
- Each VRF has its own

  Forwarding table (CEF)

  Routing process (RIP, EIGRP, OSPF, BGP)
- Interconnect options (VRF-Lite)?

  802.1q, GRE, sub-interfaces, physical cables, signaling

Path Isolation
Functional Components

- Device Virtualisation
  - Control plane Virtualisation
  - Data plane Virtualisation
  - Services Virtualisation
- Data path Virtualisation
  - Hop-by-Hop - VRF-Lite End-to-End
  - Multi-Hop - VRF-Lite GRE
  - MPLS-VPN
  - MPLS VPN over IP
  - MPLS VPN over DMVPN
  - MPLS VPN o GRE/mGRE

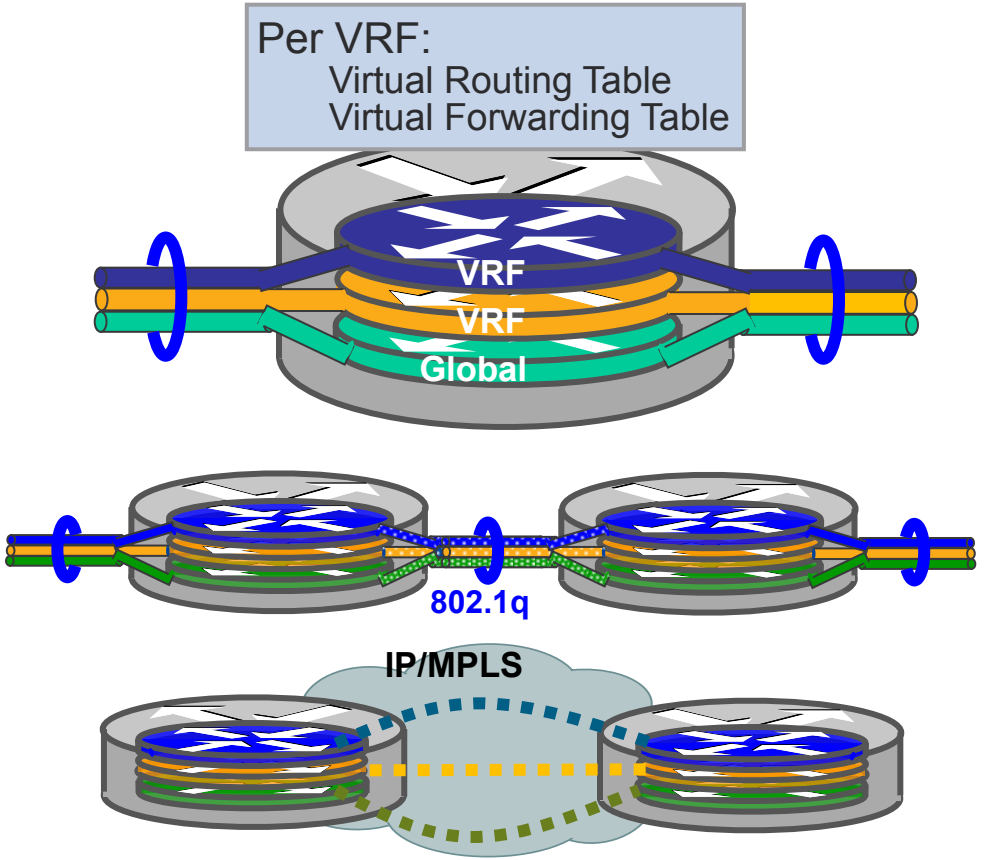

VRF-Lite

What is VRF-Lite?

Functional Components

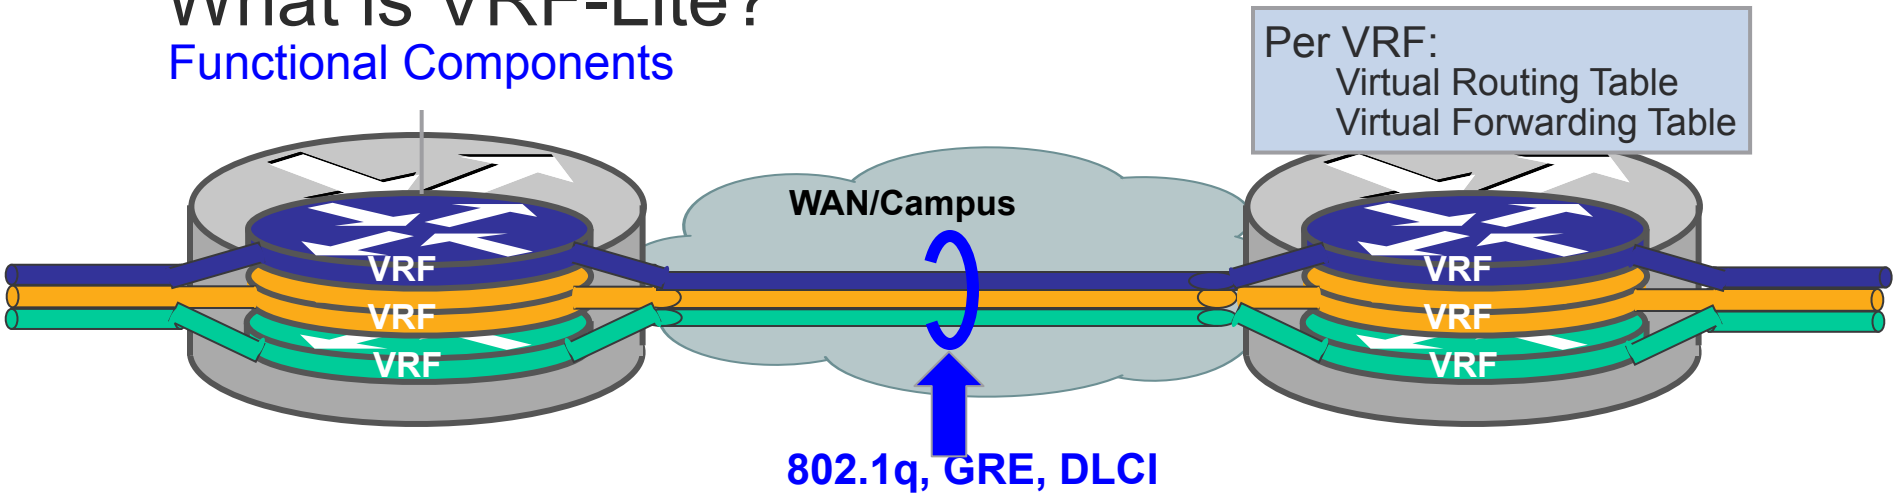

- A VRF supports it's own **Routing Information Base (RIB)** and **Forwarding Information Base (FIB)**
- Leverages "Virtual" encapsulation for separation:
  - Ethernet/802.1Q, GRE, Frame Relay
- Routing protocols are "VRF aware"
  - RIP/v2, EIGRP, OSPF, BGP, static (per VRF)
- Layer 3 interfaces can only belong to a single VRF

VRF-Lite
Things to Remember

- End-to-End segmentation is done on a per VRF and per hop basis
- MP-BGP or control plane signalling is not required
- Labels are not required (i.e. MPLS)
- Scaling should be limited to a small number of VRFs

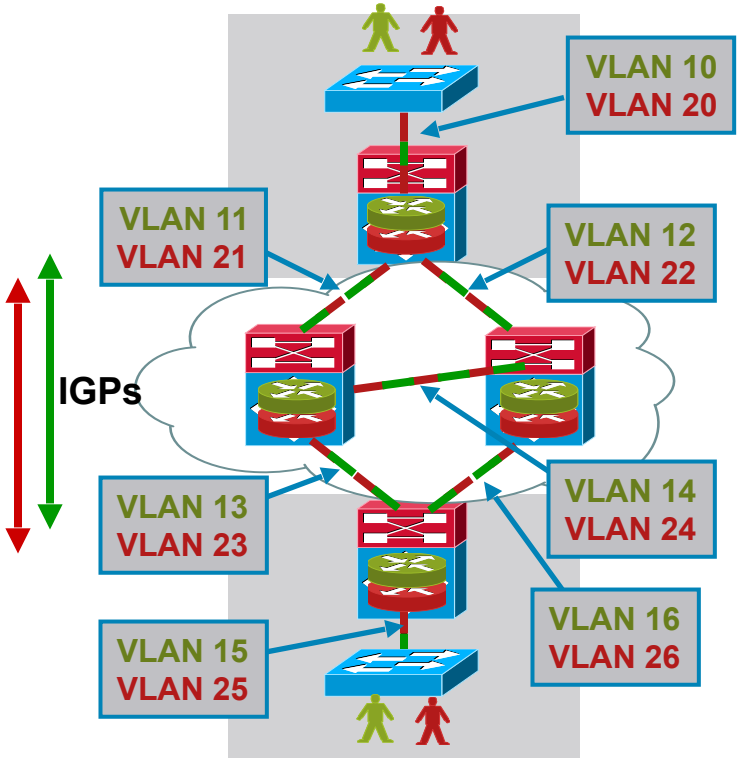

VRF-Lite
Sub-interface Example

Per VRF:
- Virtual Routing Table
- Virtual Forwarding Table
- Locally Significant

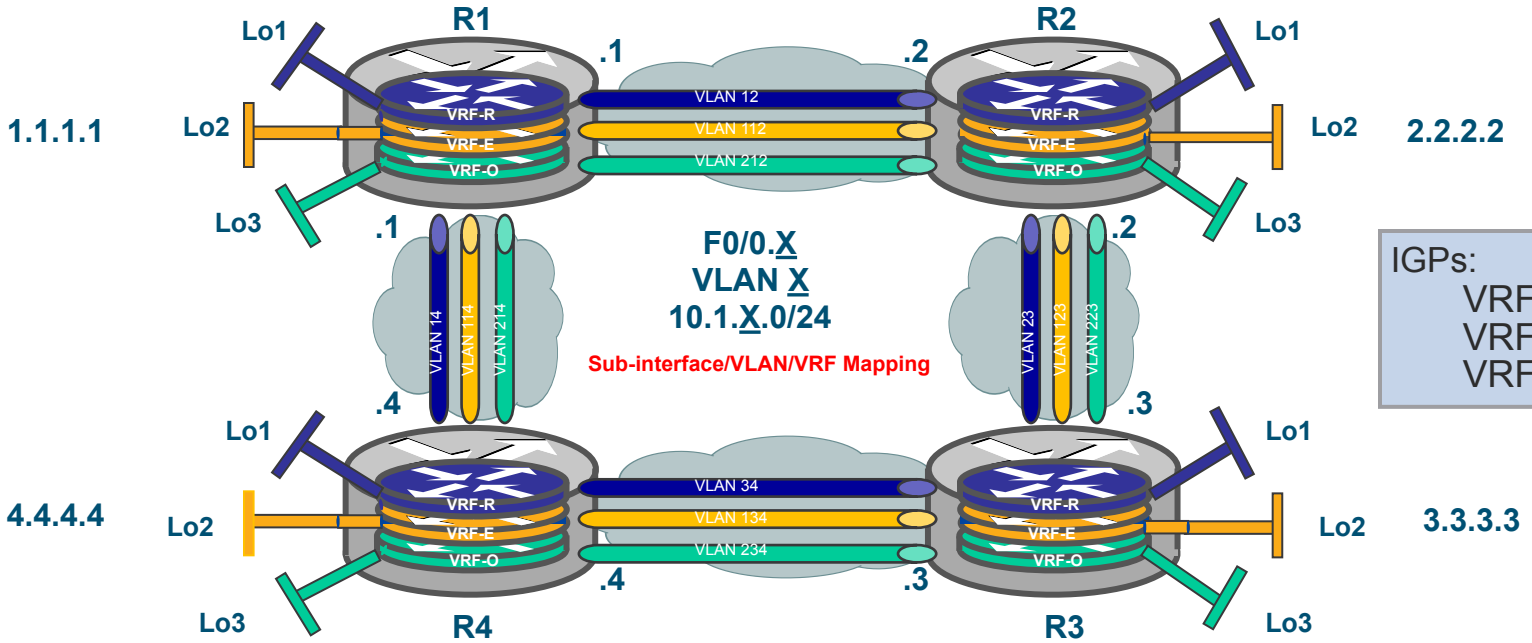

IGPs:
- VRF-R = RIP
- VRF-E = EIGRP
- VRF-O = OSPF

VRF-Lite Sub-interface Configuration

Command Line Interface (CLI) Review

```
ip vrf VRF-R
 rd 1:1

interface FastEthernet0/0.12
 ip vrf forwarding VRF-R

interface Loopback1
 ip vrf forwarding VRF-R
```

```
ip vrf VRF-E
 rd 2:2

interface FastEthernet0/0.112
 ip vrf forwarding VRF-E

interface Loopback2
 ip vrf forwarding VRF-E
```

```
ip vrf VRF-O
 rd 3:3

interface FastEthernet0/0.212
 ip vrf forwarding VRF-O

interface Loopback3
 ip vrf forwarding VRF-O
```

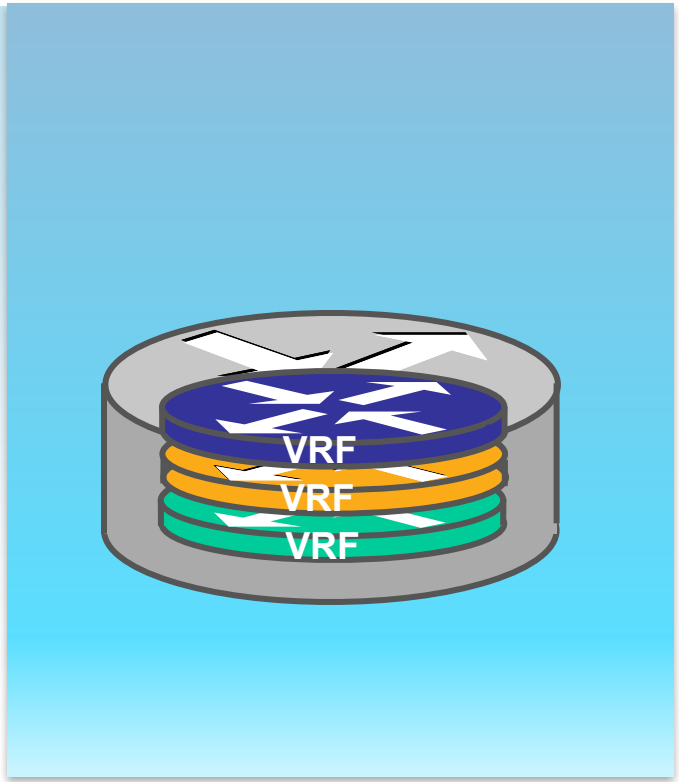

VRF-Lite Sub-interface Configuration

Command Line Interface (CLI) Review – VRF Definition Example

```
vrf definition VRF-R
 rd 1:1
 address-family ipv4

interface FastEthernet0/0.12
 vrf forwarding VRF-R

interface Loopback1
 vrf forwarding VRF-R
```

```
vrf definition VRF-O
 rd 3:3
 address-family ipv4


interface FastEthernet0/0.212
 vrf forwarding VRF-O

interface Loopback3
 vrf forwarding VRF-O
```

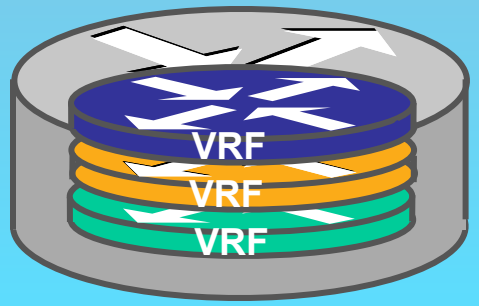

Multiprotocol VRF Conversion Configuration

Command Line Interface (CLI) Review

vrf upgrade-cli multi-af-mode {common-policies | non-common-policies} [vrf *vrf-name*]

```
PE1(config)#vrf upgrade-cli multi-af-mode common-policies
You are about to upgrade to the multi-AF VRF syntax commands.
You will lose any IPv6 addresses configured on interfaces
belonging to upgraded VRFs.

Are you sure ? [yes]:
Number of VRFs upgraded: 1
```

```
interface Ethernet0/1
 vrf forwarding VRF
 ip address 10.1.78.7 255.255.255.0

PE1(config)#do sh run | se vrf
vrf definition VRF
 rd 7:1
 route-target export 7:1
 route-target import 5:1
```

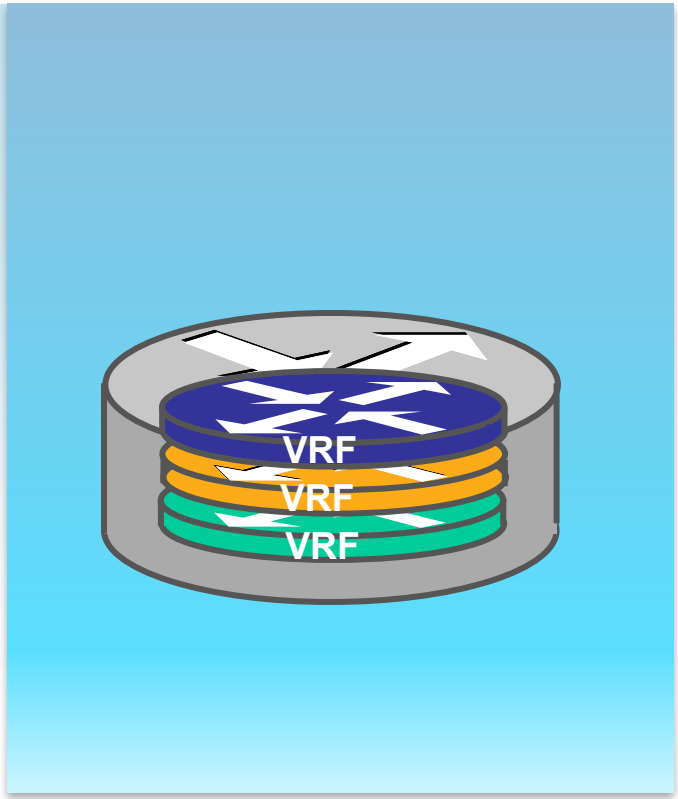

VRF Aware RIP Configuration

Command Line Interface (CLI) Review

```
router rip
 version 2
 network 1.0.0.0
 network 10.0.0.0
 no auto-summary
```

```
router rip
 !
 address-family ipv4 vrf VRF-R
  network 1.0.0.0
  network 10.0.0.0
  no auto-summary
  version 2
 exit-address-family
```

RIP leverages address-family ipv4 vrf _____

Leverage what you already know!

VRF Aware EIGRP Configuration

Command Line Interface (CLI) Review

```
router eigrp 10
 network 1.1.1.1 0.0.0.0
 network 10.1.112.0 0.0.0.255
 no auto-summary
```

```
router eigrp 10 (AS can be the same or different as one of the VRFs!!!)
 auto-summary
 !
 address-family ipv4 vrf VRF-E
  network 1.1.1.1 0.0.0.0
  network 10.1.112.0 0.0.0.255
  no auto-summary
  autonomous-system 10
 exit-address-family
```

EIGRP leverages address-family ipv4 vrf _____

Set unique autonomous system number per VRF

Leverage what you already know!

VRF Aware OSPF Configuration

Command Line Interface (CLI) Review

```
router ospf 1
 log-adjacency-changes
 network 1.1.1.1 0.0.0.0 area 1
 network 10.1.212.0 0.0.0.255 area 0
```

```
router ospf 2 vrf VRF-O
 log-adjacency-changes
 network 1.1.1.1 0.0.0.0 area 1
 network 10.1.212.0 0.0.0.255 area 0
```

OSPF leverages vrf _____ after the unique process number

Leverage what you already know!

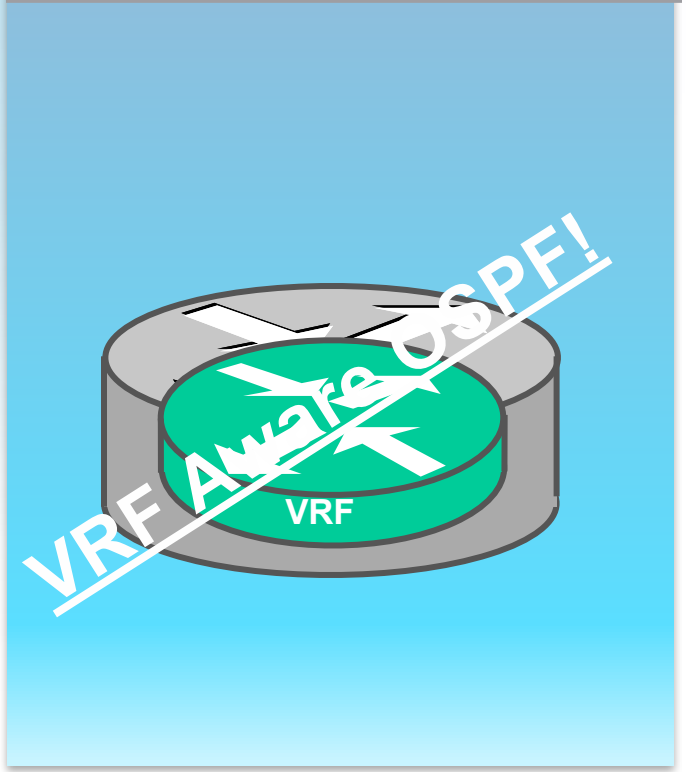

Live Exploration

No Sub-interface Support? No Problem!

GRE Example

- VRF-Lite can also leverage GRE tunnels as a segmentation technology
- Each VRF uses a unique GRE tunnel
- GRE tunnel interface is “VRF aware”

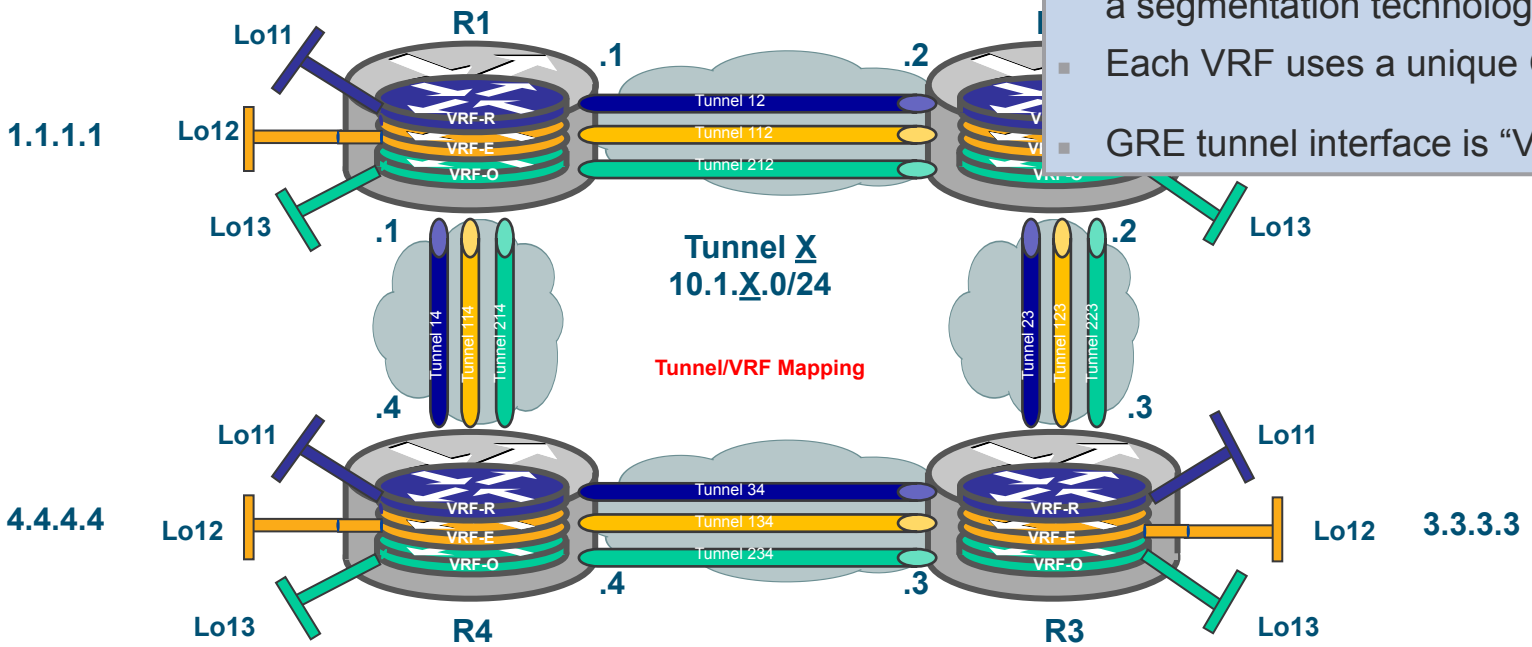

Configuration Note: Each GRE Tunnel Could Require Unique Source/Destination IP (Platform Dependent)

VRF-Lite Tunnel Configuration

Command Line Interface (CLI) Review

```
ip vrf VRF-S
 rd 11:11

interface Loopback101
 ip address 11.11.11.11 255.255.255.255 (Global Routing Table)

interface Tunnel12
 ip vrf forwarding VRF-S
 ip address 10.1.12.1 255.255.255.0
 tunnel source Loopback101
 tunnel destination 22.22.22.22
```

```
ip vrf VRF-S
 rd 22:22

interface Loopback102
 ip address 22.22.22.22 255.255.255.255 (Global Routing Table)

interface Tunnel12
 ip vrf forwarding VRF-S
 ip address 10.1.12.2 255.255.255.0
 tunnel source Loopback102
 tunnel destination 11.11.11.11
```

Leverage what you already know!

```
ip route vrf VRF-S 2.2.2.2 255.255.255.255 10.1.12.2
```

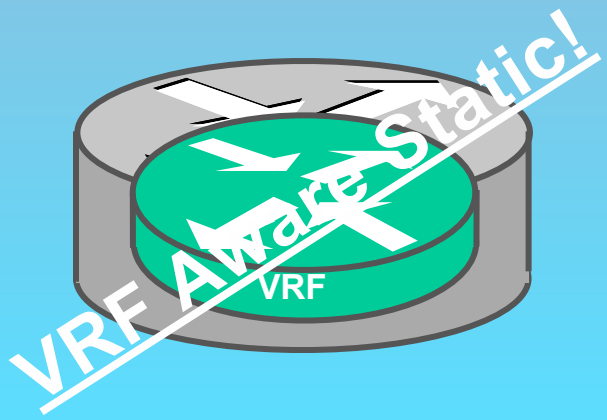


```
ip route vrf VRF-S 1.1.1.1 255.255.255.255 10.1.12.1
```

Layer 2 Serial Link? No Problem?

Back-to-Back Frame Relay Example

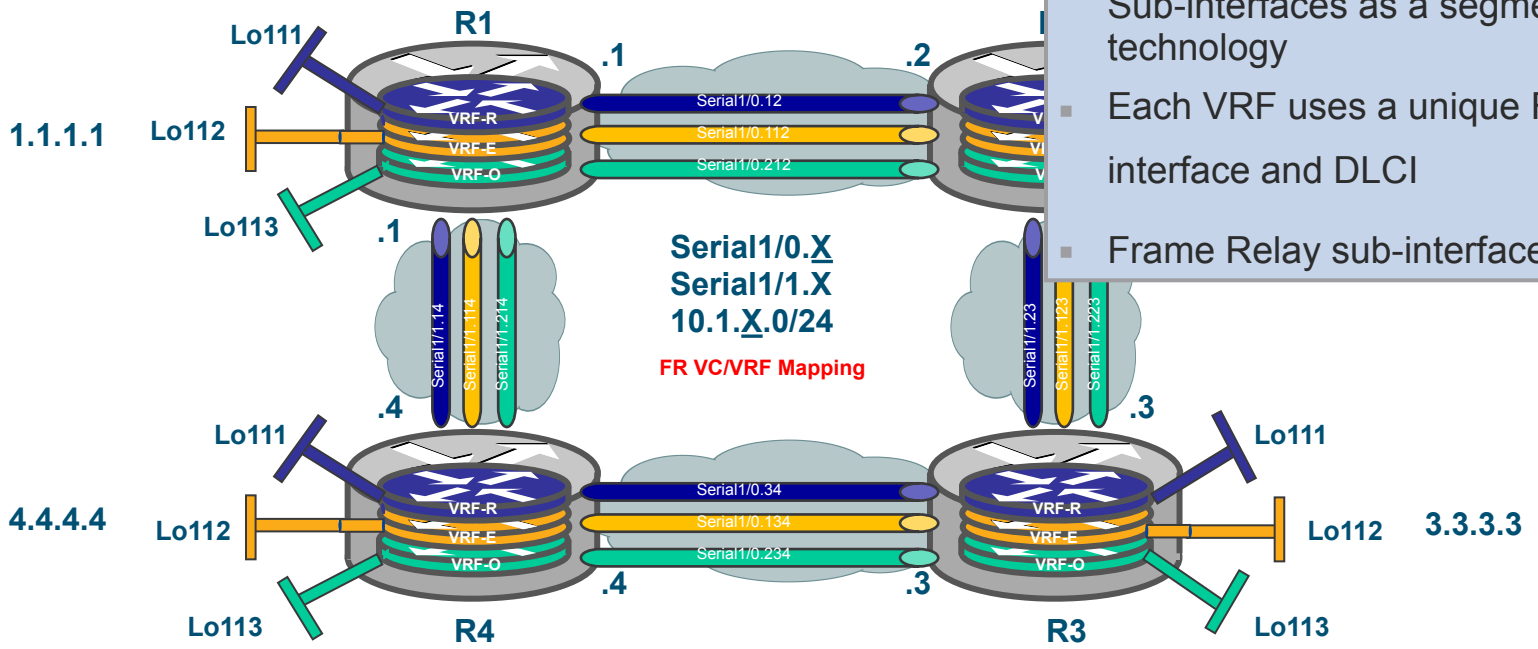

- VRF-Lite can also leverage Frame Relay Sub-interfaces as a segmentation technology
- Each VRF uses a unique Frame-Relay sub-interface and DLCI
- Frame Relay sub-interface is “VRF aware”

Configuration Note: Leveraging Back-to-Back Frame-Relay Configuration

VRF-Lite Back-to-Back Frame Relay Configuration

Command Line Interface (CLI) Review

```
ip vrf VRF-B
 rd 111:111

interface Serial1/0
 encapsulation frame-relay
 no keepalive

Interface Serial1/0.12 point-to-point
 ip vrf forwarding VRF-B
 ip address 10.1.12.1 255.255.255.0
 frame-relay interface-dlci 201
```

```
ip vrf VRF-B
 rd 222:222

interface Serial1/0
 encapsulation frame-relay
 no keepalive

Interface Serial1/0.12 point-to-point
 ip vrf forwarding VRF-B
 ip address 10.1.12.2 255.255.255.0
 frame-relay interface-dlci 201
```

Leverage what you already know!

```
router bgp 1
 address-family ipv4 vrf VRF-B
  neighbor 10.1.12.2 remote-as 2
  neighbor 10.1.12.2 activate
  no synchronization
  network 1.1.1.1 mask 255.255.255.255
 exit-address-family
```

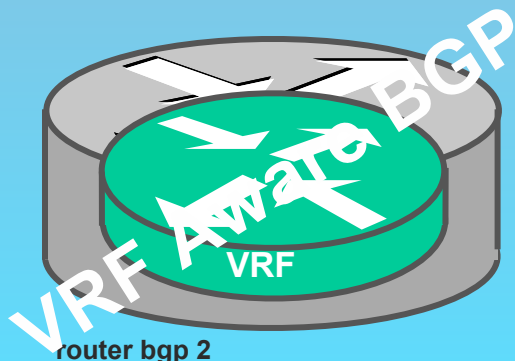


```
router bgp 2
 address-family ipv4 vrf VRF-B
  neighbor 10.1.12.1 remote-as 1
  neighbor 10.1.12.1 activate
  no synchronization
  network 2.2.2.2 mask 255.255.255.255
 exit-address-family
```

Live Exploration

VRF-Lite
Summary

- Create a VRF in router for RIB/FIB and interface segmentation
- No MPLS, LDP, or MP-BGP required
- Optimal solution when VRF count is small (~ <8)
- Supports multicast and QoS solutions
- Leverage current routing protocol knowledge and apply it to PE-CE VRF Routing

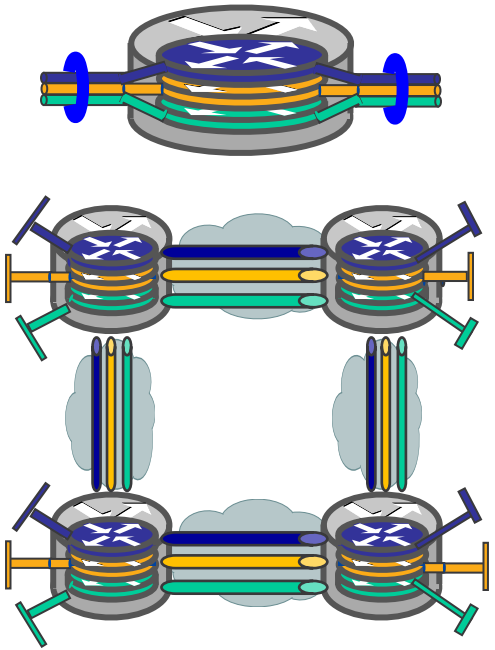

MPLS & BGP Free Core

What Is MPLS?

Most	
Painful	
Learn	
Study	

What Is

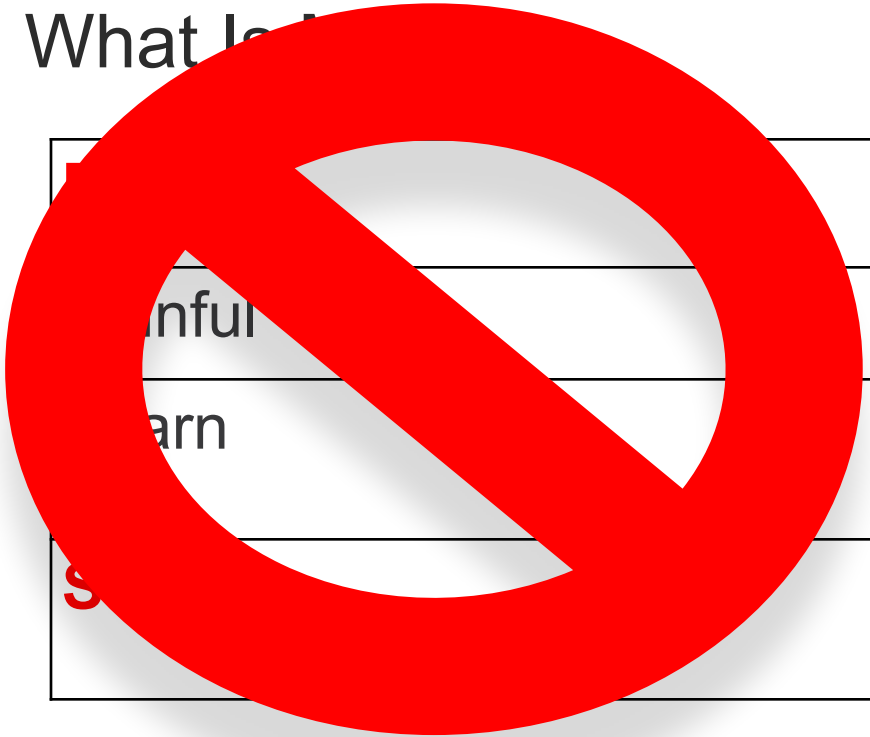

What Is MPLS?

Multi	Multi-Protocol: The ability to carry any payload Have: IPv4, IPv6, Ethernet, ATM, FR
Protocol	
Label	Uses Labels to tell a node what to do with a packet; separates forwarding (hop by hop behaviour) from routing (control plane)
Switching	Routing based on IPv4/IPv6 lookup. Everything else is label switching.

MPLS
Component Overview

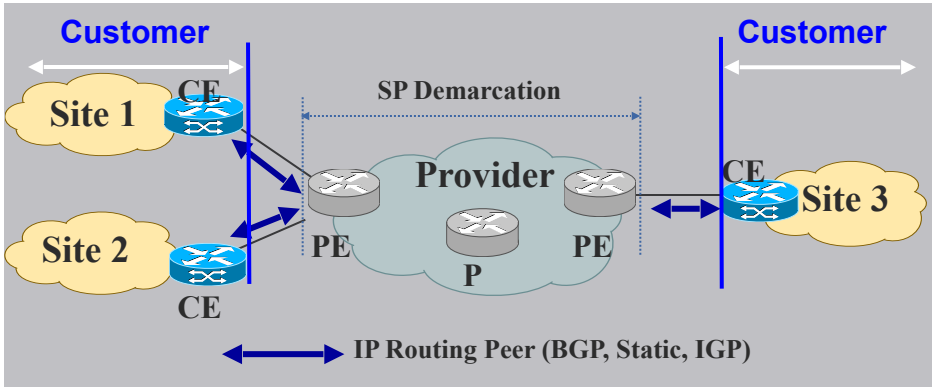

* Labels are not exchanged with the SP

MPLS
Component Overview

- CE routers owned by customer

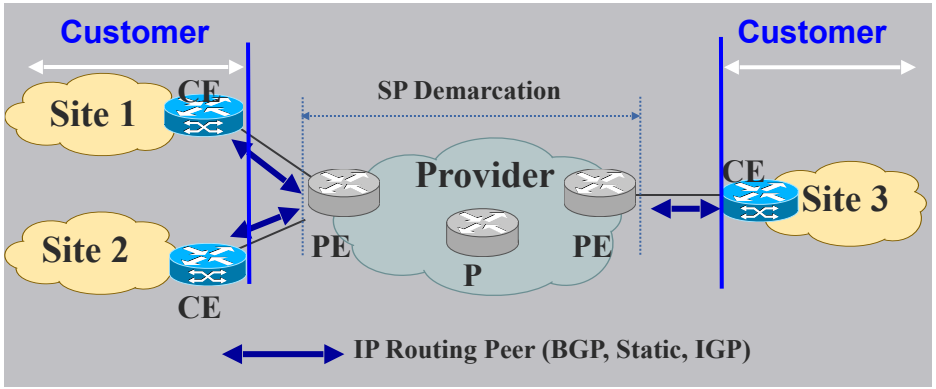

* Labels are not exchanged with the SP

MPLS
Component Overview

- CE routers owned by customer
- PE routers owned by SP

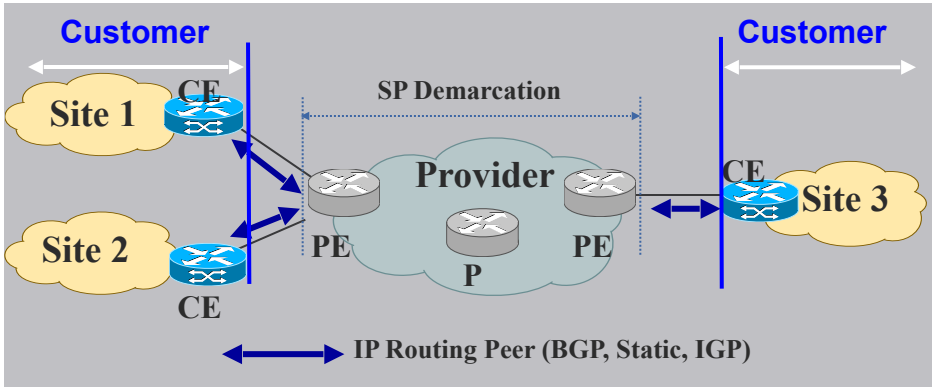

* Labels are not exchanged with the SP

MPLS
Component Overview

- CE routers owned by customer
- PE routers owned by SP
- P routers owned by SP

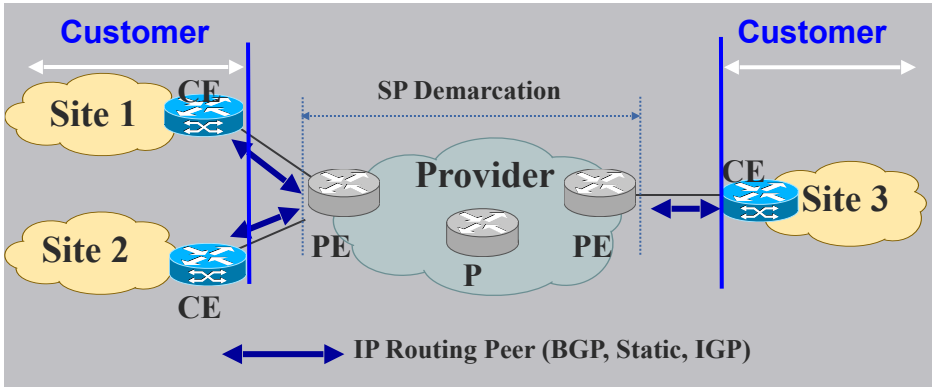

* Labels are not exchanged with the SP

MPLS
Component Overview

- CE routers owned by customer
- PE routers owned by SP
- P routers owned by SP
- Customer “peers” to “PE” via IP

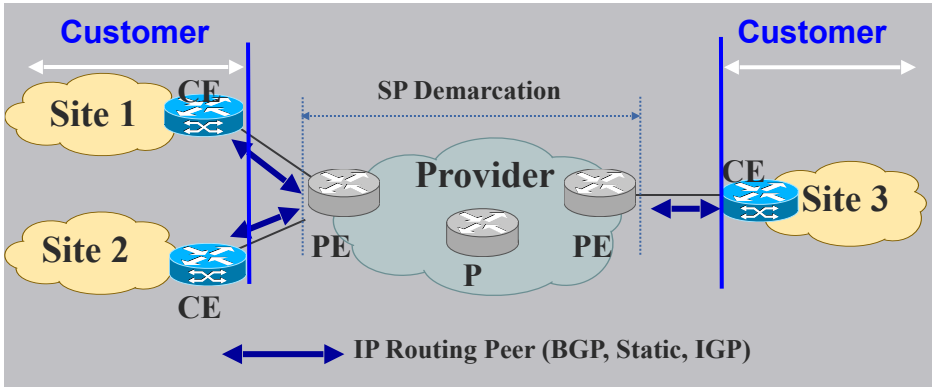

* Labels are not exchanged with the SP

MPLS
Component Overview

- CE routers owned by customer
- PE routers owned by SP
- P routers owned by SP
- Customer “peers” to “PE” via IP
- Exchanges routing with SP via routing protocol (or static route)*

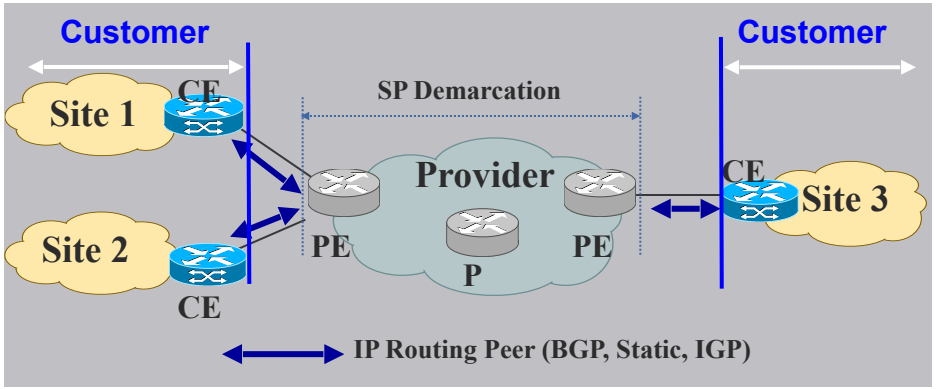

* Labels are not exchanged with the SP

MPLS
Component Overview

- CE routers owned by customer
- PE routers owned by SP
- P routers owned by SP
- Customer “peers” to “PE” via IP
- Exchanges routing with SP via routing protocol (or static route)*
- SP advertises CE routes to other CEs

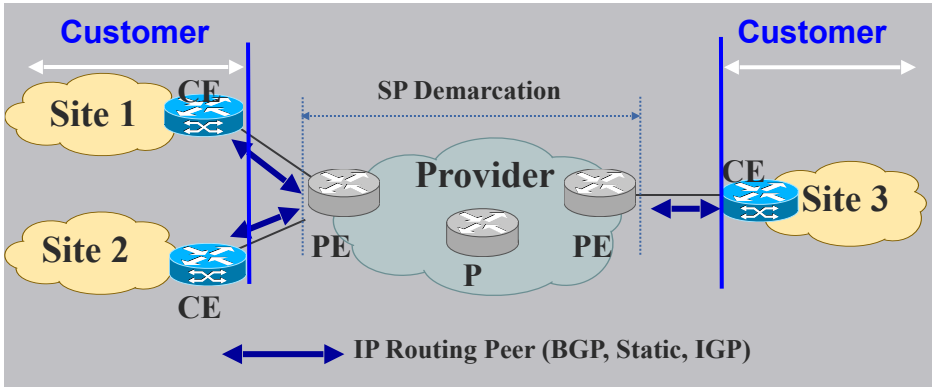

* Labels are not exchanged with the SP

IP Routing

IGP vs. BGP

- Exchange of IP routes for Loopback Reachability
  - OSPF, IS-IS, EIGRP, etc.
- iBGP neighbour peering over IGP transport
- Route towards BGP Next-Hop

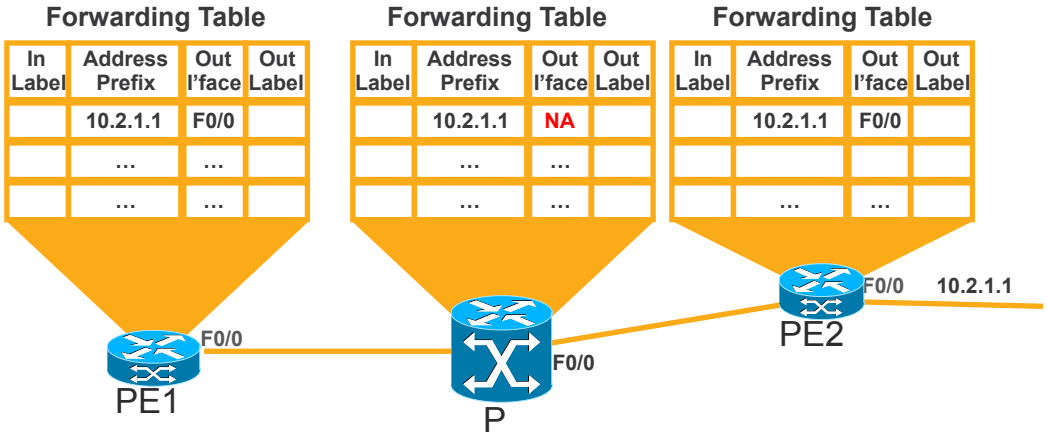

IP Routing

IGP vs. BGP

- Exchange of IP routes for Loopback Reachability
  - OSPF, IS-IS, EIGRP, etc.
- iBGP neighbour peering over IGP transport
- Route towards BGP Next-Hop

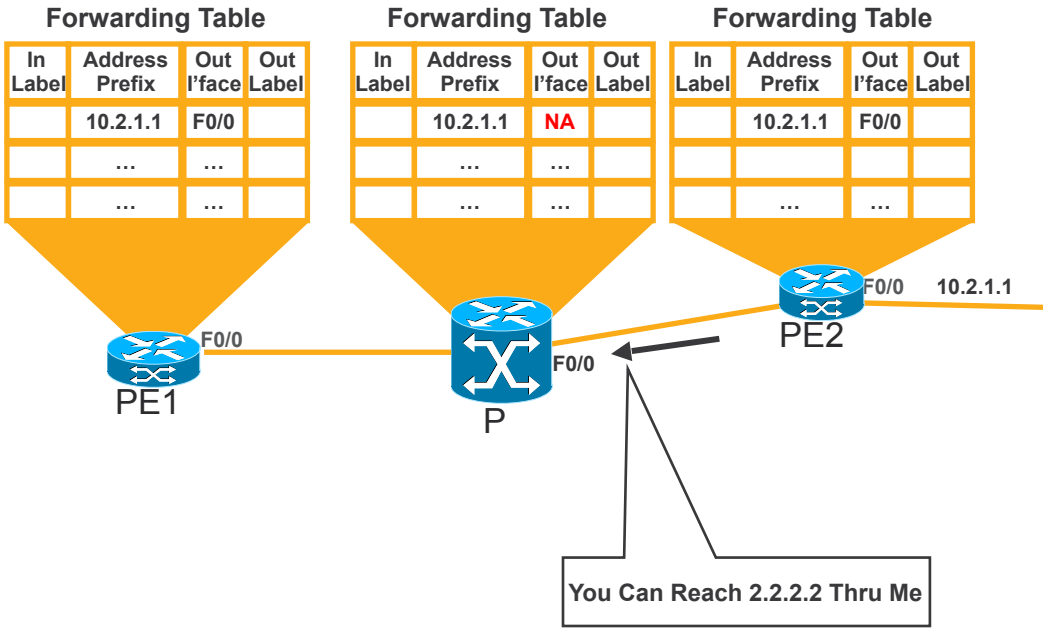

IP Routing

IGP vs. BGP

- Exchange of IP routes for Loopback Reachability
  - OSPF, IS-IS, EIGRP, etc.
- iBGP neighbour peering over IGP transport
- Route towards BGP Next-Hop

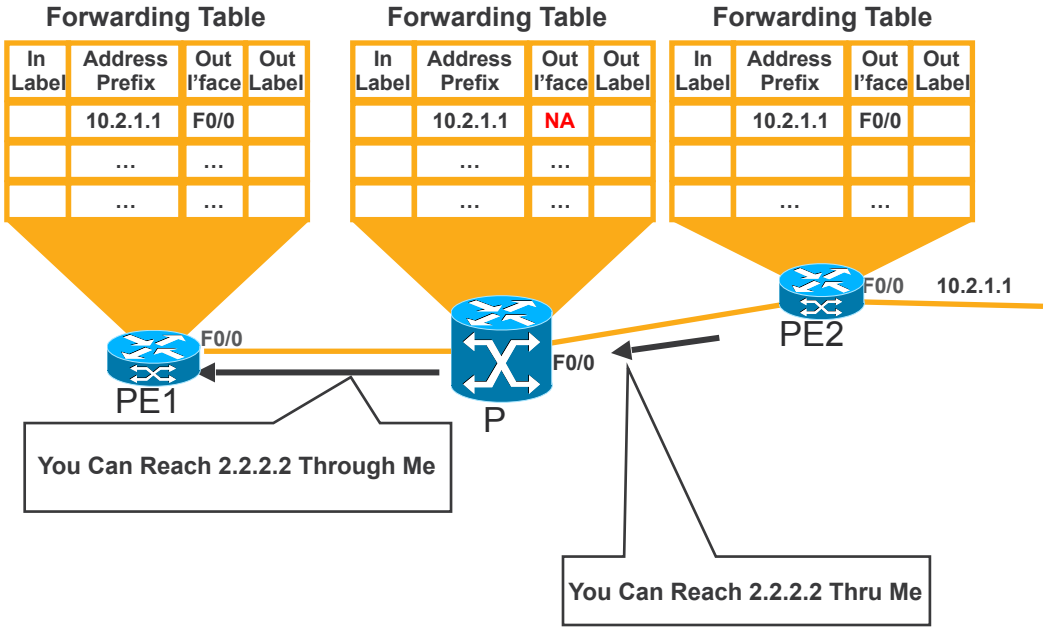

IP Routing

IGP vs. BGP

- Exchange of IP routes for Loopback Reachability
  - OSPF, IS-IS, EIGRP, etc.
- iBGP neighbour peering over IGP transport
- Route towards BGP Next-Hop

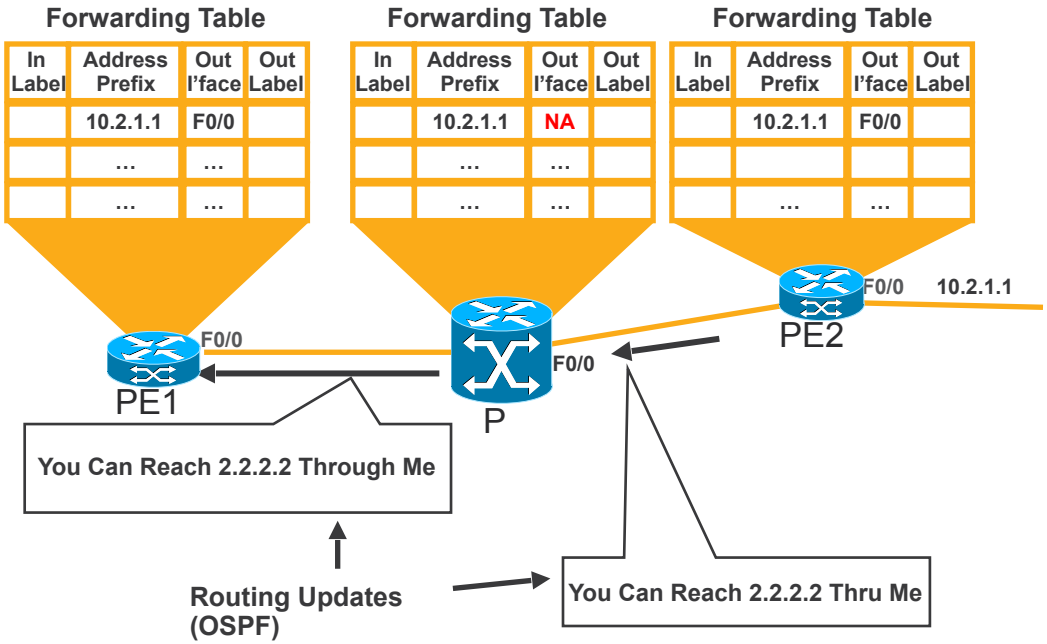

IP Routing

IGP vs. BGP

- Exchange of IP routes for Loopback Reachability
  - OSPF, IS-IS, EIGRP, etc.
- iBGP neighbour peering over IGP transport
- Route towards BGP Next-Hop

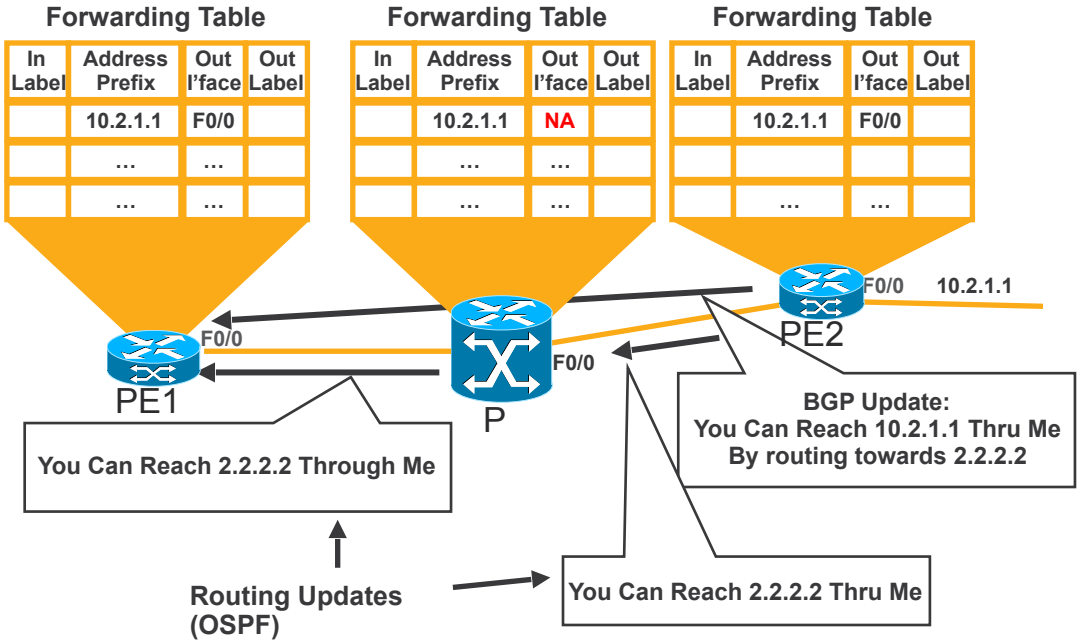

MPLS Label Switched Path (LSP) Setup with LDP

Assignment of Remote Labels

- Local label mappings are sent to connected nodes
- Receiving nodes update forwarding table
  - Out label
- LDP label advertisement happens in parallel (downstream unsolicited)

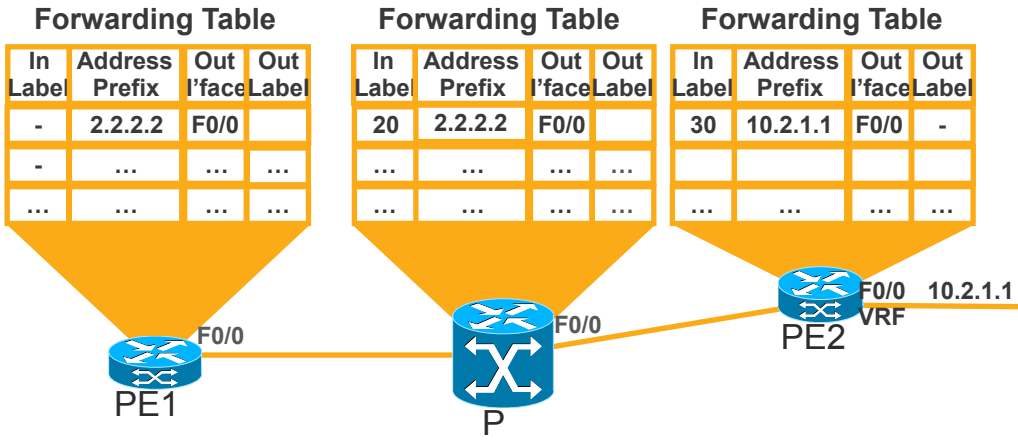

MPLS Label Switched Path (LSP) Setup with LDP

Assignment of Remote Labels

- Local label mappings are sent to connected nodes
- Receiving nodes update forwarding table
  - Out label
- LDP label advertisement happens in parallel (downstream unsolicited)

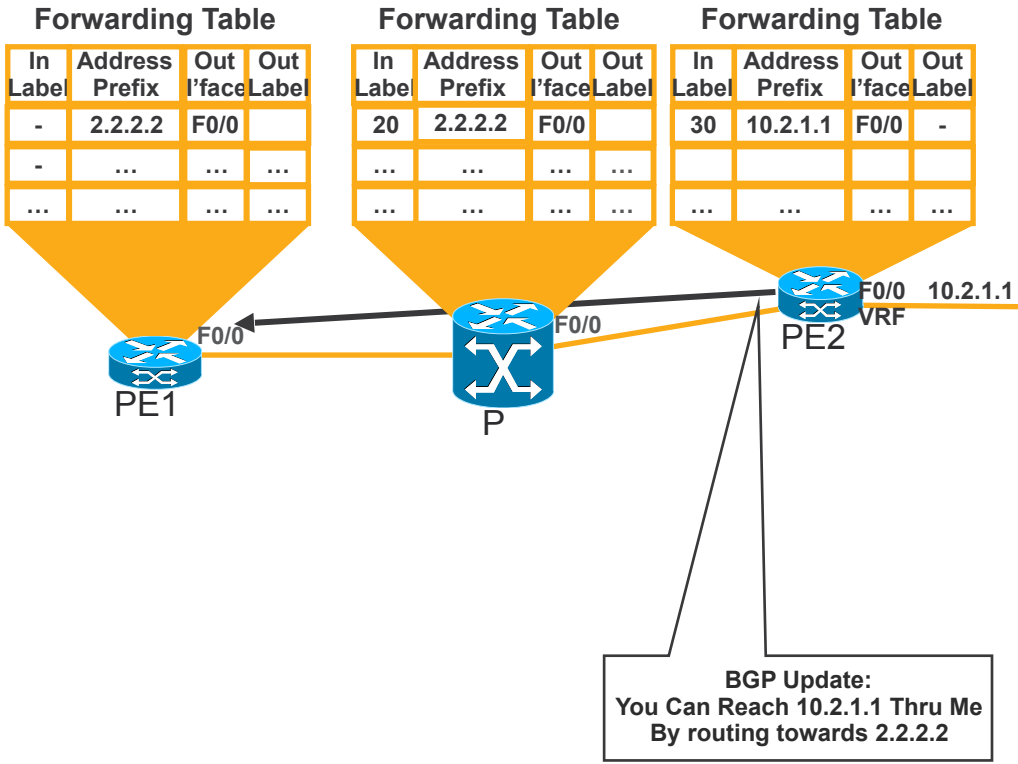

MPLS Label Switched Path (LSP) Setup with LDP

Assignment of Remote Labels

- Local label mappings are sent to connected nodes
- Receiving nodes update forwarding table
  - Out label
- LDP label advertisement happens in parallel (downstream unsolicited)

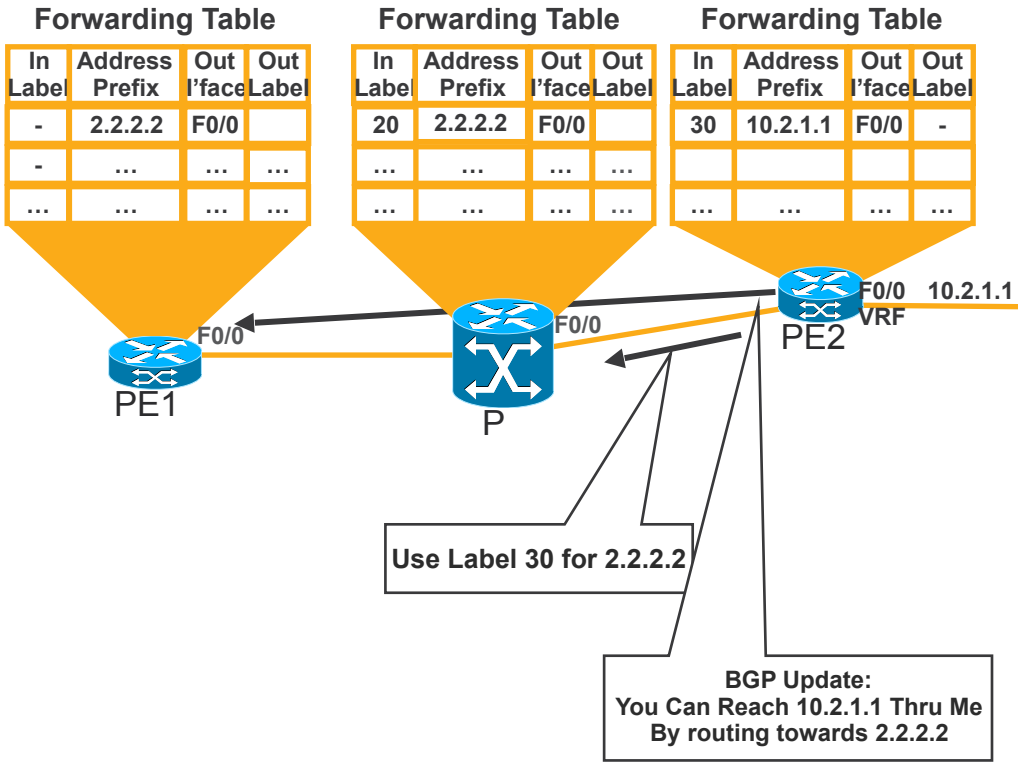

MPLS Label Switched Path (LSP) Setup with LDP

Assignment of Remote Labels

- Local label mappings are sent to connected nodes
- Receiving nodes update forwarding table
  - Out label
- LDP label advertisement happens in parallel (downstream unsolicited)

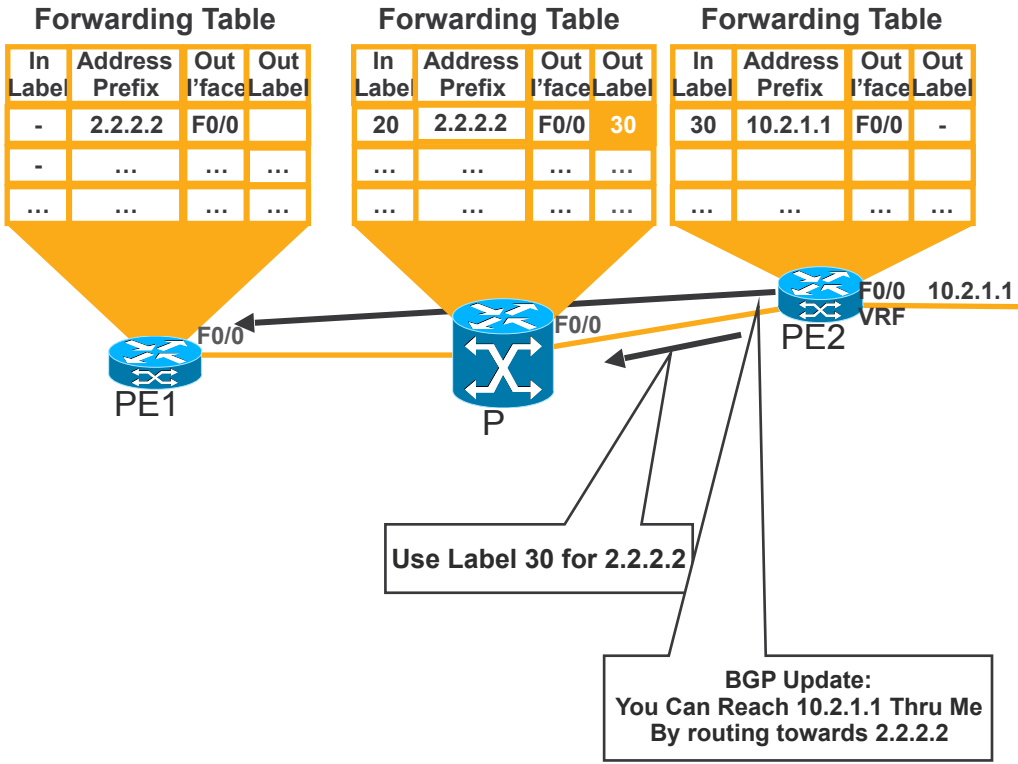

MPLS Label Switched Path (LSP) Setup with LDP

Assignment of Remote Labels

- Local label mappings are sent to connected nodes
- Receiving nodes update forwarding table
  - Out label
- LDP label advertisement happens in parallel (downstream unsolicited)

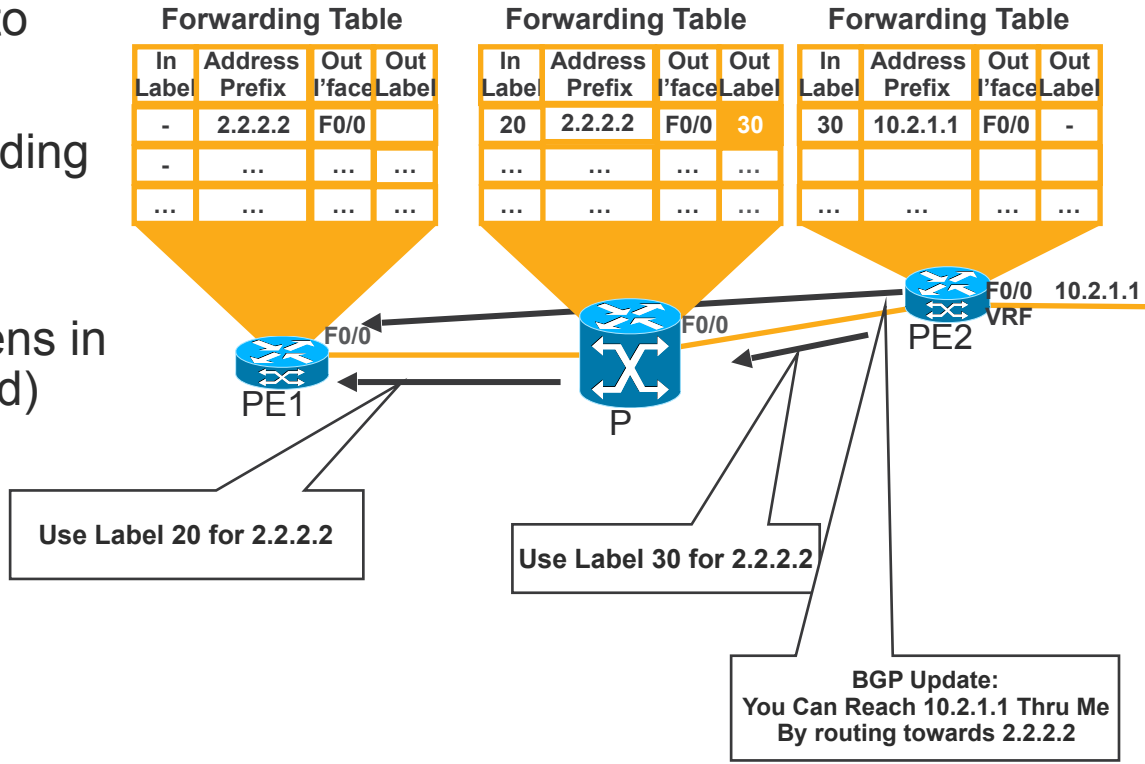

MPLS Label Switched Path (LSP) Setup with LDP

Assignment of Remote Labels

- Local label mappings are sent to connected nodes
- Receiving nodes update forwarding table
  - Out label
- LDP label advertisement happens in parallel (downstream unsolicited)

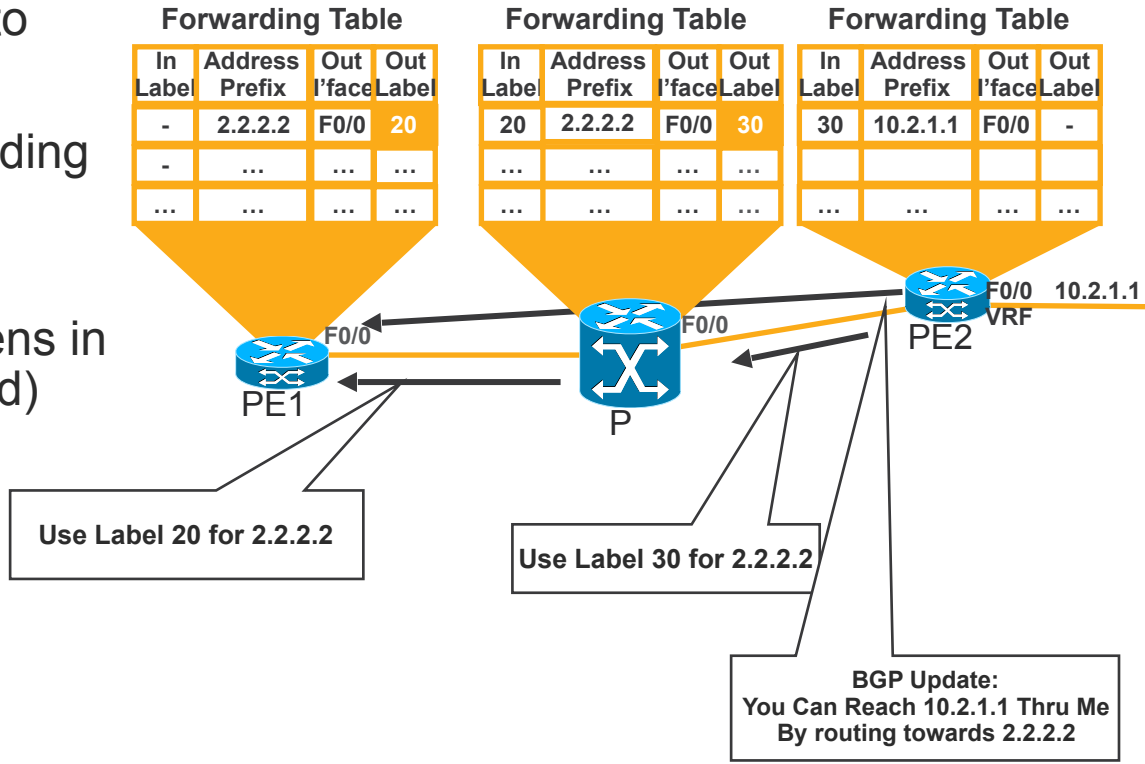

MPLS Label Switched Path (LSP) Setup with LDP

Assignment of Remote Labels

- Local label mappings are sent to connected nodes
- Receiving nodes update forwarding table
  - Out label
- LDP label advertisement happens in parallel (downstream unsolicited)

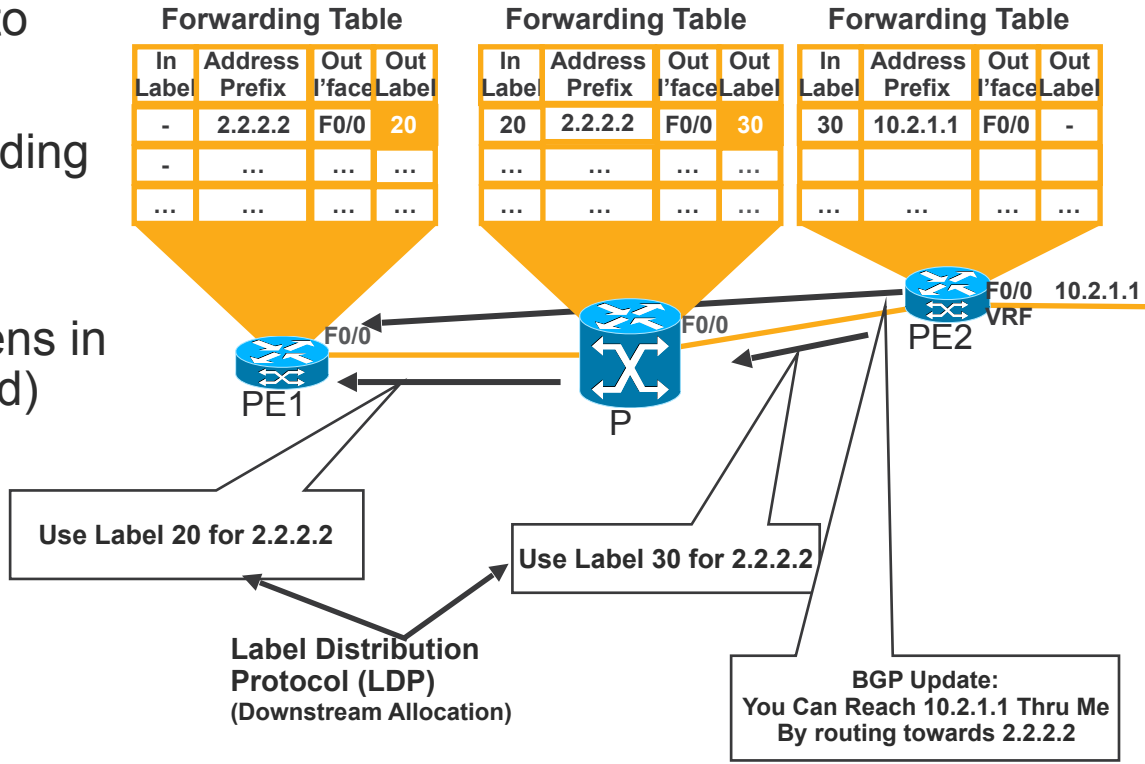

MPLS Traffic Forwarding with LDP

Hop-by-hop Traffic Forwarding Using Labels

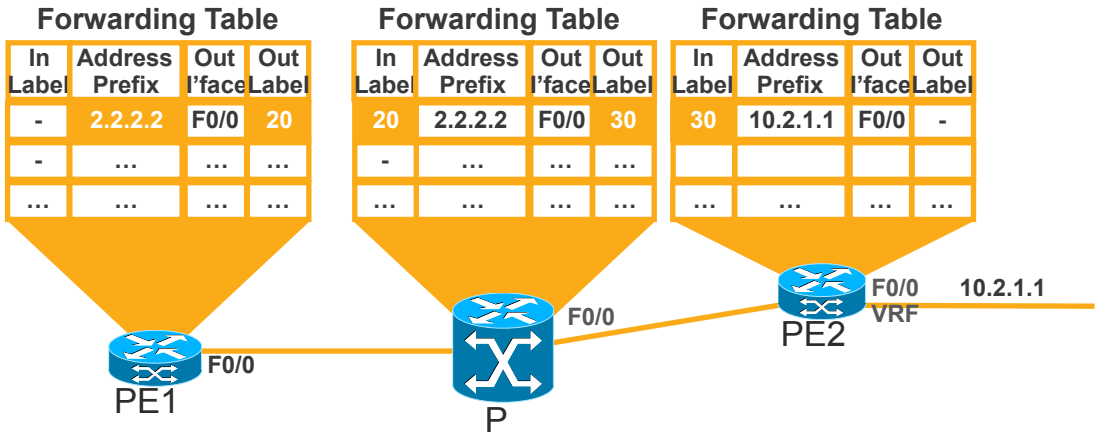

MPLS Traffic Forwarding with LDP

Hop-by-hop Traffic Forwarding Using Labels

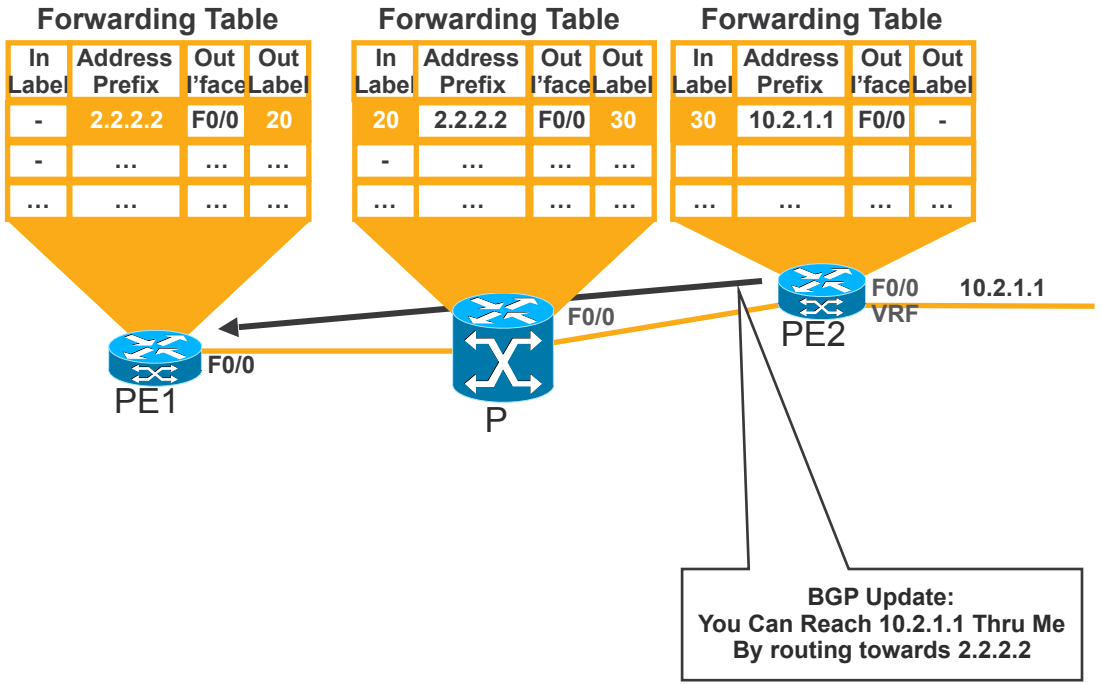

MPLS Traffic Forwarding with LDP

Hop-by-hop Traffic Forwarding Using Labels

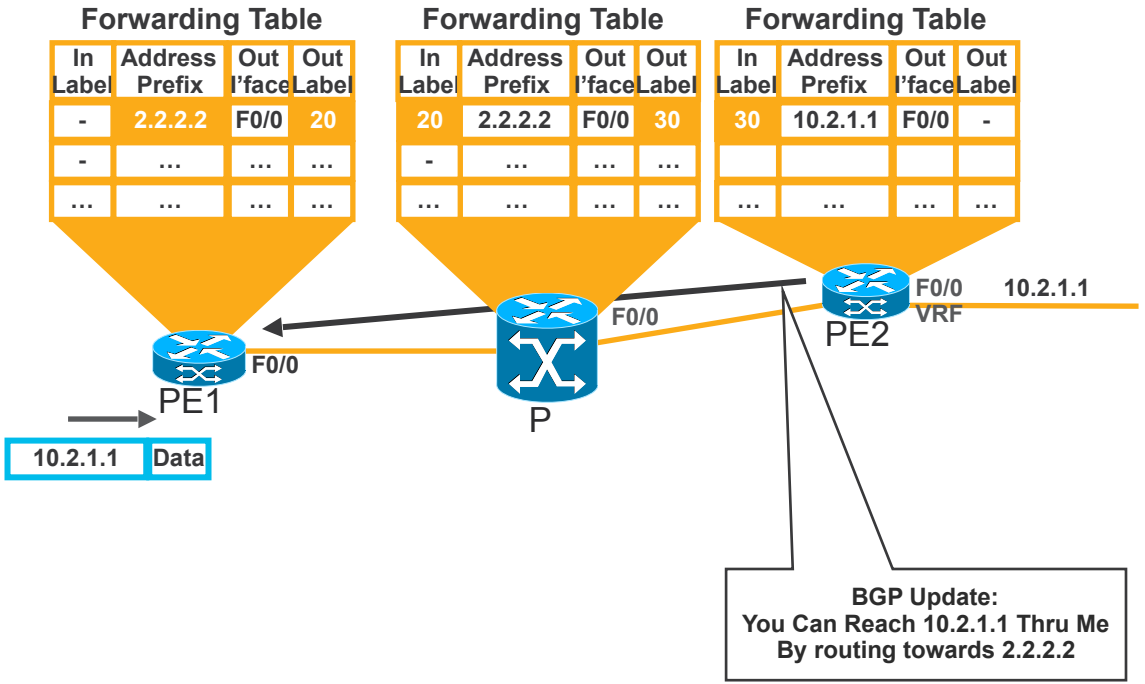

MPLS Traffic Forwarding with LDP

Hop-by-hop Traffic Forwarding Using Labels

- Ingress PE node adds labels to packet (push)
  - Via MPLS forwarding table
  - Transport label
  - VPN label (VRF)

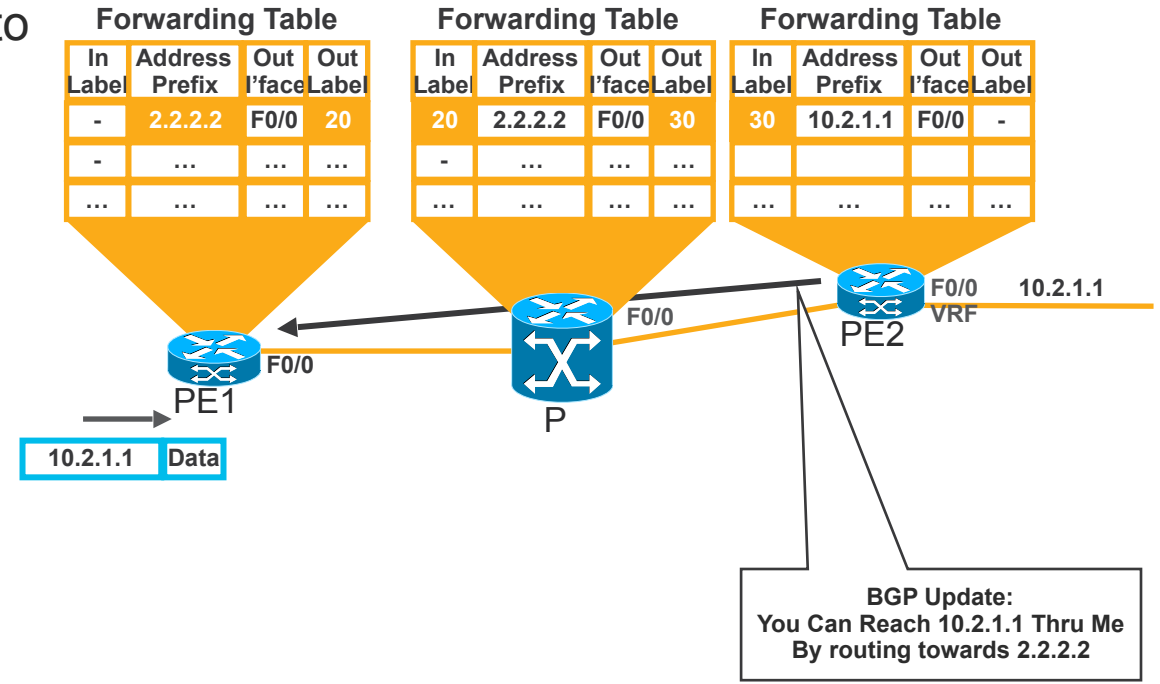

MPLS Traffic Forwarding with LDP

Hop-by-hop Traffic Forwarding Using Labels

- Ingress PE node adds labels to packet (push)
  - Via MPLS forwarding table
  - Transport label
  - VPN label (VRF)

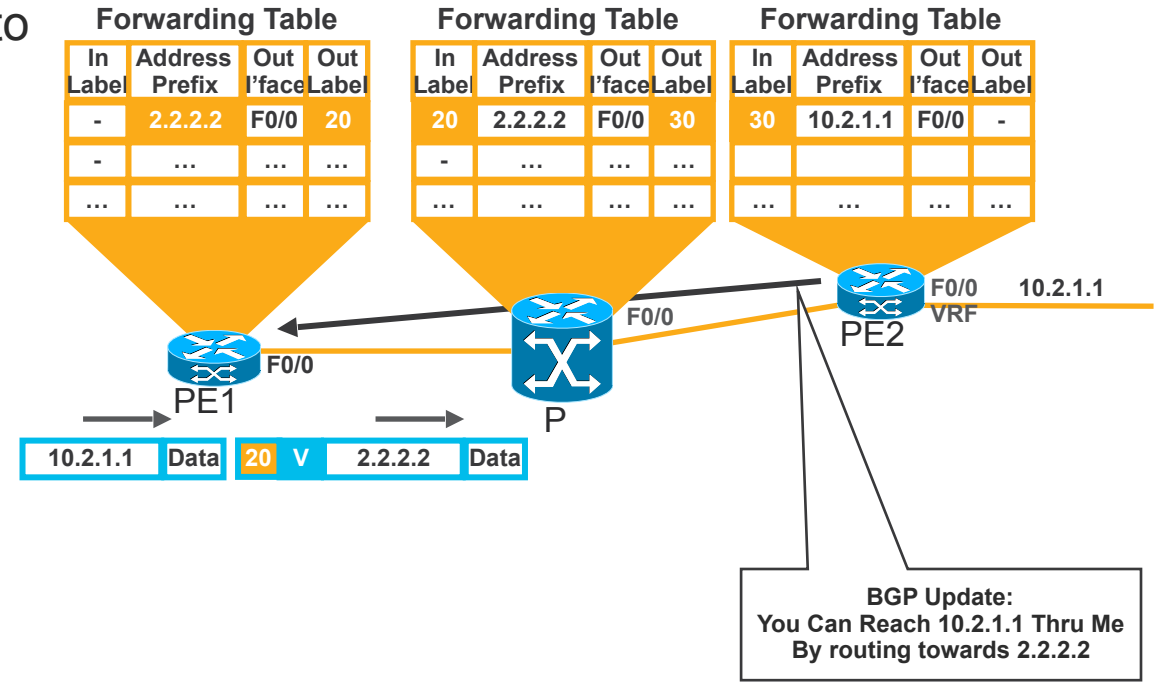

MPLS Traffic Forwarding with LDP

Hop-by-hop Traffic Forwarding Using Labels

- Ingress PE node adds labels to packet (push)
  - Via MPLS forwarding table
  - Transport label
  - VPN label (VRF)
- Downstream P node uses label forwarding decision (swap)
  - Outgoing interface
  - Out label

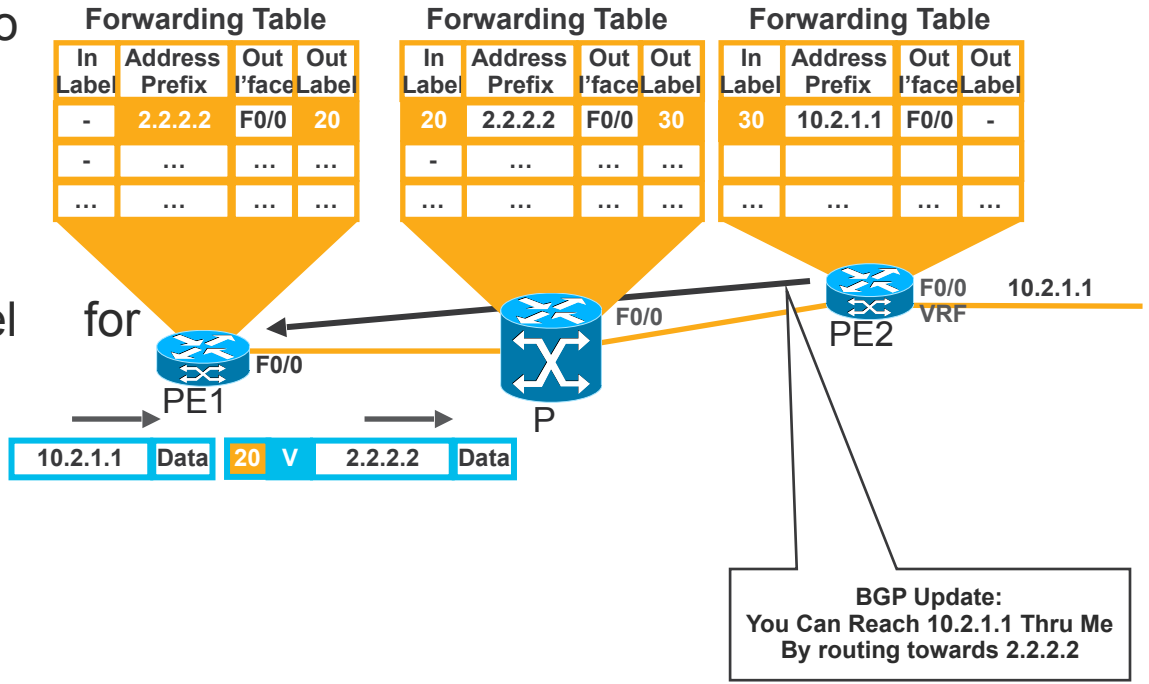

MPLS Traffic Forwarding with LDP

Hop-by-hop Traffic Forwarding Using Labels

- Ingress PE node adds labels to packet (push)
  - Via MPLS forwarding table
  - Transport label
  - VPN label (VRF)
- Downstream P node uses label forwarding decision (swap)
  - Outgoing interface
  - Out label

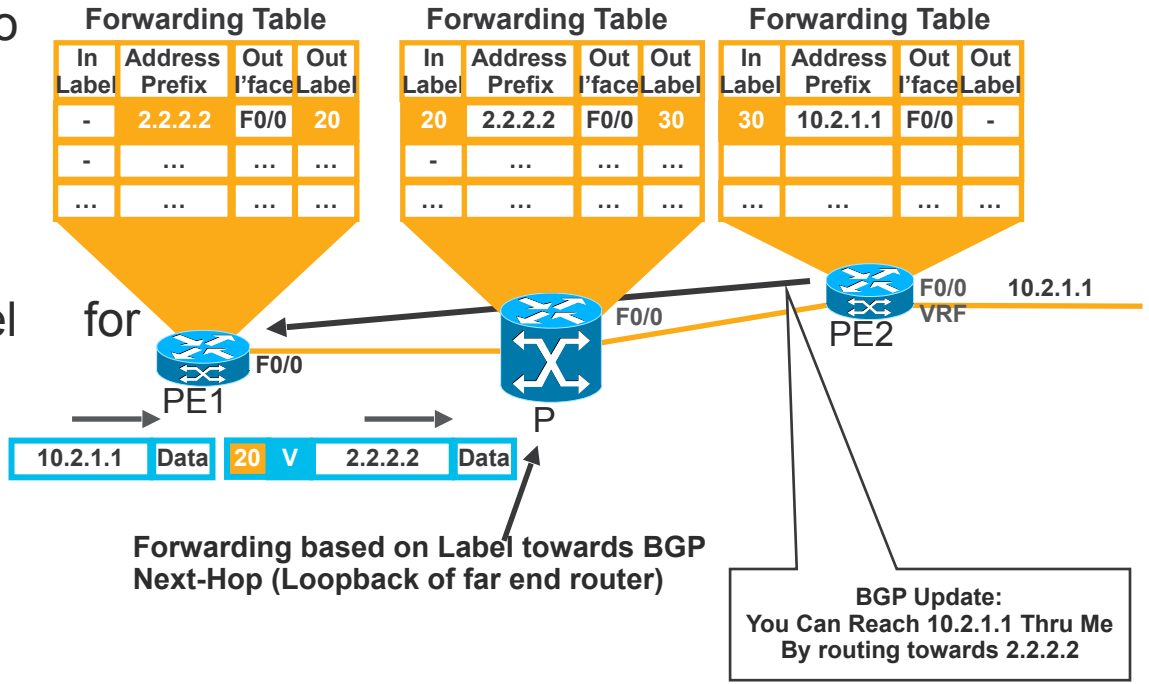

MPLS Traffic Forwarding with LDP

Hop-by-hop Traffic Forwarding Using Labels

- Ingress PE node adds labels to packet (push)
  - Via MPLS forwarding table
  - Transport label
  - VPN label (VRF)
- Downstream P node uses label forwarding decision (swap)
  - Outgoing interface
  - Out label

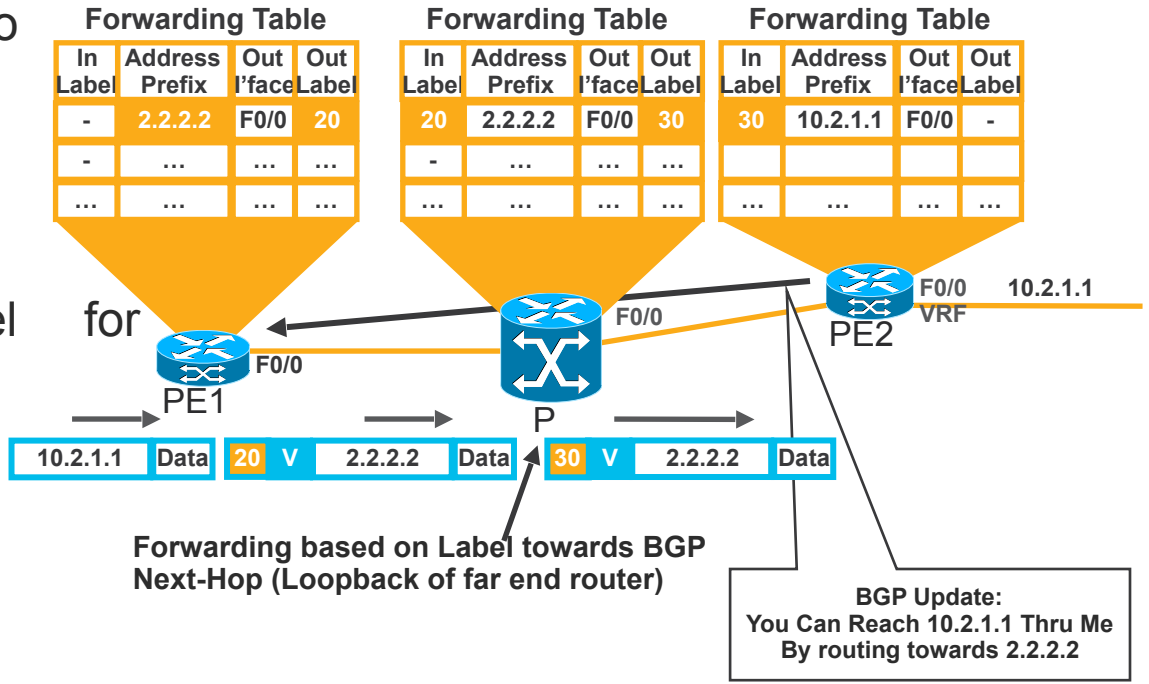

MPLS Traffic Forwarding with LDP

Hop-by-hop Traffic Forwarding Using Labels

- Ingress PE node adds labels to packet (push)
  - Via MPLS forwarding table
  - Transport label
  - VPN label (VRF)
- Downstream P node uses label forwarding decision (swap)
  - Outgoing interface
  - Out label
- Egress PE removes label and forwards original packet (pop)

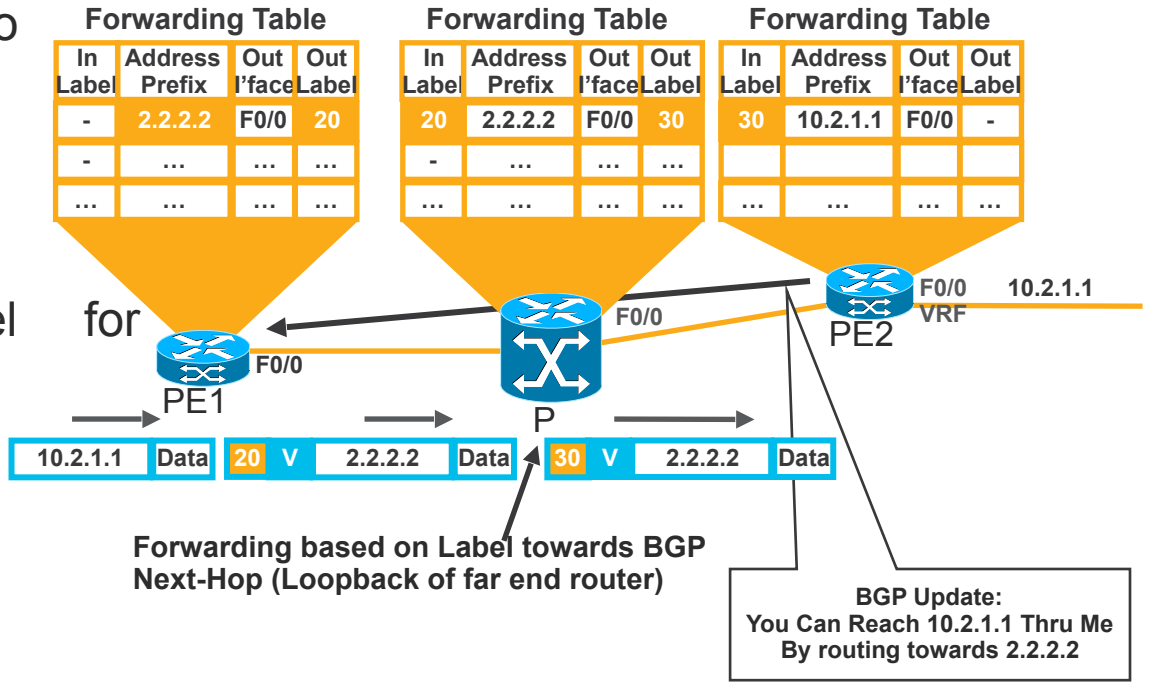

MPLS Traffic Forwarding with LDP

Hop-by-hop Traffic Forwarding Using Labels

- Ingress PE node adds labels to packet (push)
  - Via MPLS forwarding table
  - Transport label
  - VPN label (VRF)
- Downstream P node uses label forwarding decision (swap)
  - Outgoing interface
  - Out label
- Egress PE removes label and forwards original packet (pop)

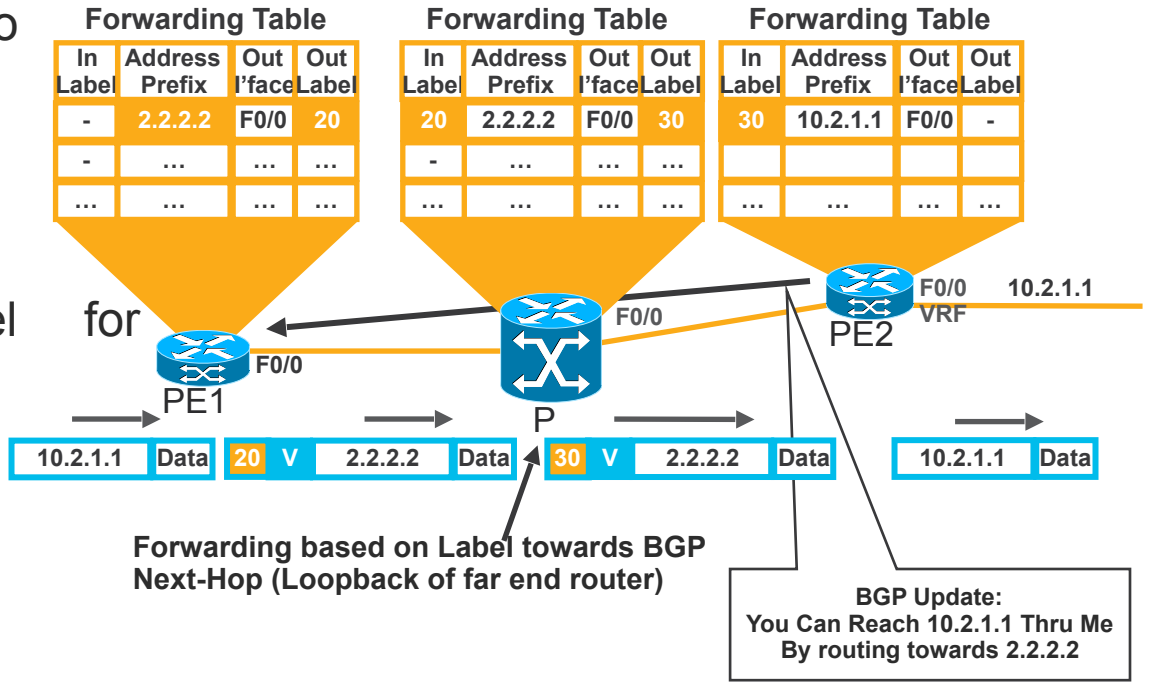

BGP Free Core
Component Overview

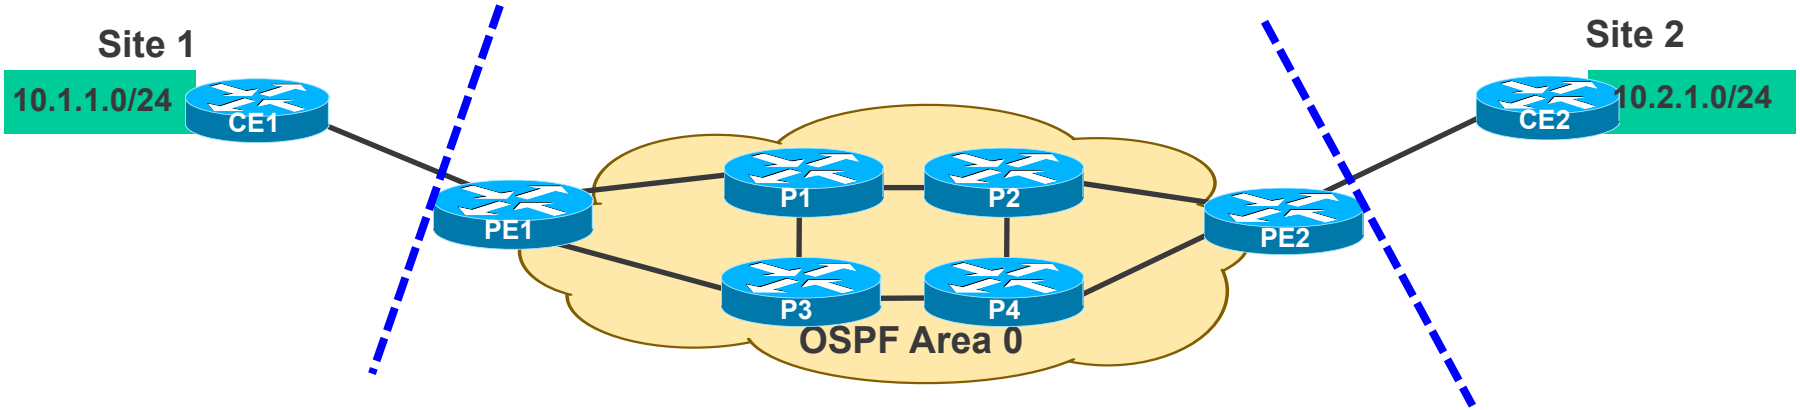

End-to-End BGP and redistribution of routes into OSPF core not necessary!

BGP Free Core
Component Overview

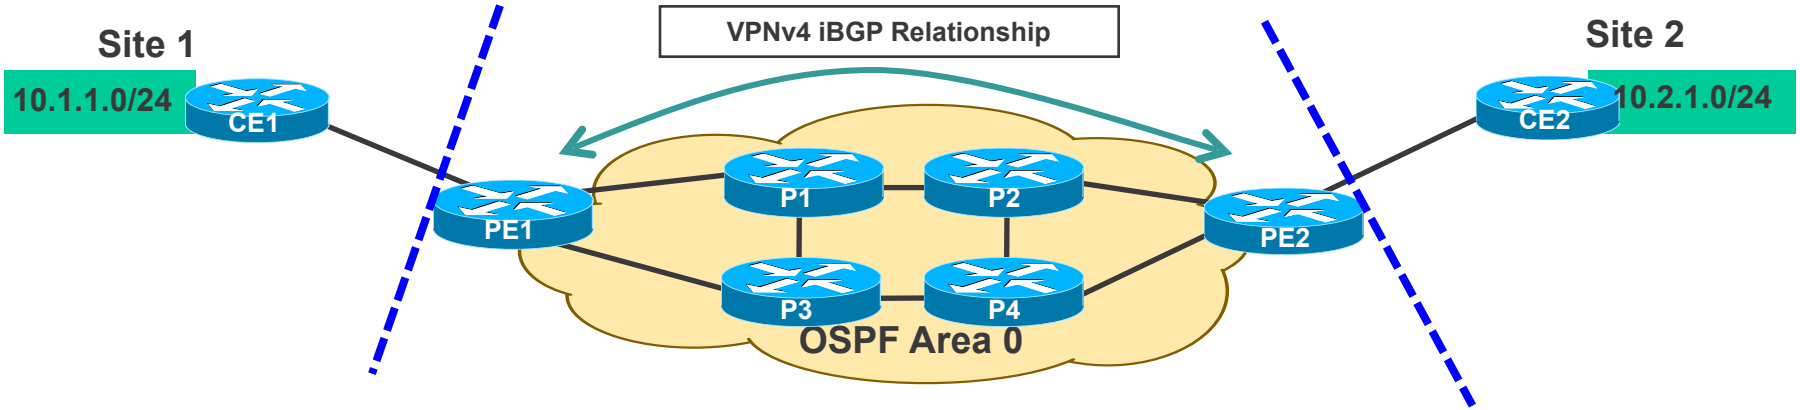

1. Always route towards BGP Next-Hop

End-to-End BGP and redistribution of routes into OSPF core not necessary!

BGP Free Core
Component Overview

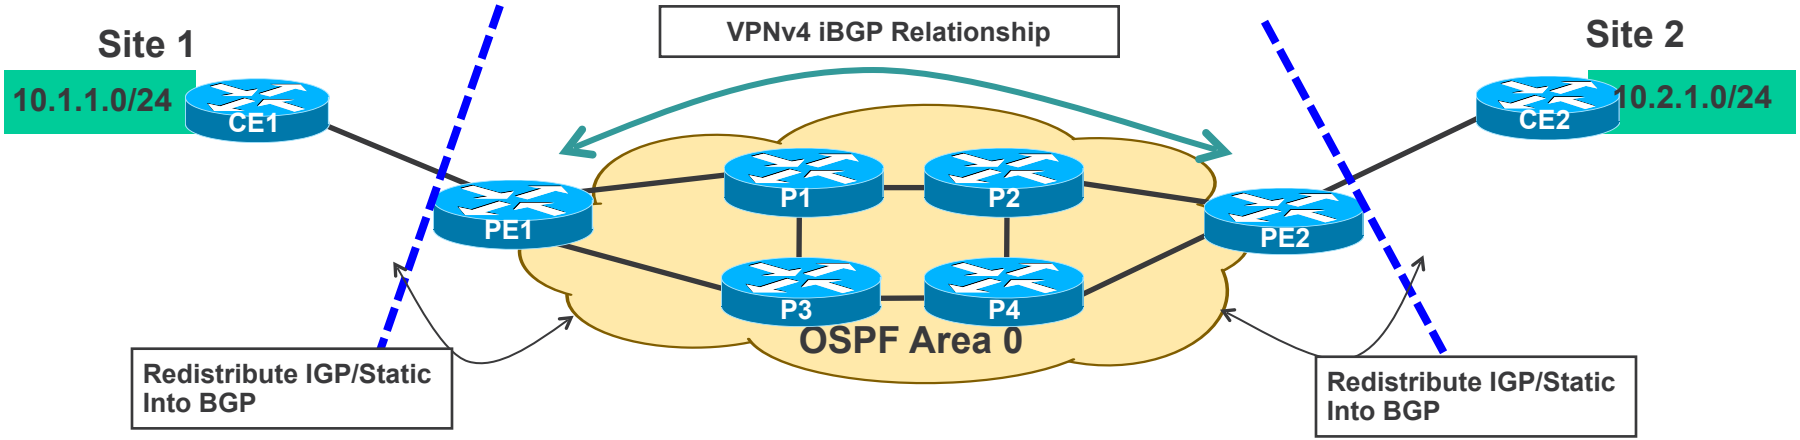

1. Always route towards BGP Next-Hop
2. Routes will be valid on PE Routers

End-to-End BGP and redistribution of routes into OSPF core not necessary!

BGP Free Core
Component Overview

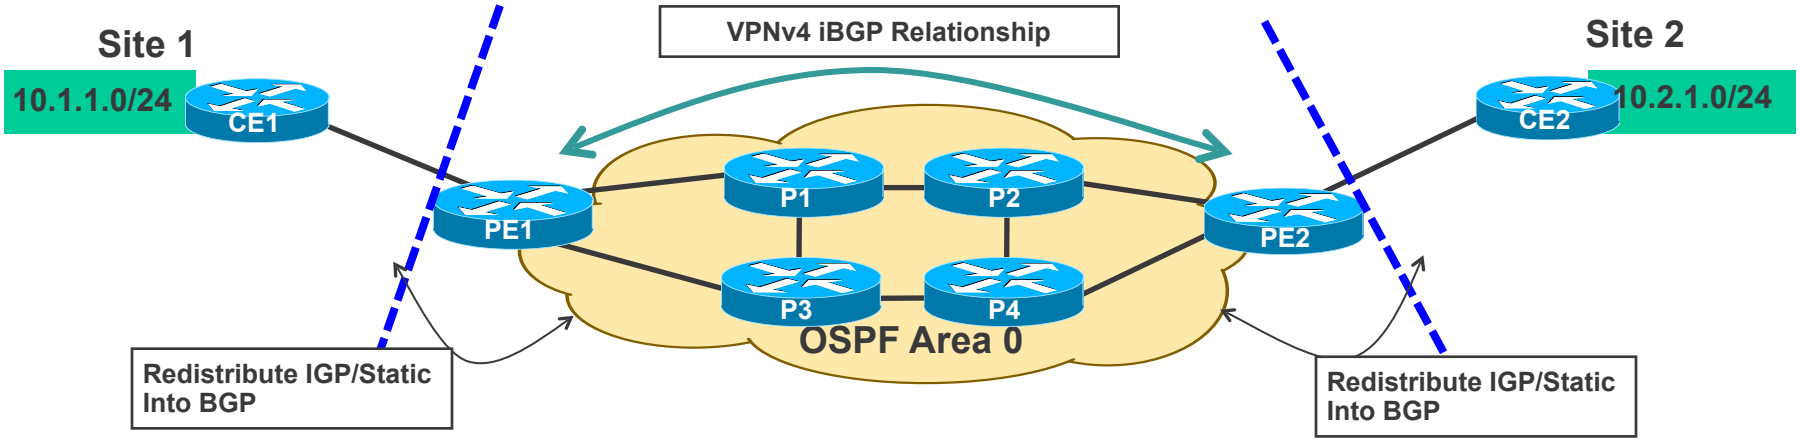

1. Always route towards BGP Next-Hop
2. Routes will be valid on PE Routers
3. Will label switch towards BGP Next-Hop of PE with MPLS enabled

End-to-End BGP and redistribution of routes into OSPF core not necessary!

Multiprotocol BGP (MP-BGP)

Multiprotocol BGP (MP-BGP)

Bringing It All Together

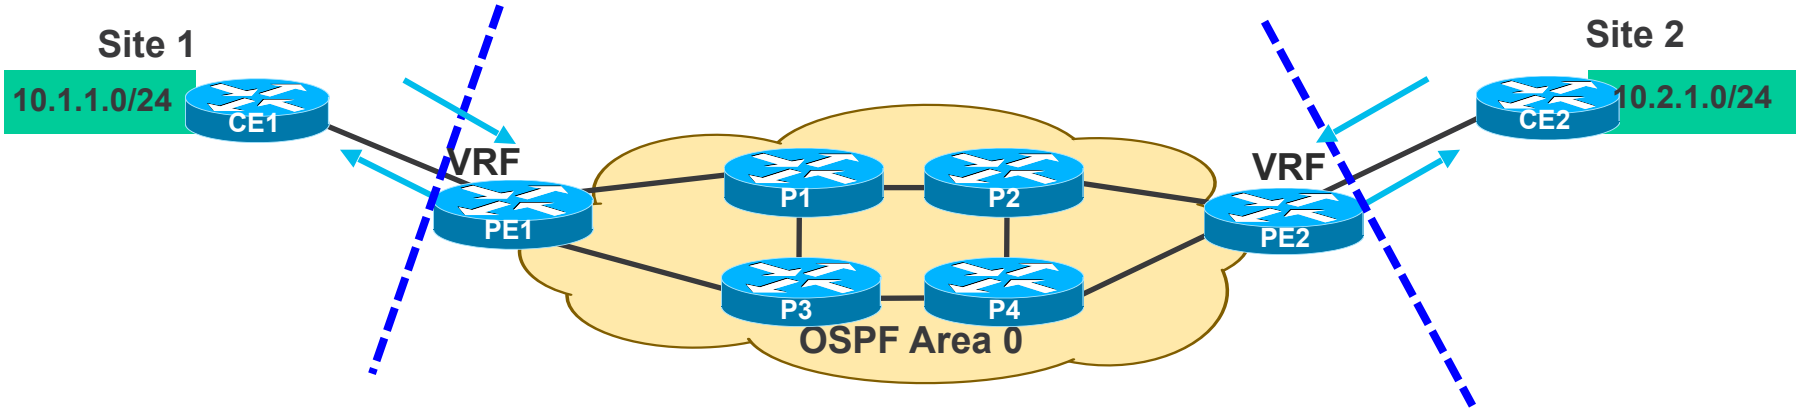

Multiprotocol BGP (MP-BGP)

Bringing It All Together

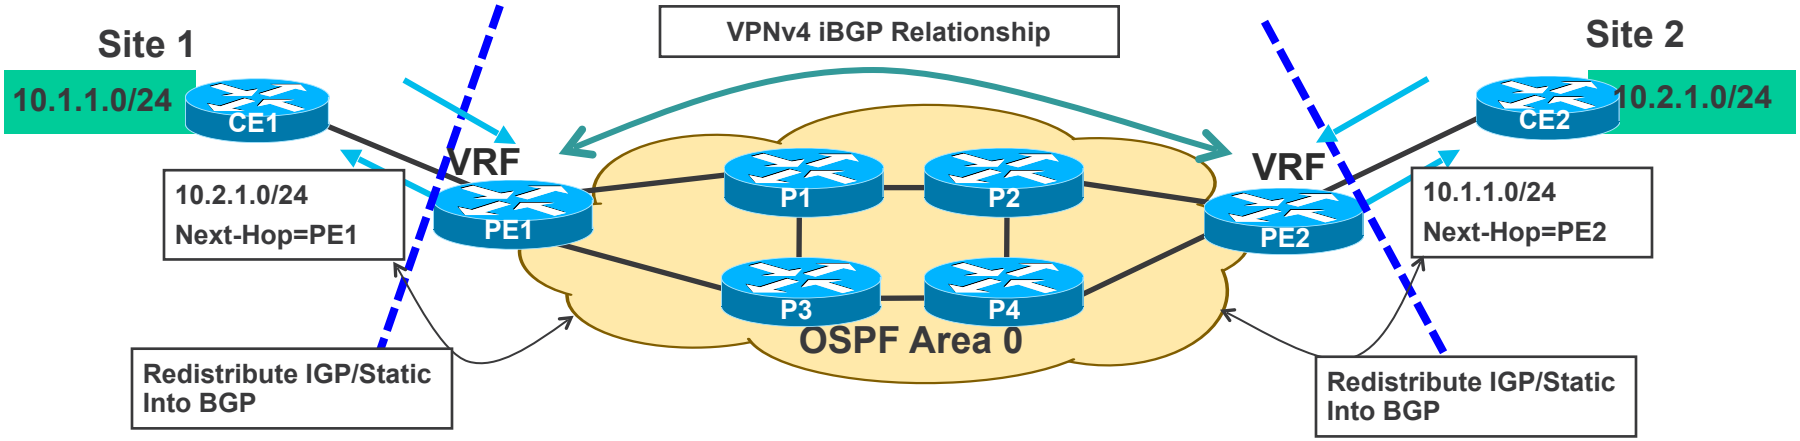

Multiprotocol BGP (MP-BGP)

Bringing It All Together

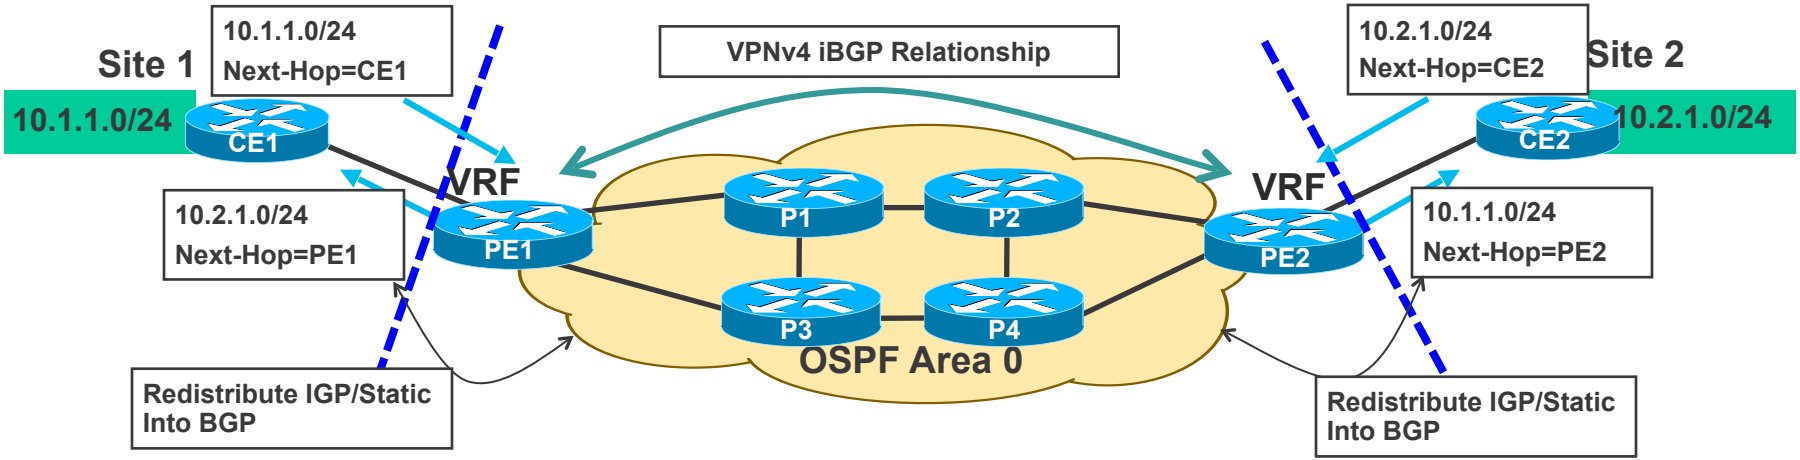

1. PE receives an IPv4 update on a VRF interface (eBGP/OSPF/RIP/EIGRP)

Multiprotocol BGP (MP-BGP)

Bringing It All Together

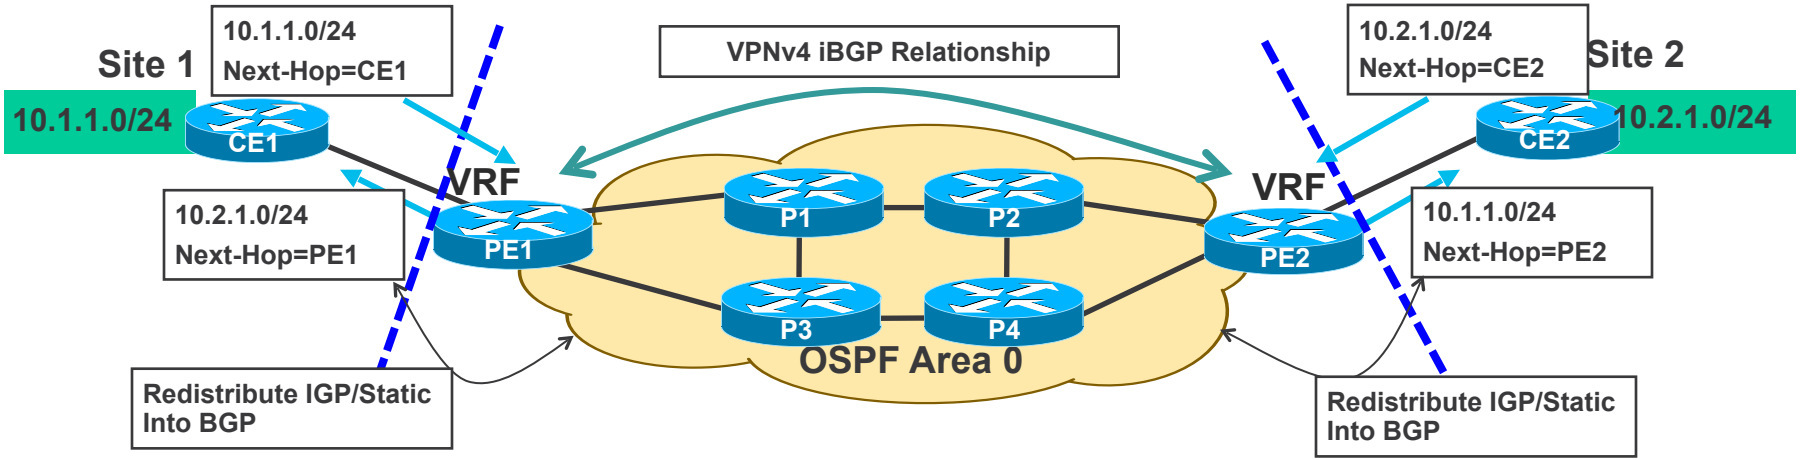

1. PE receives an IPv4 update on a VRF interface (eBGP/OSPF/RIP/EIGRP)
2. PE translates it into VPNv4 address (96-bit address) (64-bit RD + 32 bit IPv4 address)

Multiprotocol BGP (MP-BGP)

Bringing It All Together

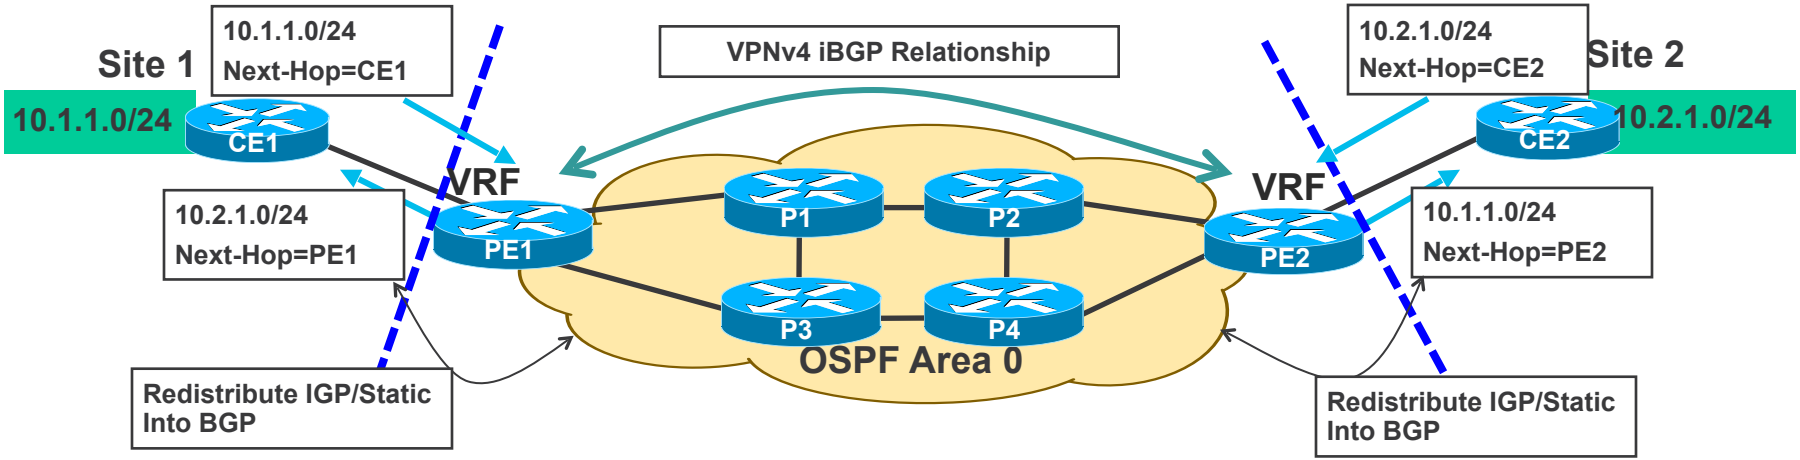

1. PE receives an IPv4 update on a VRF interface (eBGP/OSPF/RIP/EIGRP)
2. PE translates it into VPNv4 address (96-bit address) (64-bit RD + 32 bit IPv4 address)
   - Assigns an RT per VRF configuration

Multiprotocol BGP (MP-BGP)

Bringing It All Together

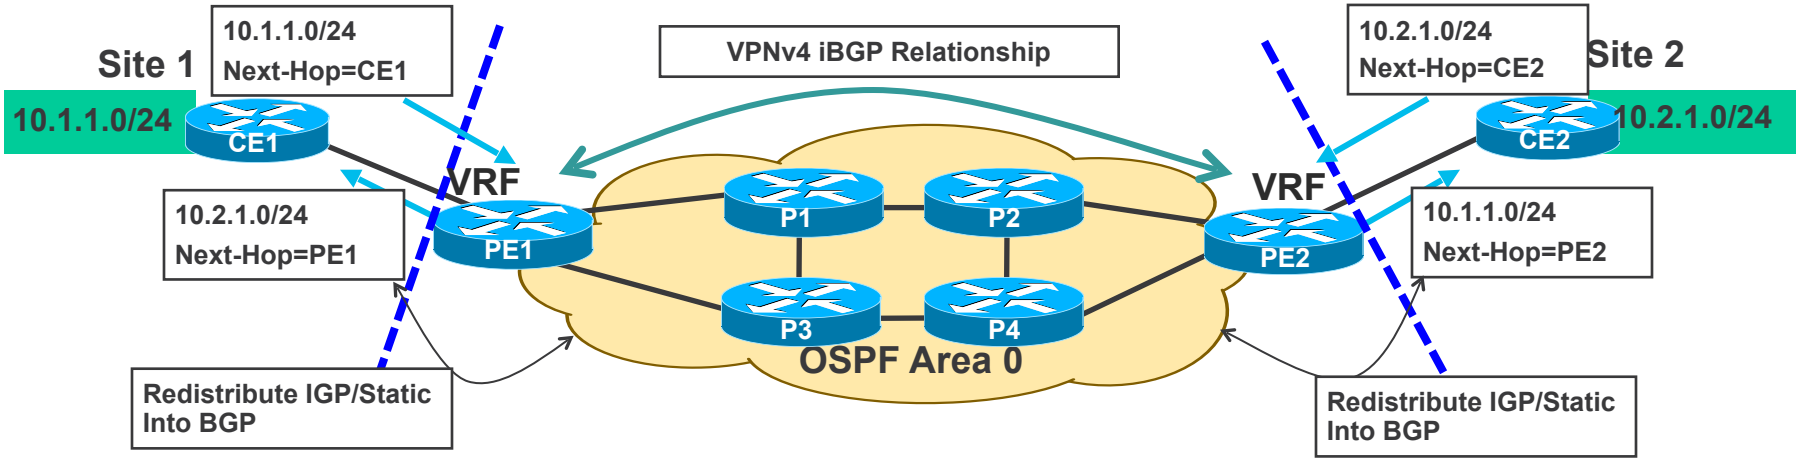

1. PE receives an IPv4 update on a VRF interface (eBGP/OSPF/RIP/EIGRP)
2. PE translates it into VPNv4 address (96-bit address) (64-bit RD + 32 bit IPv4 address)
   - Assigns an RT per VRF configuration
   - Rewrites next-hop attribute to itself

Multiprotocol BGP (MP-BGP)

Bringing It All Together

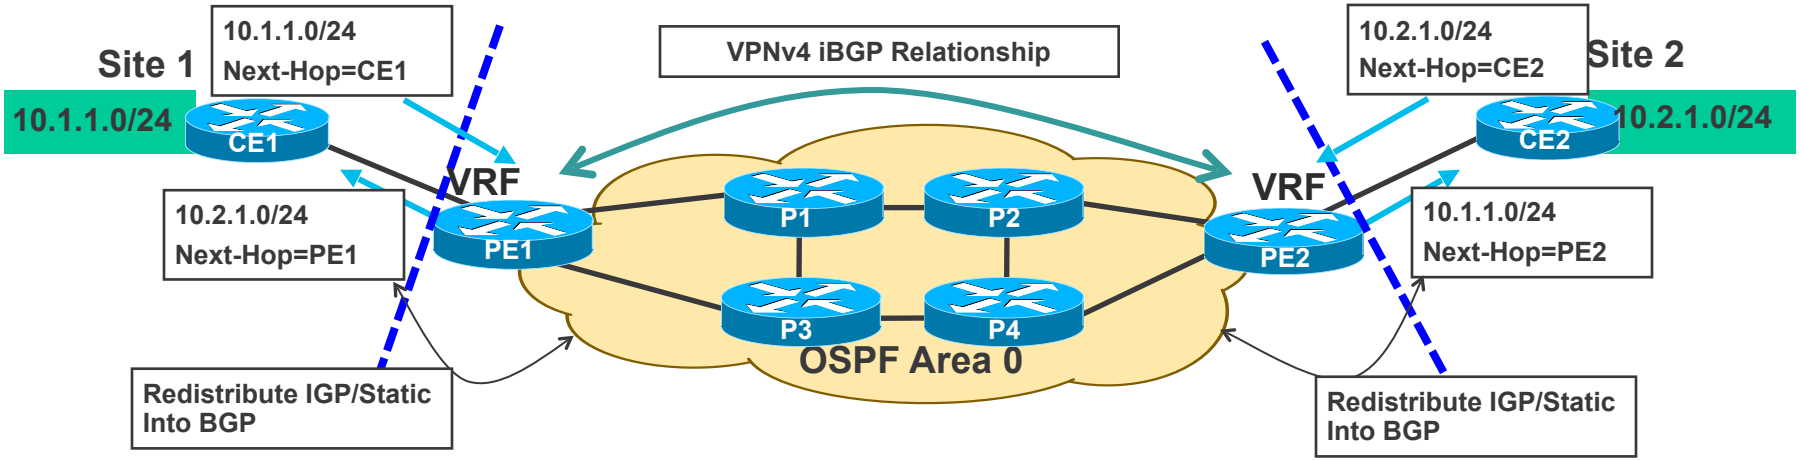

1. PE receives an IPv4 update on a VRF interface (eBGP/OSPF/RIP/EIGRP)
2. PE translates it into VPNv4 address (96-bit address) (64-bit RD + 32 bit IPv4 address)
   - Assigns an RT per VRF configuration
   - Rewrites next-hop attribute to itself
   - Assigns a label based on VRF and/or interface

Multiprotocol BGP (MP-BGP)

Bringing It All Together

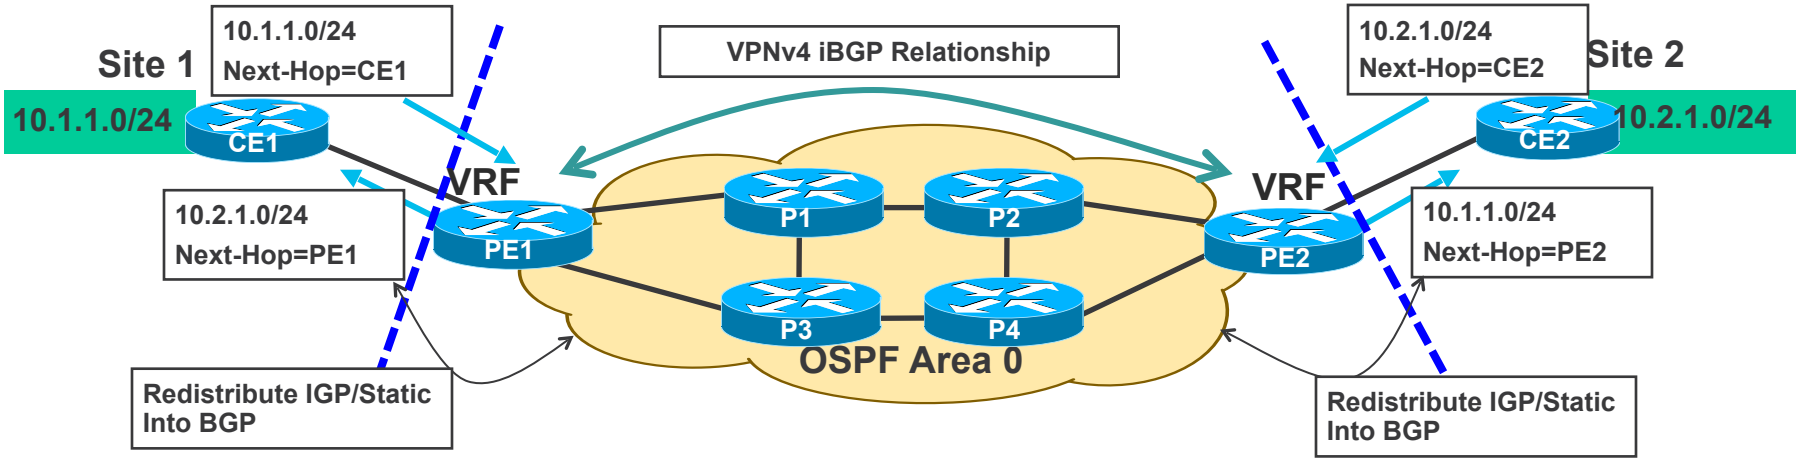

1. PE receives an IPv4 update on a VRF interface (eBGP/OSPF/RIP/EIGRP)
2. PE translates it into VPNv4 address (96-bit address) (64-bit RD + 32 bit IPv4 address)
   - Assigns an RT per VRF configuration
   - Rewrites next-hop attribute to itself
   - Assigns a label based on VRF and/or interface
3. PE sends MP-iBGP update to other PE routers

Why an RD and VPNv4 Address?

Use Case

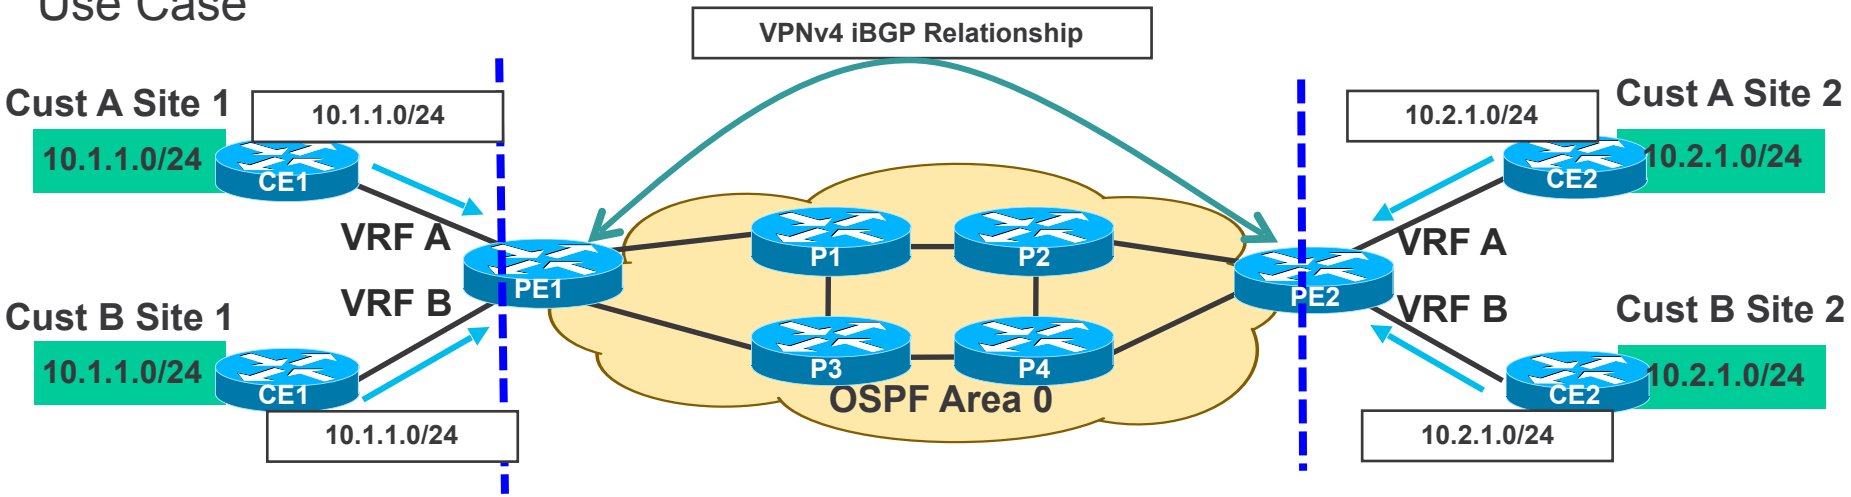

VPNv4 prefixes are the combination of a 64-bit RD and a 32-bit IPv4 prefix. VPNv4 prefixes are 96-bits in length

Why an RD and VPNv4 Address?

Use Case

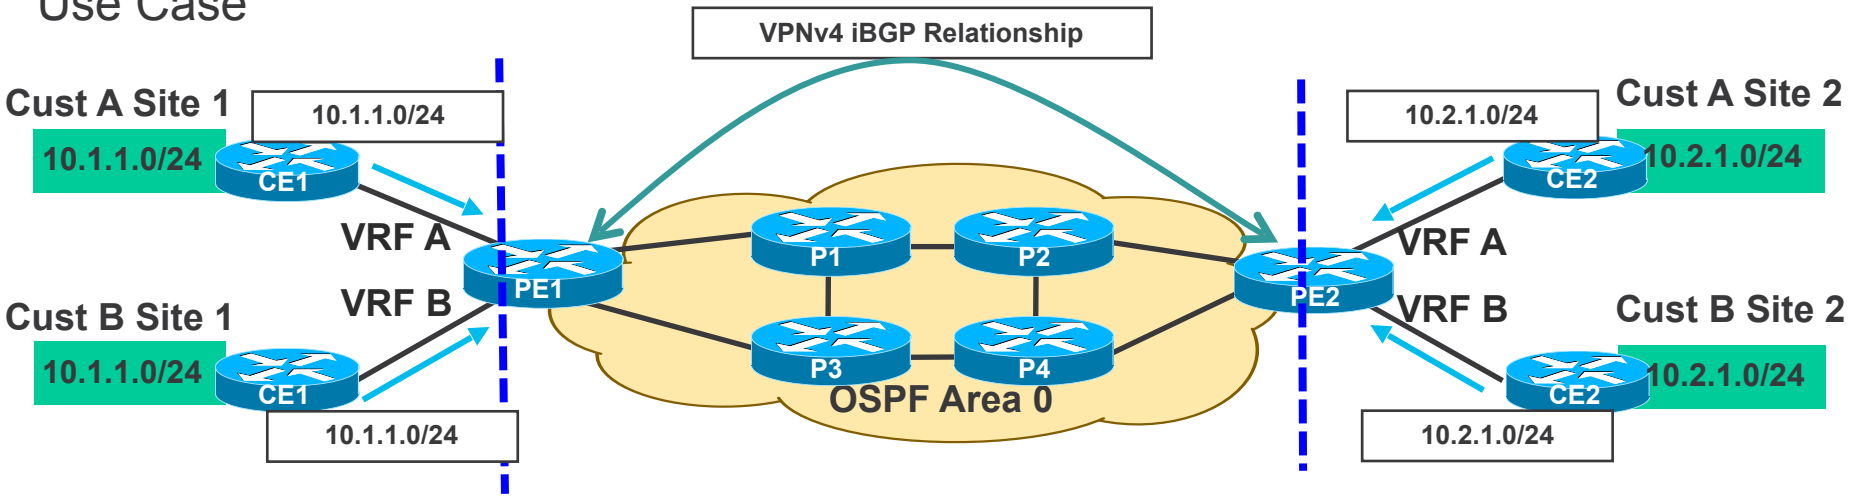

1. PE routers service multiple customers

VPNv4 prefixes are the combination of a 64-bit RD and a 32-bit IPv4 prefix. VPNv4 prefixes are 96-bits in length

Why an RD and VPNv4 Address?

Use Case

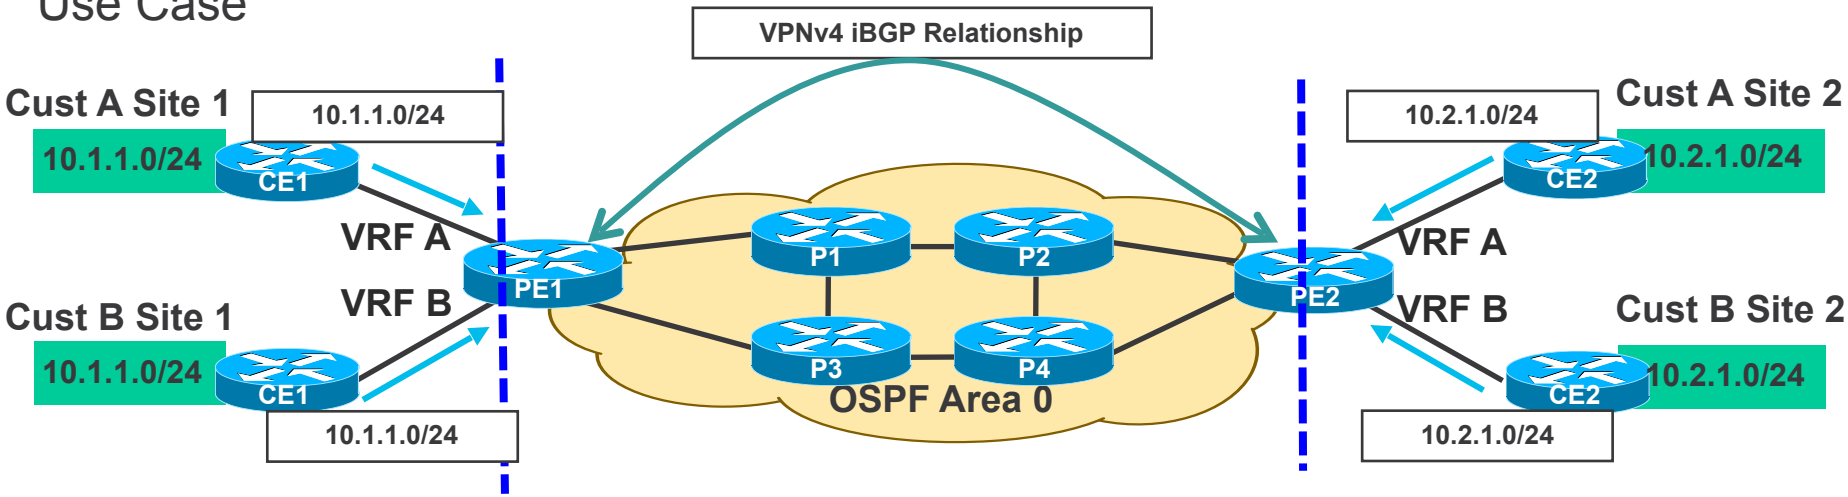

1. PE routers service multiple customers
2. Once PE redistributes customer routes into MP-BGP, they must be unique

VPNv4 prefixes are the combination of a 64-bit RD and a 32-bit IPv4 prefix. VPNv4 prefixes are 96-bits in length

Why an RD and VPNv4 Address?

Use Case

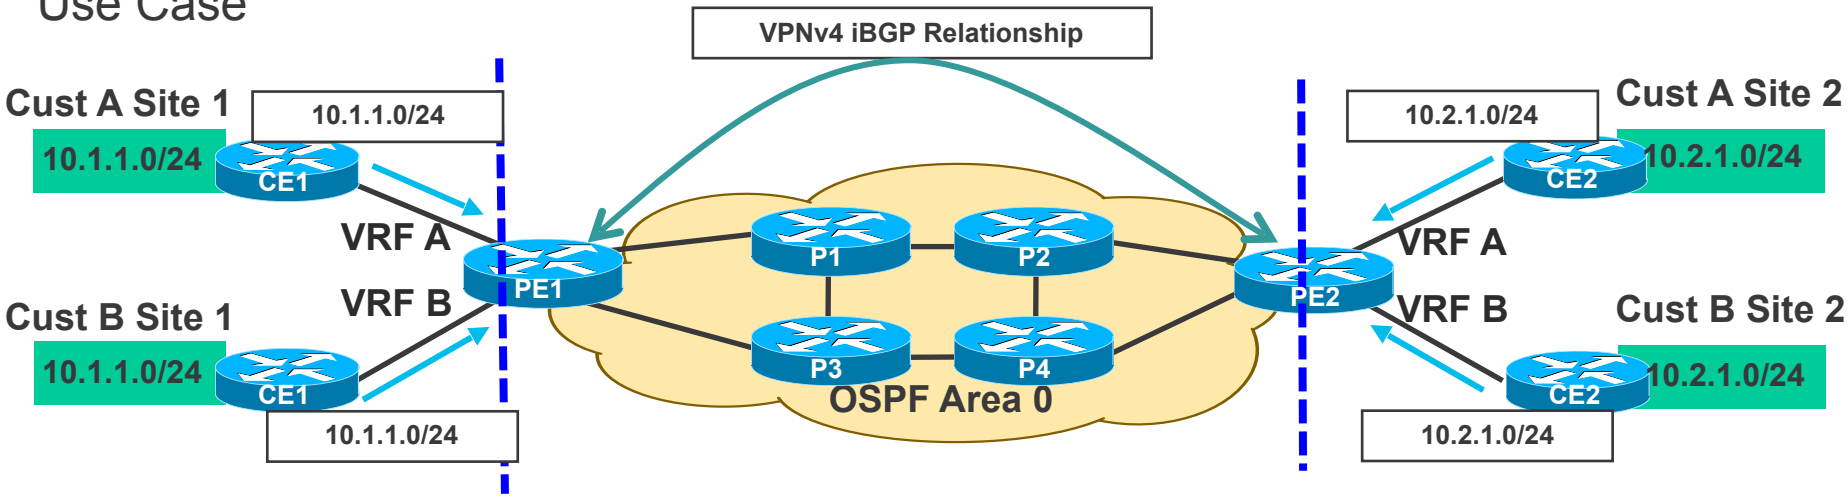

1. PE routers service multiple customers
2. Once PE redistributes customer routes into MP-BGP, they must be unique
3. RD is prepended to each prefix to make routes unique

VPNv4 prefixes are the combination of a 64-bit RD and a 32-bit IPv4 prefix. VPNv4 prefixes are 96-bits in length

Why an RD and VPNv4 Address?

Use Case

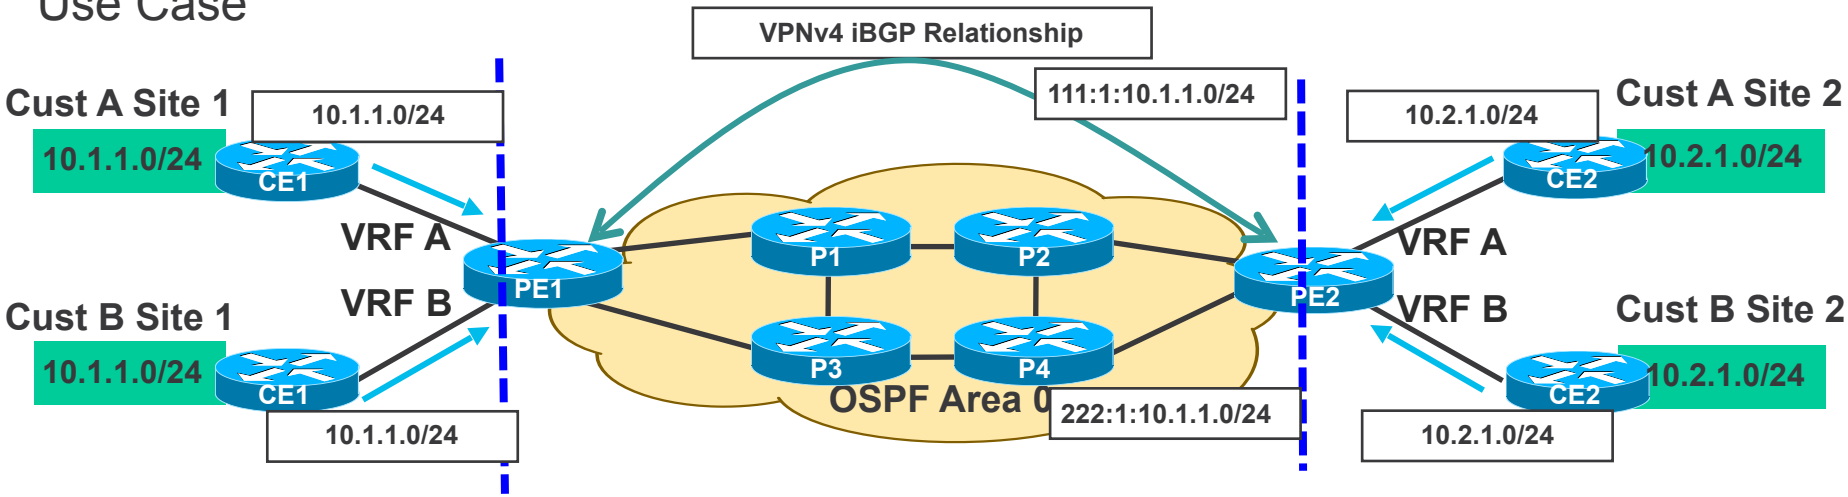

1. PE routers service multiple customers
2. Once PE redistributes customer routes into MP-BGP, they must be unique
3. RD is prepended to each prefix to make routes unique

VPNv4 prefixes are the combination of a 64-bit RD and a 32-bit IPv4 prefix. VPNv4 prefixes are 96-bits in length

Why an RD and VPNv4 Address?

Use Case

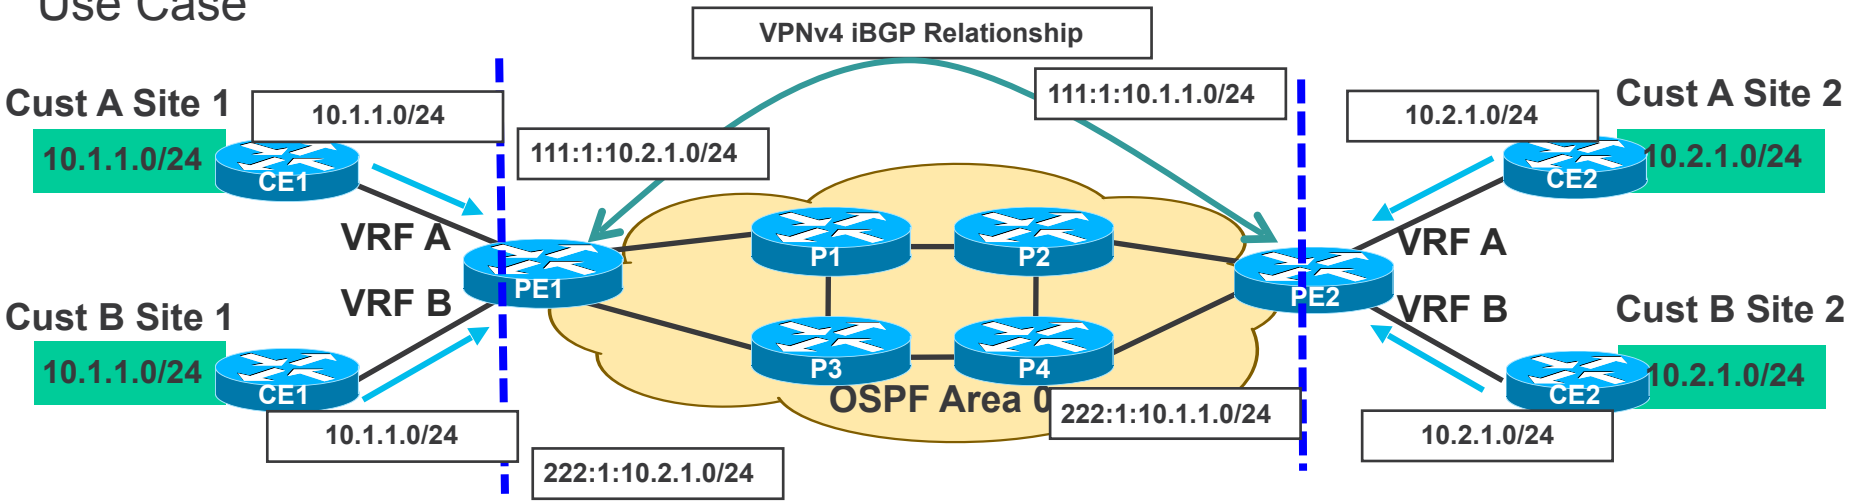

1. PE routers service multiple customers
2. Once PE redistributes customer routes into MP-BGP, they must be unique
3. RD is prepended to each prefix to make routes unique

VPNv4 prefixes are the combination of a 64-bit RD and a 32-bit IPv4 prefix. VPNv4 prefixes are 96-bits in length

Why are Route Targets Important?

Use Case

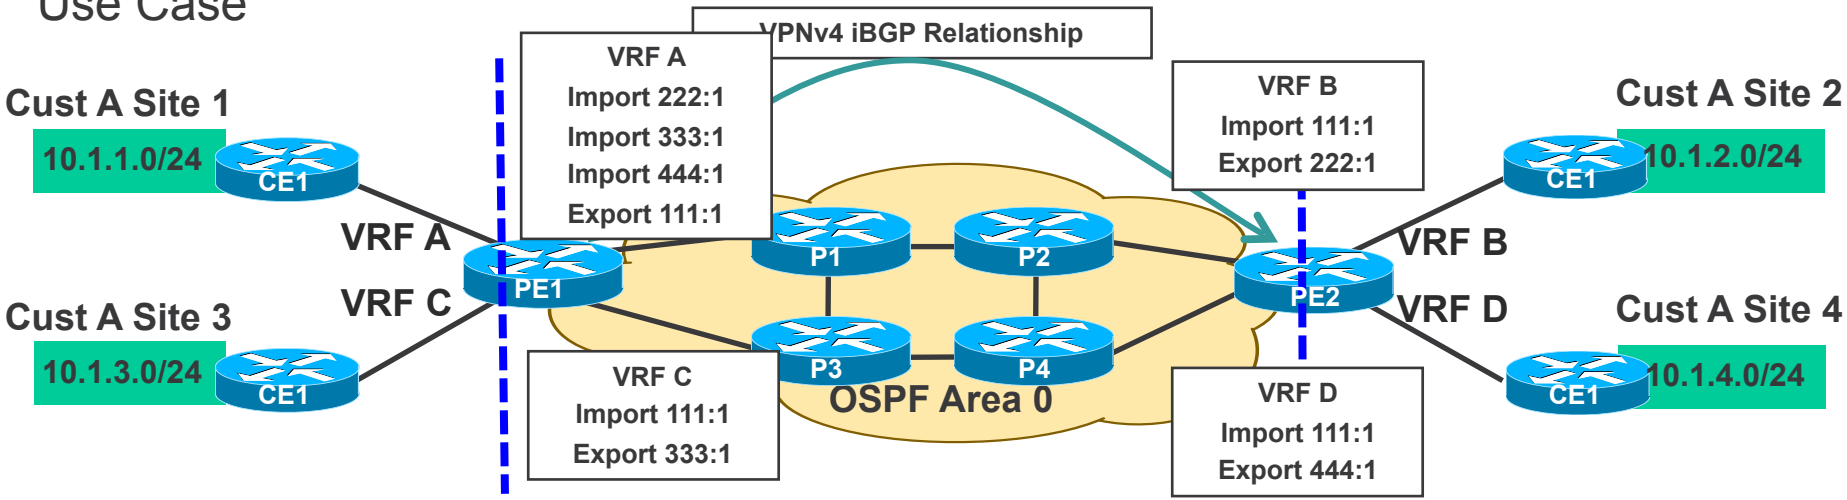

Route Targets are a 64-bit value and are carried in BGP as an extended community

Why are Route Targets Important?

Use Case

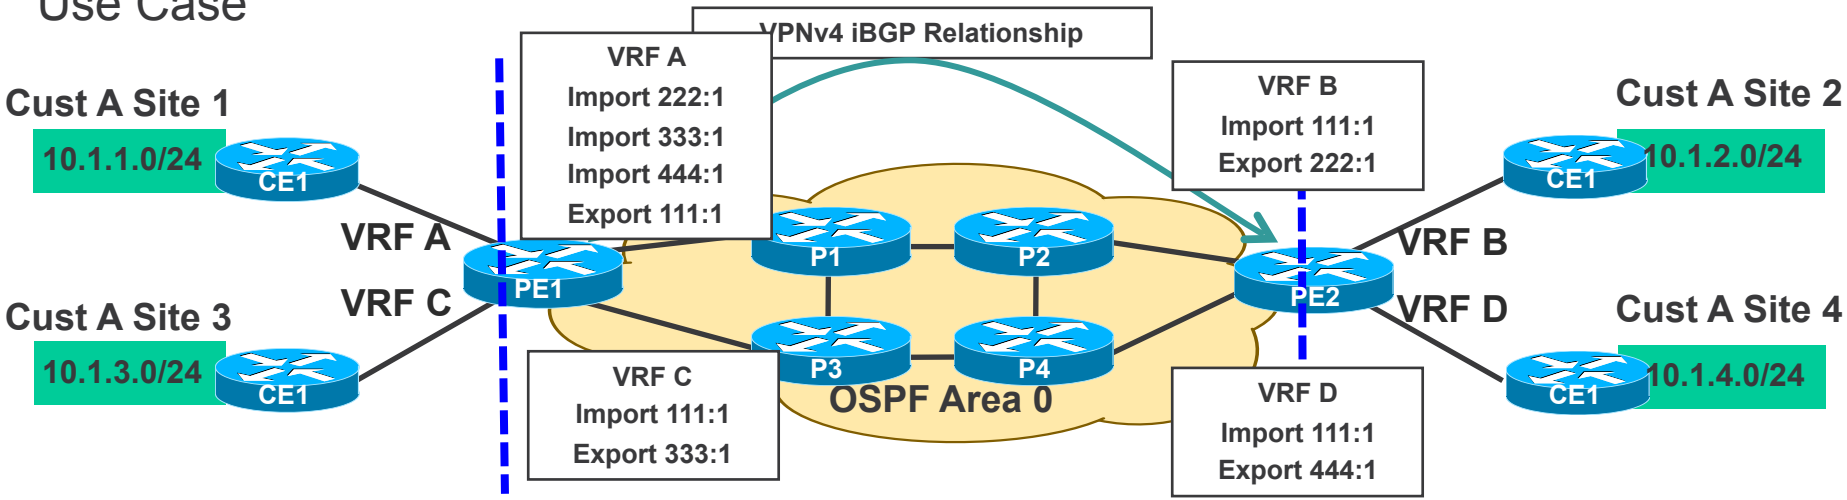

1. Route Targets dictate which VRF will receive what routes

Route Targets are a 64-bit value and are carried in BGP as an extended community

Why are Route Targets Important?

Use Case

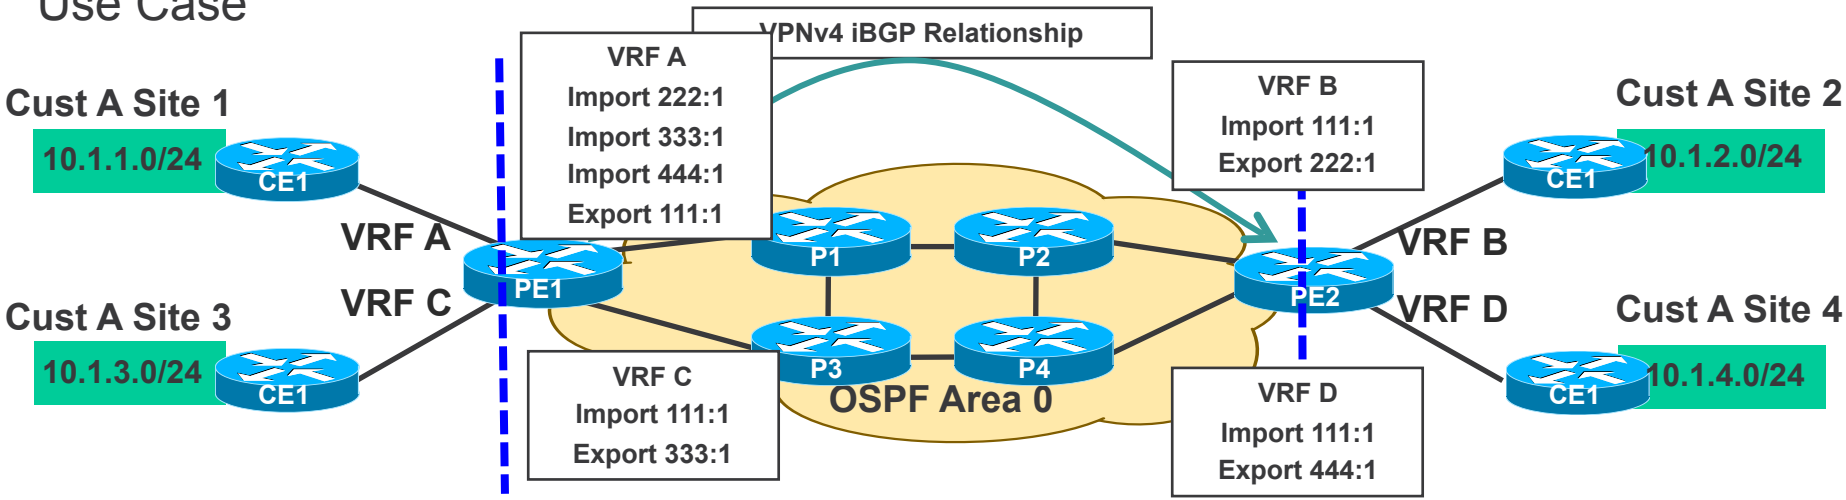

1. Route Targets dictate which VRF will receive what routes
2. Can be used to allow specific sites access to centralised services

Route Targets are a 64-bit value and are carried in BGP as an extended community

Why are Route Targets Important?

Use Case

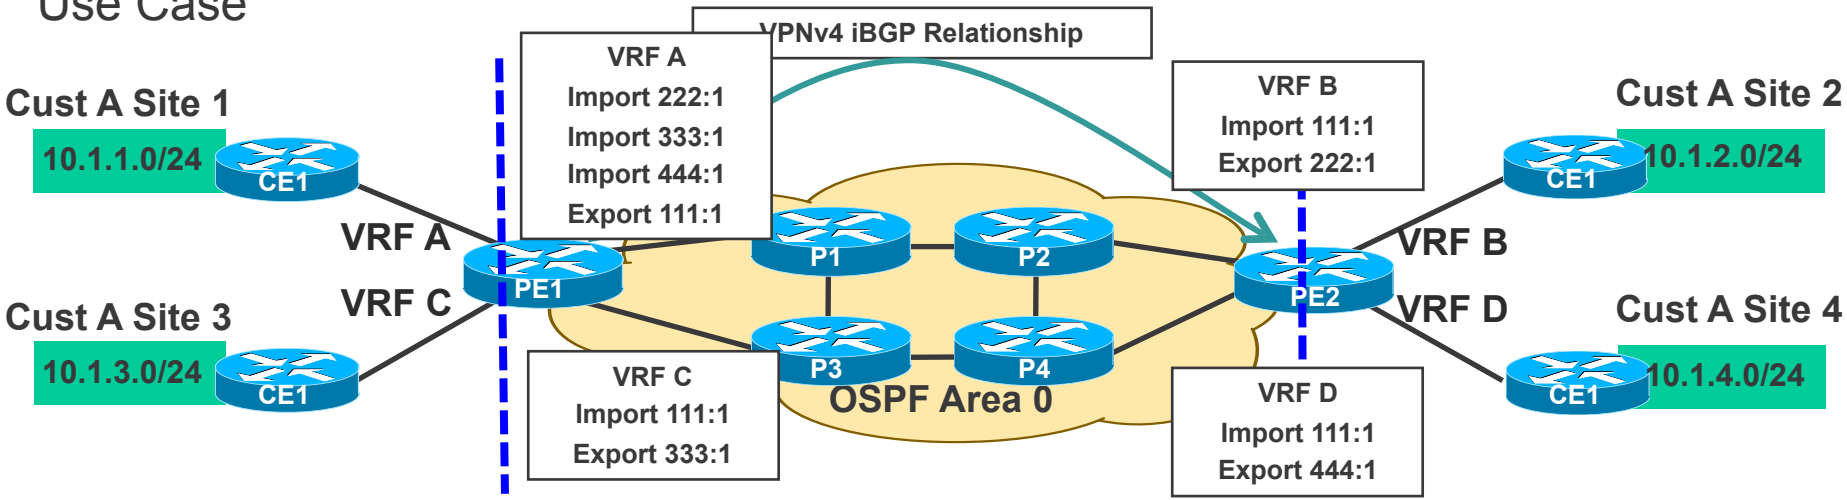

1. Route Targets dictate which VRF will receive what routes
2. Can be used to allow specific sites access to centralised services
3. Cust A Site 2, Site 3 and Site 4 will not be able to exchange routes with each other

Route Targets are a 64-bit value and are carried in BGP as an extended community

MPLS VPN and MP-BGP

Command Line Interface (CLI) Review

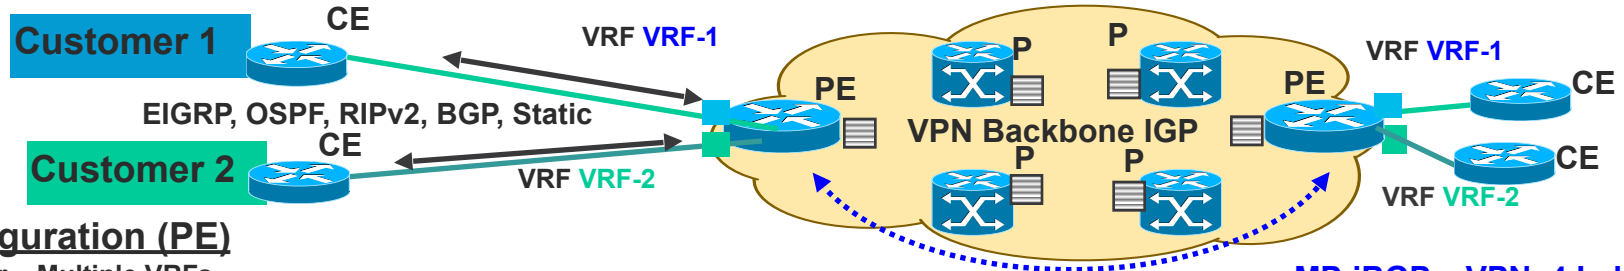

VRF Configuration (PE)

```
! PE Router – Multiple VRFs
ip vrf VRF-1
 rd 65100:10
 route-target import 65102:10
 route-target export 65102:10
ip vrf VRF-2
 rd 65100:20
 route-target import 65102:20
 route-target export 65102:20
!
Interface FastEthernet0/1.10
 ip vrf forwarding VRF-1
Interface FastEthernet0/1.20
 ip vrf forwarding VRF-2
```

MP-iBGP Configuration (PE)

MP-iBGP – VPNv4 Label Exchange

```
! PE router
router bgp 65102
 no bgp default ipv4-unicast
 neighbor 2.2.2.2 remote-as 65102
 !
 address-family vpnv4
 neighbor 2.2.2.2 activate
 neighbor 2.2.2.2 send-community extended
 exit-address-family
 !
 address-family ipv4 vrf VRF-1
 redistribute rip
 exit-address-family
```

Live Exploration

MPLS VPN Technology Summary

MPLS VPN Connection Model

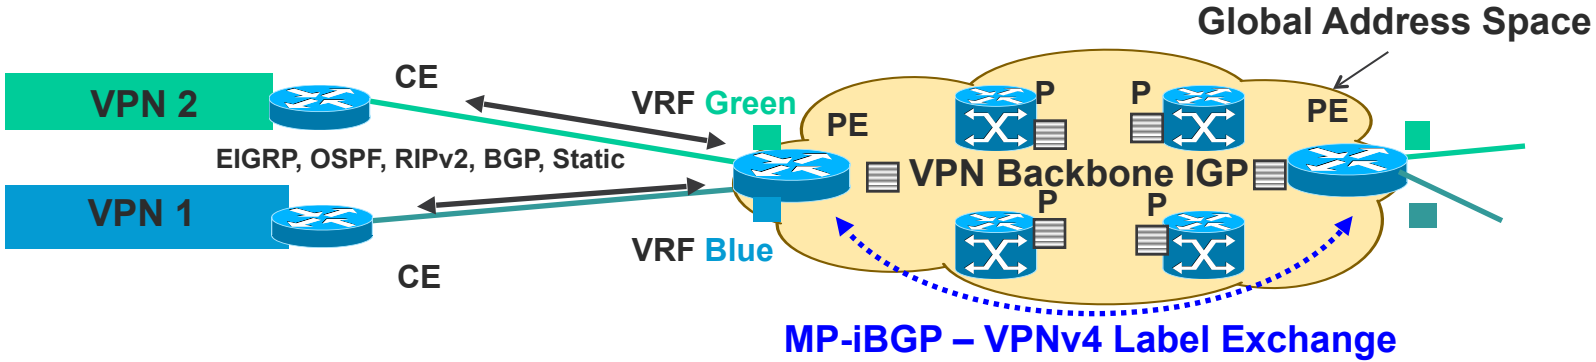

MPLS VPN Technology Summary

MPLS VPN Connection Model

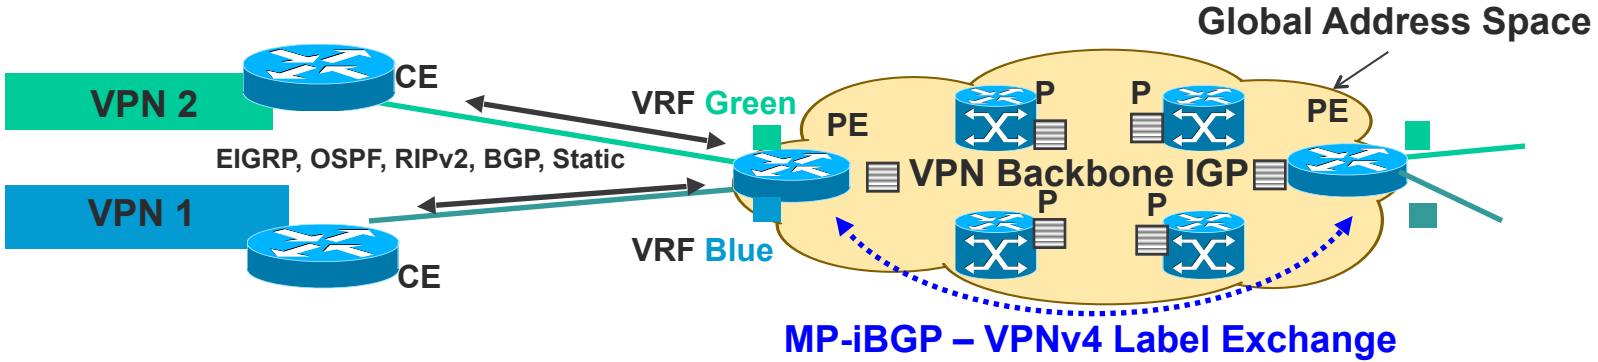

CE Routers

- Sends routes to PE
  - via (static, RIP, BGP, EIGRP, OSPF)
  - Unlabeled IP packets

MPLS VPN Technology Summary

MPLS VPN Connection Model

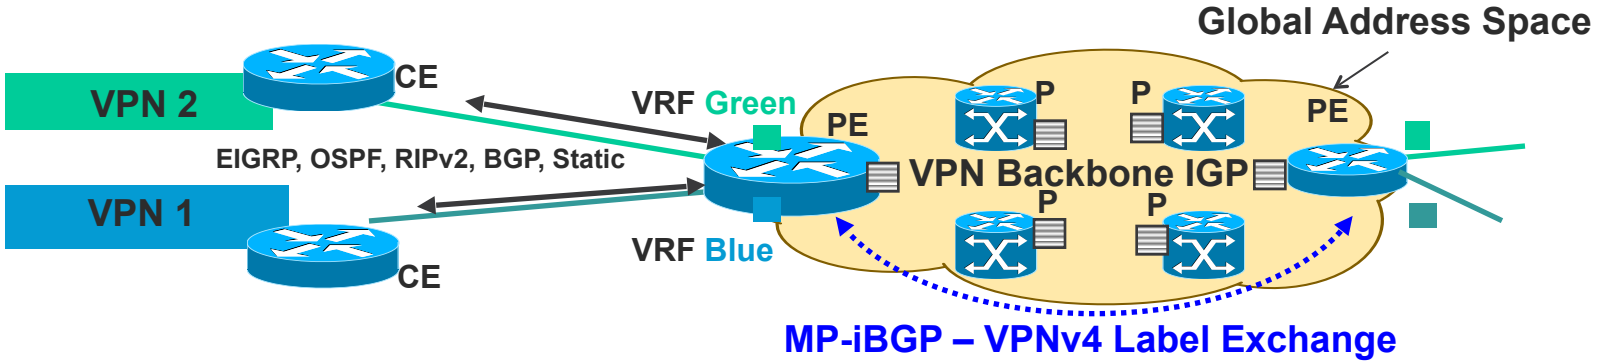

CE Routers

- Sends routes to PE
  - via (static, RIP, BGP, EIGRP, OSPF)
  - Unlabeled IP packets

PE Routers

MPLS VPN Technology Summary

MPLS VPN Connection Model

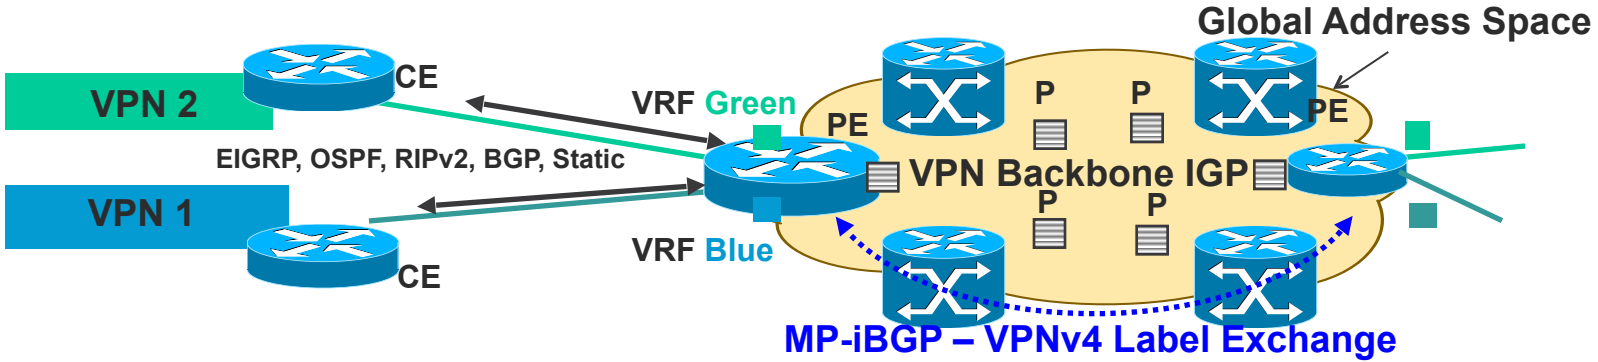

CE Routers

- Sends routes to PE
  - via (static, RIP, BGP, EIGRP, OSPF)
  - Unlabeled IP packets

PE Routers

- MPLS Edge routers with VRF(s)
- MPLS forwarding to P routers
- IGP/BGP – IP to CE routers
- Distributes VPN information through MP-BGP to other PE routers with VPNv4 addresses, extended community, VPN labels
- Push labels onto incoming IP packets

P Routers

- P routers are in the core of the MPLS cloud
- P routers do not need to run BGP
- Do not have knowledge of VPNs
- Switch packets based on labels (swap/pop) not IP

Closing Thoughts

- Break MPLS into smaller, more manageable chunks to accelerate learning
- Leverage current routing protocol knowledge learning PE-CE VRF routing
- MP-BGP and traditional IPv4 BGP configuration is very similar
- If routes are not present on CE routers check route-target import/export, communities and redistribution between IPv4 VRF address-families under IGP and BGP
- If routes are present but you are having problems with reachability, check MPLS configuration
- Remember on PE devices you are living in a VRF world (Ping, Traceroute etc.)
- HAVE FUN !!!!! Remember, it's a journey not a destination!

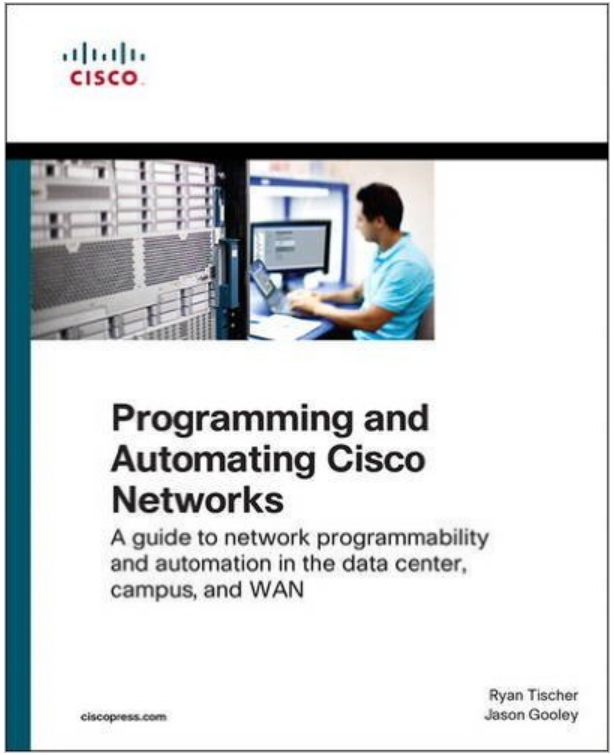

Order a copy of Programming and Automating Cisco Networks here:
http://www.amazon.com/Programming-Automating-Cisco-Networks-programmability/dp/1587144654/ref=sr_1_1?ie=UTF8&qid=1450796338&sr=8-1&keywords=Jason+Gooley

Continue Your Education

- Demos in the Cisco campus
- Walk-in Self-Paced Labs
- Lunch & Learn
- Meet the Engineer 1:1 meetings
- Related sessions

What Is MPLS?

What Is MPLS?

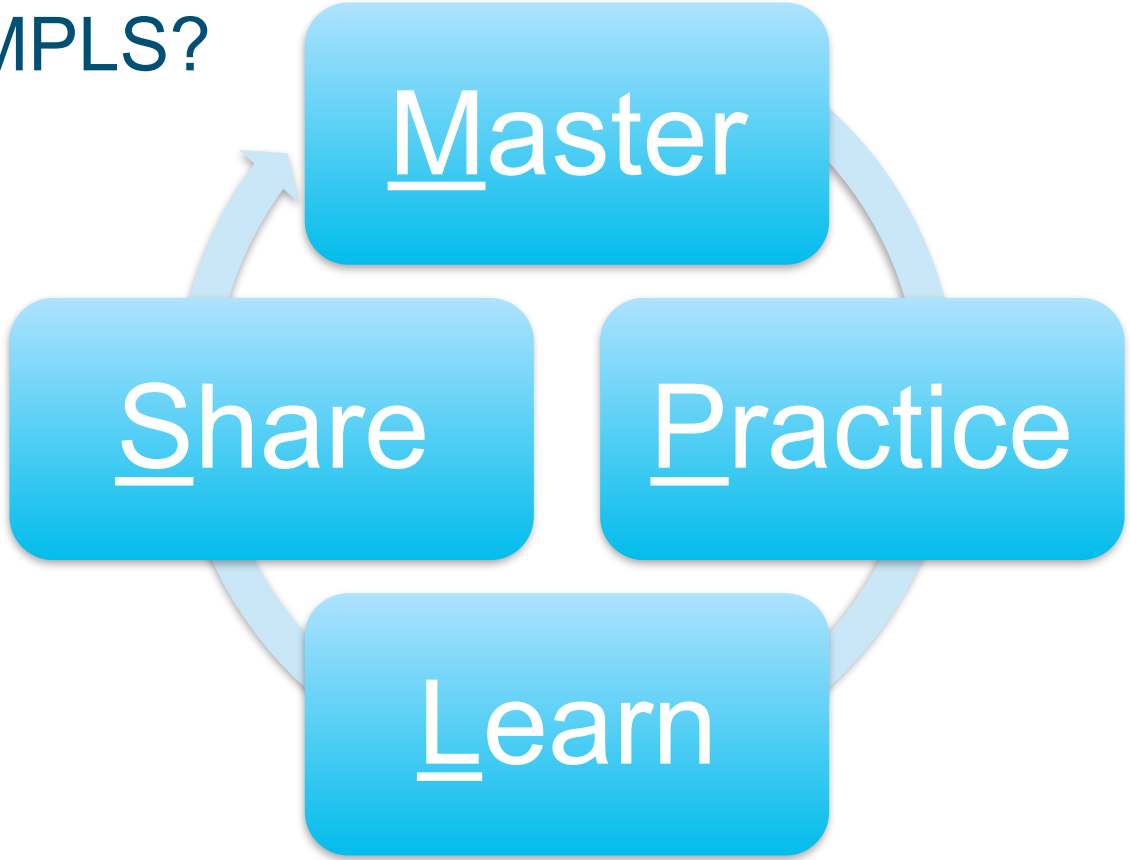

Live Exploration Diagrams

VRF-Lite
Sub-interface Example

Per VRF:
- Virtual Routing Table
- Virtual Forwarding Table
- Locally Significant

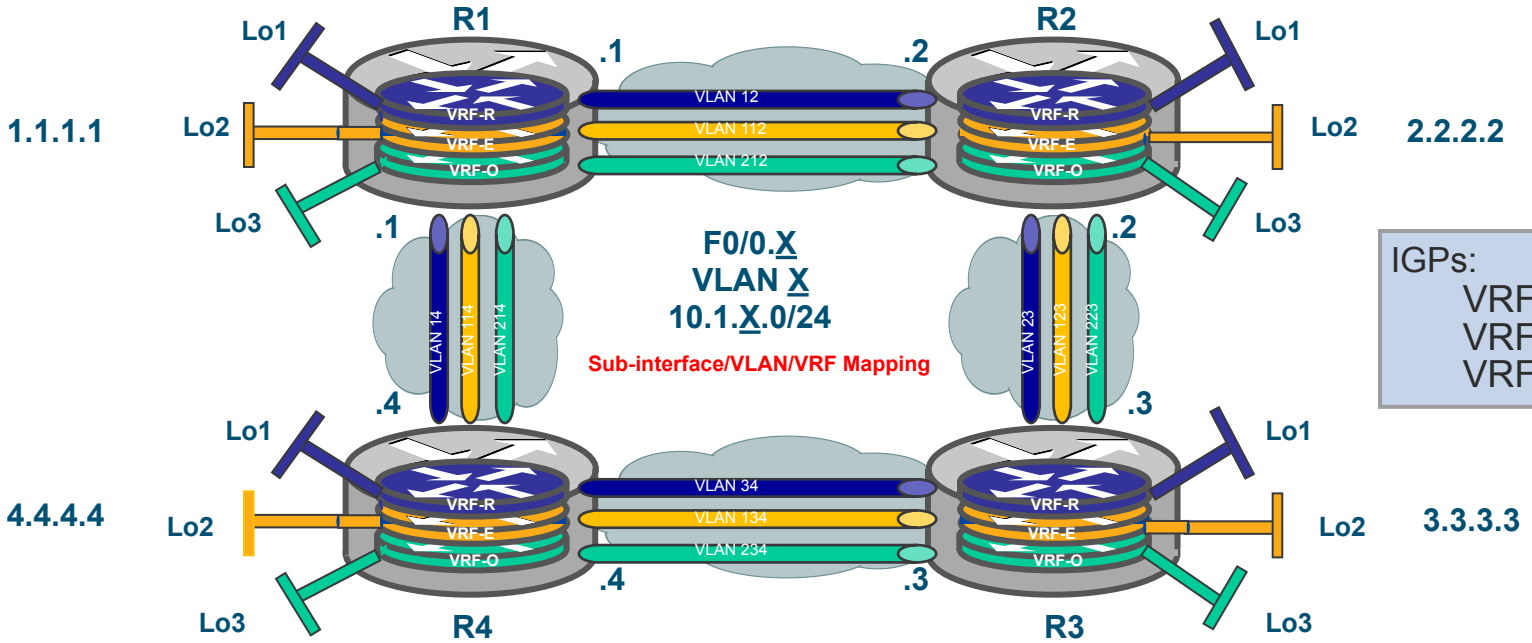

IGPs:
- VRF-R = RIP
- VRF-E = EIGRP
- VRF-O = OSPF

No Sub-interface Support? No Problem!

GRE Example

- VRF-Lite can also leverage GRE tunnels as a segmentation technology
- Each VRF uses a unique GRE tunnel
- GRE tunnel interface is “VRF aware”

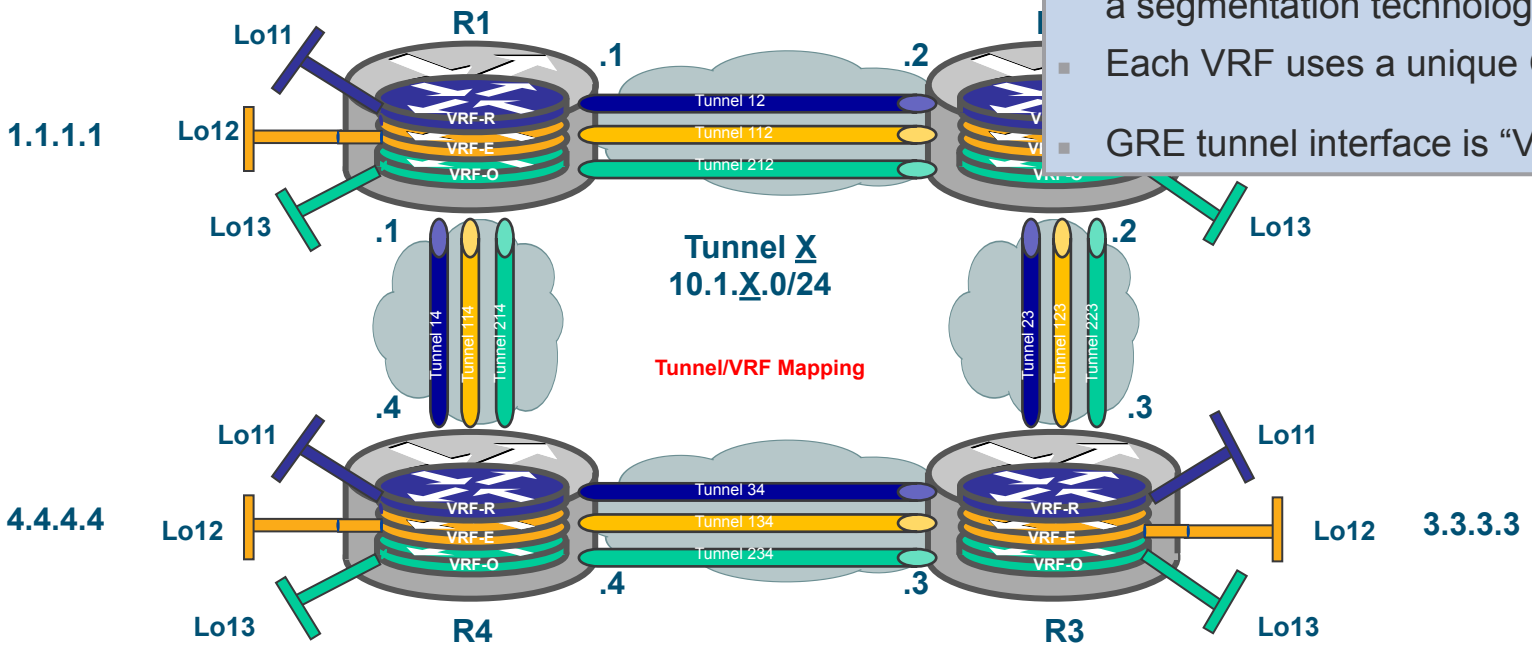

Configuration Note: Each GRE Tunnel Could Require Unique Source/Destination IP (Platform Dependent)

Layer 2 Serial Link? No Problem?

Back-to-Back Frame Relay Example

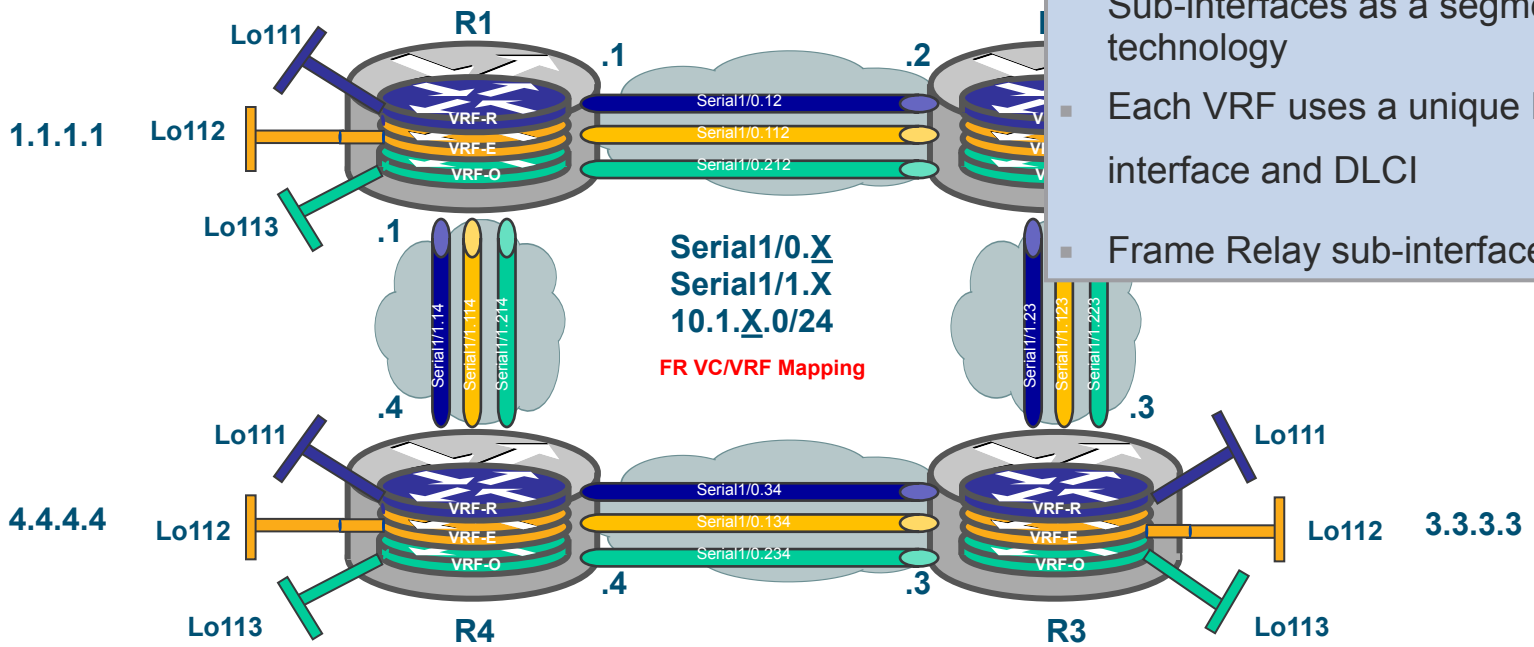

- VRF-Lite can also leverage Frame Relay Sub-interfaces as a segmentation technology
- Each VRF uses a unique Frame-Relay sub-interface and DLCI
- Frame Relay sub-interface is “VRF aware”

Configuration Note: Leveraging Back-to-Back Frame-Relay Configuration

Multiprotocol BGP (MP-BGP)

Bringing It All Together

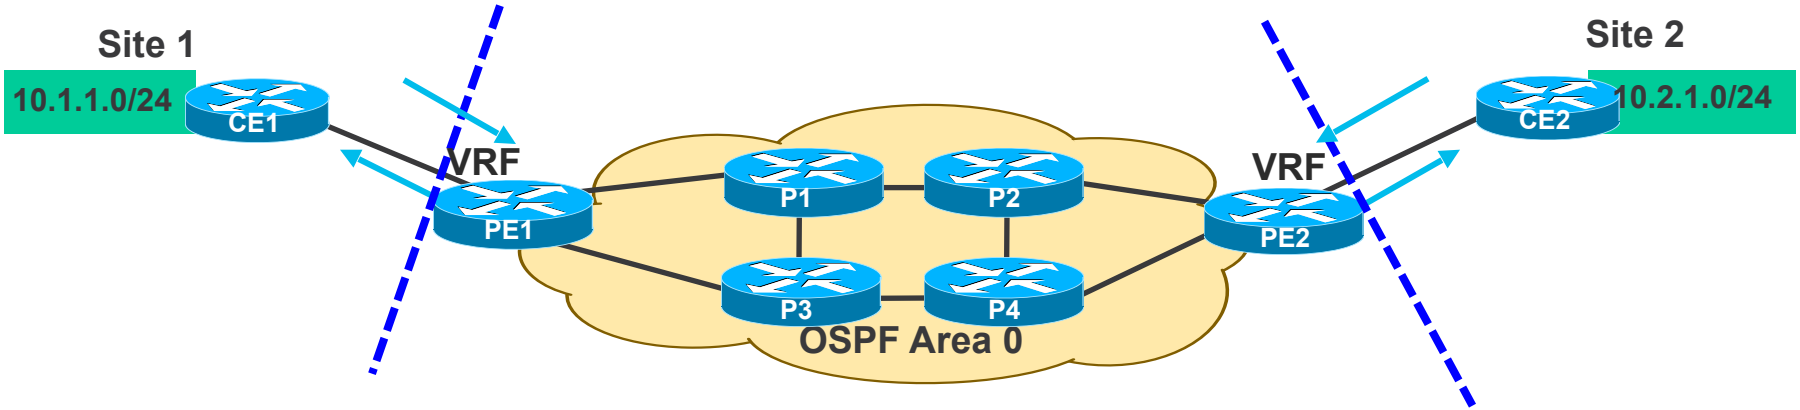

Multiprotocol BGP (MP-BGP)

Bringing It All Together

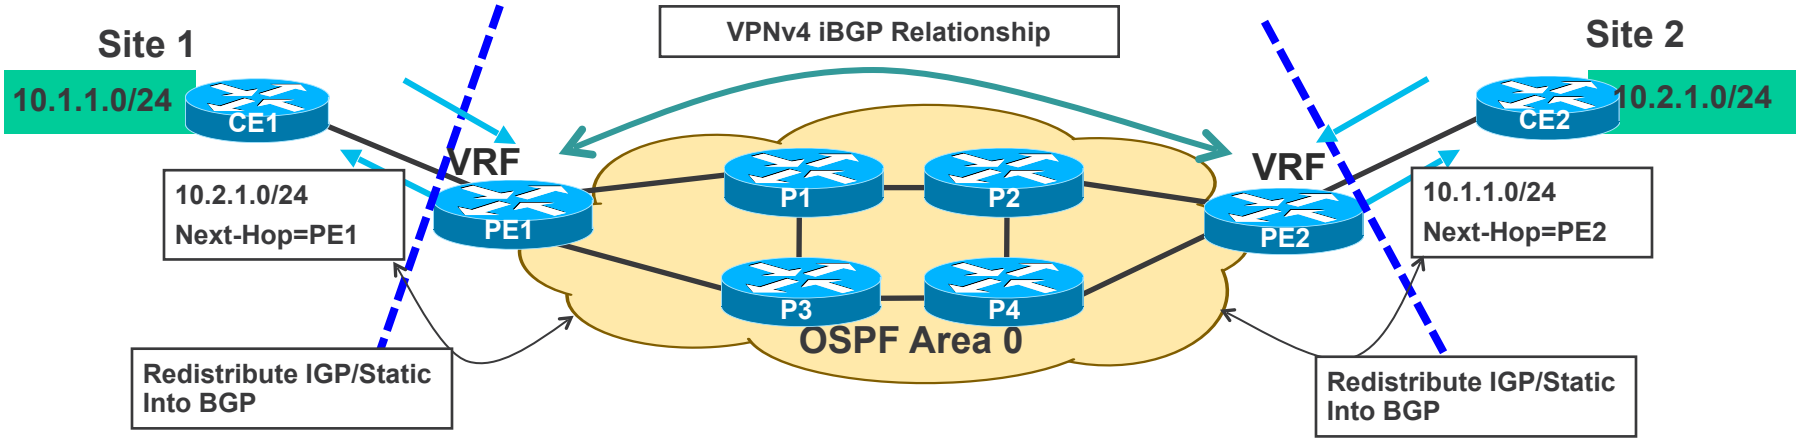

Multiprotocol BGP (MP-BGP)

Bringing It All Together

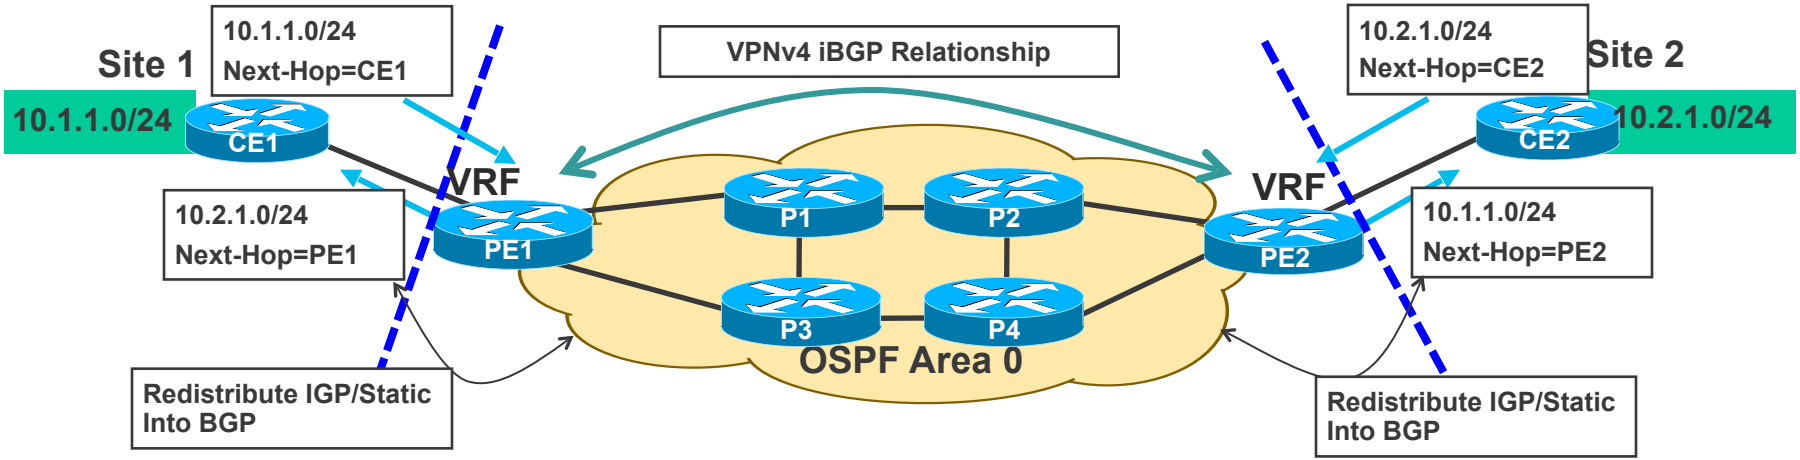

1. PE receives an IPv4 update on a VRF interface (eBGP/OSPF/RIP/EIGRP)

Multiprotocol BGP (MP-BGP)

Bringing It All Together

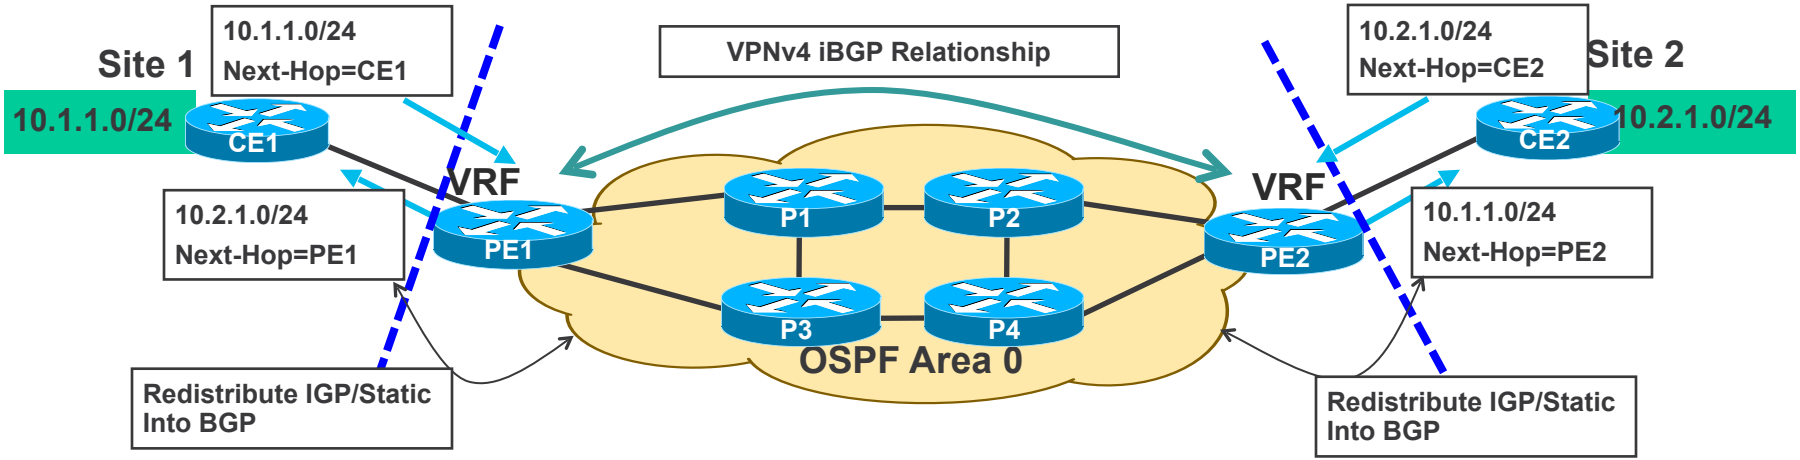

1. PE receives an IPv4 update on a VRF interface (eBGP/OSPF/RIP/EIGRP)
2. PE translates it into VPNv4 address (96-bit address) (64-bit RD + 32 bit IPv4 address)

Multiprotocol BGP (MP-BGP)

Bringing It All Together

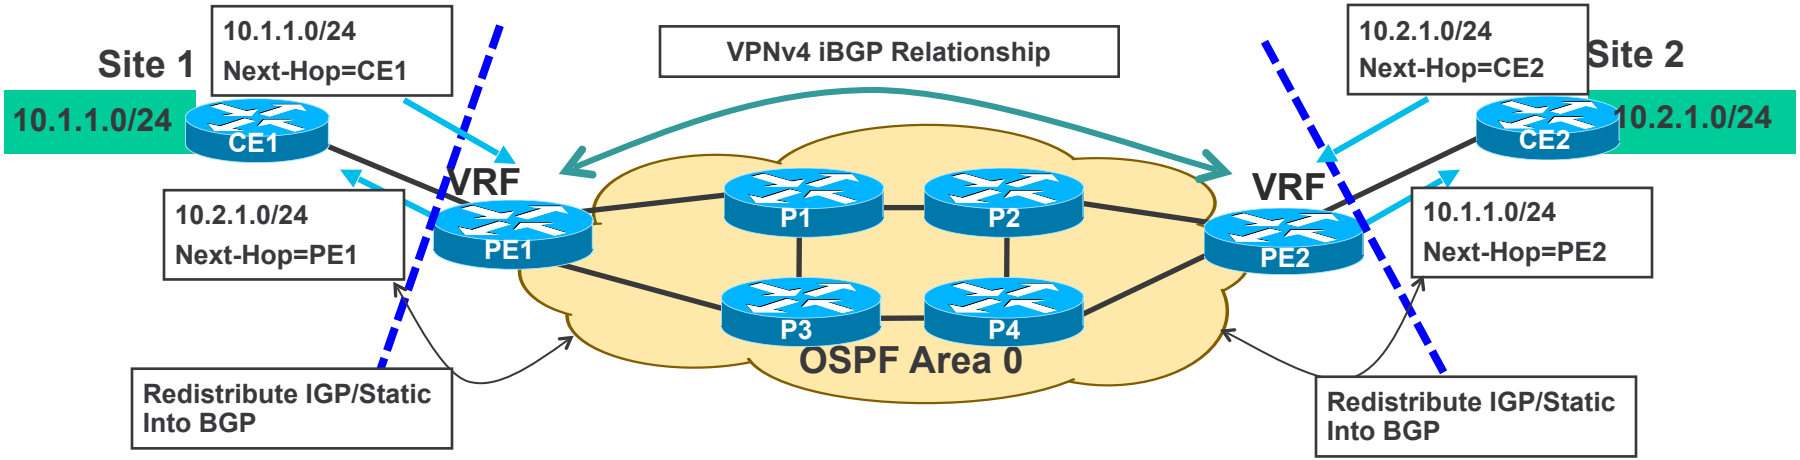

1. PE receives an IPv4 update on a VRF interface (eBGP/OSPF/RIP/EIGRP)
2. PE translates it into VPNv4 address (96-bit address) (64-bit RD + 32 bit IPv4 address)
   - Assigns an RT per VRF configuration

Multiprotocol BGP (MP-BGP)

Bringing It All Together

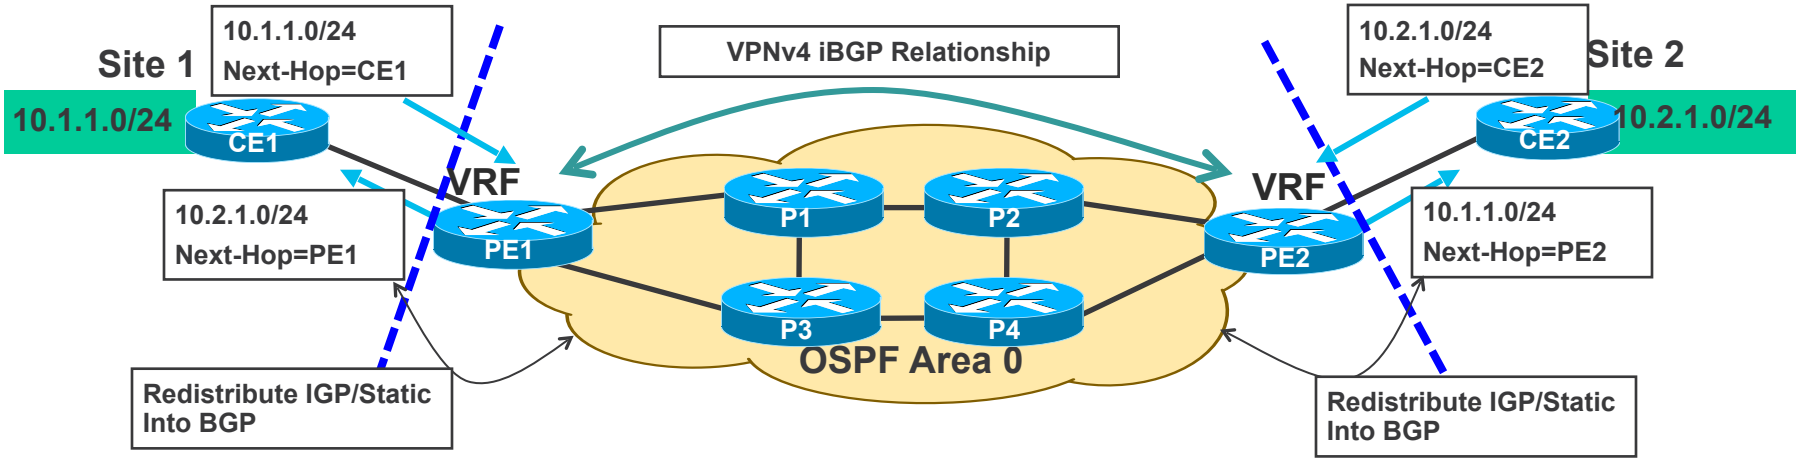

1. PE receives an IPv4 update on a VRF interface (eBGP/OSPF/RIP/EIGRP)
2. PE translates it into VPNv4 address (96-bit address) (64-bit RD + 32 bit IPv4 address)
   - Assigns an RT per VRF configuration
   - Rewrites next-hop attribute to itself

Multiprotocol BGP (MP-BGP)

Bringing It All Together

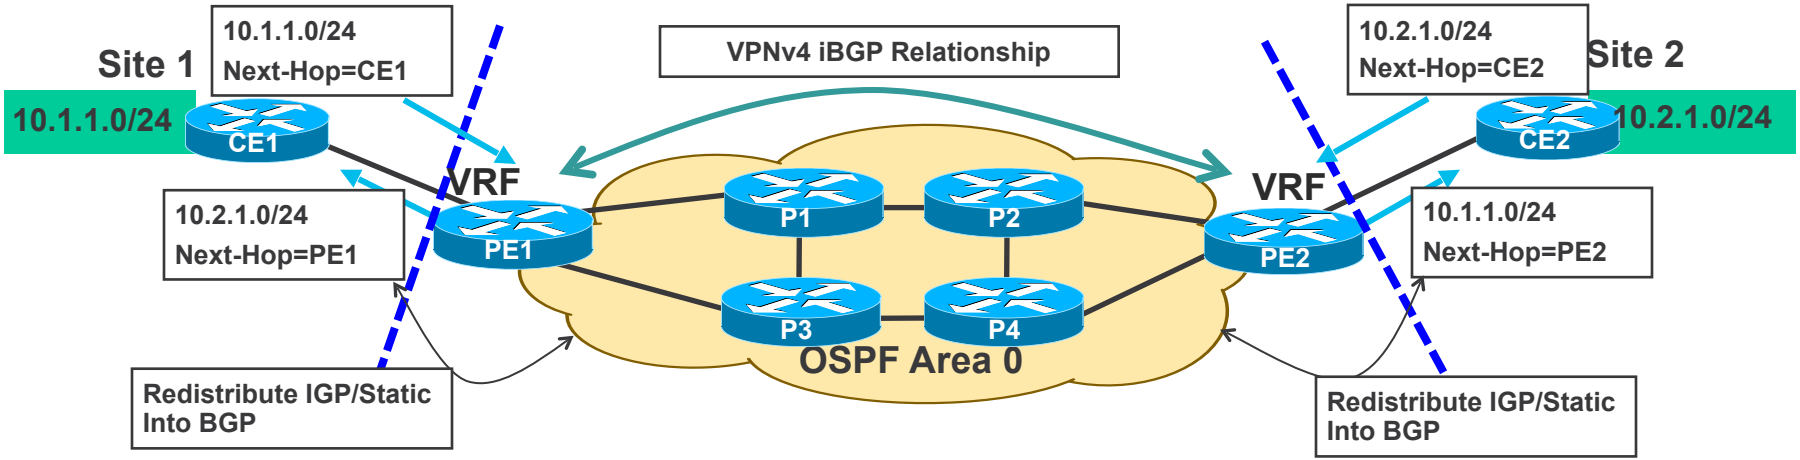

1. PE receives an IPv4 update on a VRF interface (eBGP/OSPF/RIP/EIGRP)
2. PE translates it into VPNv4 address (96-bit address) (64-bit RD + 32 bit IPv4 address)
   - Assigns an RT per VRF configuration
   - Rewrites next-hop attribute to itself
   - Assigns a label based on VRF and/or interface

Multiprotocol BGP (MP-BGP)

Bringing It All Together

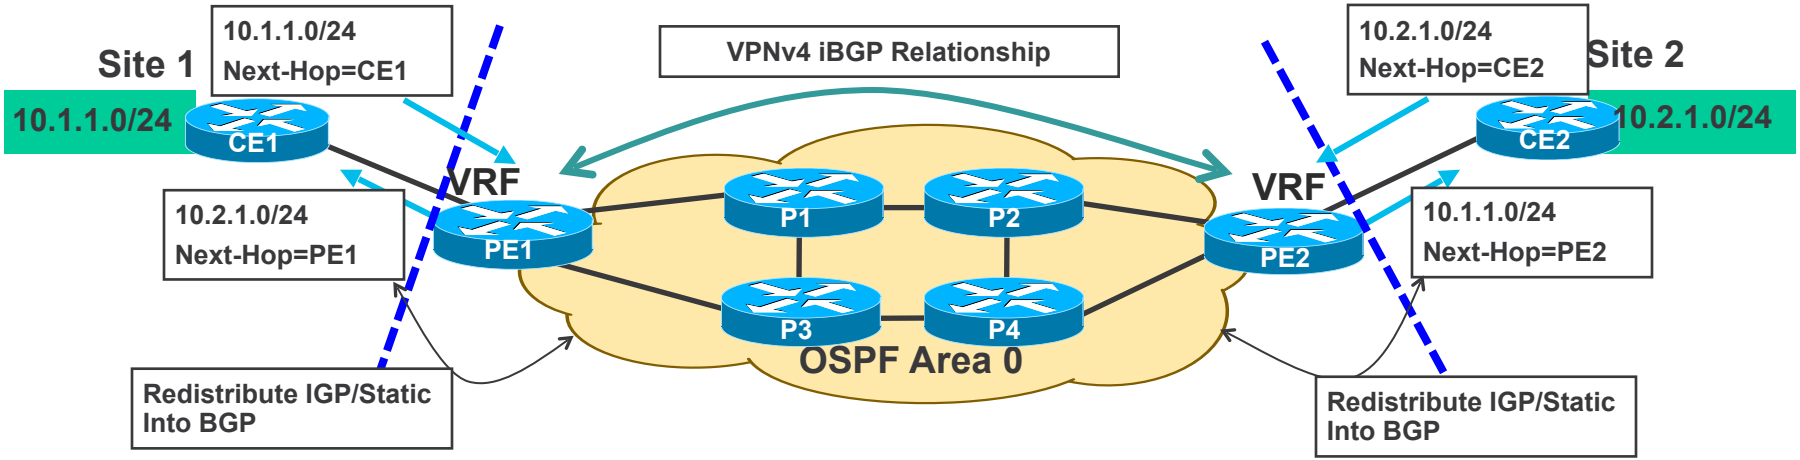

1. PE receives an IPv4 update on a VRF interface (eBGP/OSPF/RIP/EIGRP)
2. PE translates it into VPNv4 address (96-bit address) (64-bit RD + 32 bit IPv4 address)
   - Assigns an RT per VRF configuration
   - Rewrites next-hop attribute to itself
   - Assigns a label based on VRF and/or interface
3. PE sends MP-iBGP update to other PE routers

R&S Related Cisco Education Offerings

Course	Description	Cisco Certification
CCIE R&S Advanced Workshops (CIERS-1 & CIERS-2) plus Self Assessments, Workbooks & Labs	Expert level trainings including: instructor led workshops, self assessments, practice labs and CCIE Lab Builder to prepare candidates for the CCIE R&S practical exam.	CCIE® Routing & Switching
• Implementing Cisco IP Routing v2.0 • Implementing Cisco IP Switched Networks V2.0 • Troubleshooting and Maintaining Cisco IP Networks v2.0	Professional level instructor led trainings to prepare candidates for the CCNP R&S exams (ROUTE, SWITCH and TSHOOT). Also available in self study eLearning formats with Cisco Learning Labs.	CCNP® Routing & Switching
Interconnecting Cisco Networking Devices: Part 2 (or combined)	Configure, implement and troubleshoot local and wide-area IPv4 and IPv6 networks. Also available in self study eLearning format with Cisco Learning Lab.	CCNA® Routing & Switching
Interconnecting Cisco Networking Devices: Part 1	Installation, configuration, and basic support of a branch network. Also available in self study eLearning format with Cisco Learning Lab.	CCENT® Routing & Switching

For more details, please visit: http://learningnetwork.cisco.com
Questions? Visit the Learning@Cisco Booth

Network Programmability Cisco Education Offerings

Course	Description	Cisco Certification
Developing with Cisco Network Programmability (NPDEV)	Provides Application Developers with comprehensive curriculum to develop infrastructure programming skills; Addresses needs of software engineers who automate network infrastructure and/or utilise APIs and toolkits to interface with SDN controllers and individual devices	Cisco Network Programmability Developer (NPDEV) Specialist Certification
Designing and Implementing Cisco Network Programmability (NPDESI)	Provides network engineers with comprehensive soup-to-nuts curriculum to develop and validate automation and programming skills; Directly addresses the evolving role of network engineers towards more programmability, automation and orchestration	Cisco Network Programmability Design and Implementation (NPDESI) Specialist Certification
Programming for Network Engineers (PRNE)	Learn the fundamentals of Python programming – within the context of performing functions relevant to network engineers. Use Network Programming to simplify or automate tasks	Recommended pre-requisite for NPDESI and NPDEV Specialist Certifications
Cisco Digital Network Architecture Implementation Essentials (DNAIE)	This training provides students with the guiding principles and core elements of Cisco's Digital Network Architecture (DNA) architecture and its solution components including; APIC-EM, NFV, Analytics, Security and Fabric.	None

For more details, please visit: http://learningnetwork.cisco.com
Questions? Visit the Learning@Cisco Booth

Q & A

Complete Your Online Session Evaluation

- Give us your feedback and receive a **Cisco Live 2018 Cap** by completing the overall event evaluation and 5 session evaluations.
- All evaluations can be completed via the Cisco Live Mobile App.

Don’t forget: Cisco Live sessions will be available for viewing on demand after the event at www.CiscoLive.com/Global.

Ciscolive!

Thank you


cisco
You're
Cisco live!